

Elektra
0.8.4

Generated by Doxygen 1.8.1.2

Tue Jan 14 2014 15:37:46

Contents

1	The Elektra API	1
1.1	Elektra Initiative Overview	1
1.2	API docu	1
1.3	Using the Elektra Library	1
1.4	Elektra API	1
1.5	Key Names and Namespaces	2
1.6	Rules for Key Names	2
1.7	Backend Overview	3
1.8	Glossary	3
2	Deprecated List	5
3	Module Index	7
3.1	Modules	7
4	Namespace Index	9
4.1	Namespace List	9
5	Data Structure Index	11
5.1	Data Structures	11
6	Module Documentation	13
6.1	KDB	13
6.1.1	Detailed Description	13
6.1.2	Enumeration Type Documentation	14
6.1.2.1	option_t	14
6.1.3	Function Documentation	15
6.1.3.1	kdbClose	15
6.1.3.2	kdbGet	16
6.1.4	Example	16
6.1.5	Details	16
6.1.6	Updating	16
6.1.6.1	kdbOpen	17

6.1.6.2	kdbSet	17
6.1.7	parentKey	18
6.1.8	Update	18
6.1.9	Error Situations	18
6.2	Key	20
6.2.1	Detailed Description	21
6.2.2	Enumeration Type Documentation	21
6.2.2.1	keyswitch_t	21
6.2.3	Function Documentation	21
6.2.3.1	keyClear	21
6.2.3.2	keyCopy	22
6.2.3.3	keyDecRef	23
6.2.3.4	keyDel	24
6.2.3.5	keyDup	24
6.2.3.6	keyGetRef	25
6.2.3.7	keyIncRef	25
6.2.3.8	keyNew	26
6.3	Basic Methods	29
6.4	Meta Info Manipulation Methods	30
6.4.1	Detailed Description	30
6.4.2	Function Documentation	31
6.4.2.1	keyCopyAllMeta	31
6.4.2.2	keyCopyMeta	32
6.4.2.3	keyCurrentMeta	33
6.4.2.4	keyGetATime	33
6.4.2.5	keyGetCTime	34
6.4.2.6	keyGetGID	34
6.4.3	GID	34
6.4.3.1	keyGetMeta	35
6.4.3.2	keyGetMode	35
6.4.3.3	keyGetMTime	36
6.4.3.4	keyGetUID	36
6.4.4	UID	36
6.4.4.1	keyNextMeta	37
6.4.4.2	keyRewindMeta	37
6.4.4.3	keySetATime	38
6.4.4.4	keySetCTime	38
6.4.4.5	keySetDir	39
6.4.4.6	keySetGID	39
6.4.4.7	keySetMeta	40

6.4.4.8	keySetMode	40
6.4.5	Modes	41
6.4.5.1	keySetMTime	41
6.4.5.2	keySetUID	42
6.5	Name Manipulation Methods	43
6.5.1	Detailed Description	43
6.5.2	Function Documentation	44
6.5.2.1	keyAddBaseName	44
6.5.2.2	keyBaseName	44
6.5.2.3	keyGetBaseName	45
6.5.2.4	keyGetBaseNameSize	45
6.5.2.5	keyGetFullName	46
6.5.2.6	keyGetFullNameSize	46
6.5.2.7	keyGetName	47
6.5.2.8	keyGetNameSize	47
6.5.2.9	keyGetOwner	47
6.5.2.10	keyGetOwnerSize	48
6.5.2.11	keyName	49
6.5.2.12	keyOwner	49
6.5.2.13	keySetBaseName	50
6.5.2.14	keySetName	50
6.5.2.15	keySetOwner	51
6.6	KeySet	52
6.6.1	Detailed Description	52
6.6.2	Function Documentation	53
6.6.2.1	ksAppend	53
6.6.2.2	ksAppendKey	53
6.6.2.3	ksCopy	54
6.6.2.4	ksCurrent	55
6.6.2.5	ksCut	55
6.6.2.6	ksDel	55
6.6.2.7	ksDup	56
6.6.2.8	ksGetCursor	56
6.6.3	Read ahead	56
6.6.4	Restoring state	57
6.6.4.1	ksGetSize	57
6.6.4.2	ksHead	58
6.6.4.3	ksLookup	58
6.6.5	Introduction	58
6.6.6	Usage	58

6.6.6.1	KDB_O_NOALL	59
6.6.6.2	KDB_O_POP	59
6.6.6.3	ksLookupByName	60
6.6.7	Cascading	60
6.6.8	Full Search	60
6.6.8.1	ksNew	61
6.6.8.2	ksNext	62
6.6.8.3	ksPop	63
6.6.8.4	ksRewind	63
6.6.8.5	ksSetCursor	64
6.6.8.6	ksTail	64
6.7	Methods for Making Tests	66
6.7.1	Detailed Description	66
6.7.2	Function Documentation	66
6.7.2.1	keyCmp	66
6.7.2.2	keyCompare	67
6.7.2.3	keyIsBelow	68
6.7.2.4	keyIsBinary	69
6.7.2.5	keyIsDir	69
6.7.2.6	keyIsDirectBelow	70
6.7.2.7	keyIsInactive	71
6.7.2.8	keyIsString	71
6.7.2.9	keyIsSystem	72
6.7.2.10	keyIsUser	72
6.7.2.11	keyNeedSync	72
6.7.2.12	keyRel	73
6.8	Value Manipulation Methods	75
6.8.1	Detailed Description	75
6.8.2	Function Documentation	75
6.8.2.1	keyComment	75
6.8.2.2	keyGetBinary	76
6.8.2.3	keyGetComment	77
6.8.3	Comments	77
6.8.3.1	keyGetCommentSize	77
6.8.3.2	keyGetString	78
6.8.3.3	keyGetValueSize	79
6.8.3.4	keySetBinary	79
6.8.3.5	keySetComment	80
6.8.3.6	keySetString	80
6.8.3.7	keyString	81

6.8.3.8	keyValue	81
6.8.4	String Handling	82
6.8.5	Binary Data Handling	82
6.9	Plugins	84
6.9.1	Detailed Description	84
6.9.2	Function Documentation	85
6.9.2.1	docClose	85
6.9.2.2	docError	85
6.9.2.3	docGet	85
6.9.3	Introduction	86
6.9.3.1	Storage Plugins	86
6.9.3.2	Filter Plugins	86
6.9.4	Conditions	87
6.9.5	Updating	87
6.9.5.1	docOpen	88
6.9.5.2	docSet	89
6.9.5.3	ELEKTRA_PLUGIN_EXPORT	90
7	Namespace Documentation	93
7.1	kdb Namespace Reference	93
7.1.1	Detailed Description	93
7.1.2	Function Documentation	93
7.1.2.1	operator<<	93
7.1.2.2	operator<<	94
7.1.2.3	operator>>	94
7.1.2.4	operator>>	94
8	Data Structure Documentation	95
8.1	kdb::KDB Class Reference	95
8.1.1	Detailed Description	95
8.1.2	Constructor & Destructor Documentation	96
8.1.2.1	KDB	96
8.1.2.2	KDB	97
8.1.2.3	~KDB	98
8.1.3	Member Function Documentation	99
8.1.3.1	close	99
8.1.3.2	get	99
8.1.3.3	get	100
8.1.3.4	open	100
8.1.3.5	set	101
8.1.3.6	set	102

8.2	kdb::Key Class Reference	102
8.2.1	Detailed Description	105
8.2.2	Constructor & Destructor Documentation	105
8.2.2.1	Key	105
8.2.2.2	Key	106
8.2.2.3	Key	106
8.2.2.4	Key	106
8.2.2.5	Key	106
8.2.2.6	Key	109
8.2.2.7	Key	112
8.2.2.8	~Key	115
8.2.3	Member Function Documentation	115
8.2.3.1	addBaseName	115
8.2.3.2	clear	116
8.2.3.3	copy	116
8.2.3.4	copyAllMeta	118
8.2.3.5	copyMeta	118
8.2.3.6	currentMeta	119
8.2.3.7	dup	120
8.2.3.8	get	121
8.2.3.9	getBaseName	121
8.2.3.10	getBaseNameSize	122
8.2.3.11	getBinary	122
8.2.3.12	getBinarySize	124
8.2.3.13	getDirName	124
8.2.3.14	getFullName	124
8.2.3.15	getFullNameSize	125
8.2.3.16	getFunc	125
8.2.3.17	getKey	125
8.2.3.18	getMeta	125
8.2.3.19	getName	127
8.2.3.20	getNameSize	127
8.2.3.21	getReferenceCounter	128
8.2.3.22	getString	128
8.2.3.23	getStringSize	129
8.2.3.24	getValue	129
8.2.3.25	isBelow	130
8.2.3.26	isBelowOrSame	130
8.2.3.27	isBinary	130
8.2.3.28	isDirectBelow	131

8.2.3.29	isInactive	131
8.2.3.30	isString	131
8.2.3.31	isSystem	132
8.2.3.32	isUser	132
8.2.3.33	isValid	132
8.2.3.34	nextMeta	132
8.2.3.35	operator bool	133
8.2.3.36	operator!=	133
8.2.3.37	operator*	135
8.2.3.38	operator++	135
8.2.3.39	operator++	135
8.2.3.40	operator+=	136
8.2.3.41	operator+=	136
8.2.3.42	operator--	136
8.2.3.43	operator--	137
8.2.3.44	operator-=	137
8.2.3.45	operator-=	137
8.2.3.46	operator<	138
8.2.3.47	operator<=	139
8.2.3.48	operator=	140
8.2.3.49	operator=	140
8.2.3.50	operator=	140
8.2.3.51	operator=	140
8.2.3.52	operator==	140
8.2.3.53	operator>	142
8.2.3.54	operator>=	143
8.2.3.55	release	144
8.2.3.56	rewindMeta	144
8.2.3.57	set	145
8.2.3.58	setBaseName	146
8.2.3.59	setBinary	146
8.2.3.60	setMeta	147
8.2.3.61	setName	147
8.2.3.62	setString	148
8.3	kdb::KeySet Class Reference	149
8.3.1	Detailed Description	150
8.3.2	Constructor & Destructor Documentation	151
8.3.2.1	KeySet	151
8.3.2.2	KeySet	152
8.3.2.3	KeySet	152

8.3.2.4	KeySet	153
8.3.2.5	KeySet	153
8.3.2.6	~KeySet	153
8.3.3	Member Function Documentation	153
8.3.3.1	append	153
8.3.3.2	append	154
8.3.3.3	clear	155
8.3.3.4	copy	155
8.3.3.5	current	156
8.3.3.6	cut	156
8.3.3.7	dup	157
8.3.3.8	getCursor	157
8.3.3.9	getKeySet	157
8.3.3.10	head	158
8.3.3.11	lookup	158
8.3.3.12	lookup	158
8.3.3.13	next	159
8.3.3.14	operator=	159
8.3.3.15	pop	159
8.3.3.16	release	160
8.3.3.17	rewind	160
8.3.3.18	setCursor	160
8.3.3.19	setKeySet	161
8.3.3.20	size	161
8.3.3.21	tail	161

Chapter 1

The Elektra API

1.1 Elektra Initiative Overview

Elektra provides a universal and secure framework to store configuration parameters in a global, hierarchical key database. The core is a small library implemented in C. The plugin-based framework fulfills many configuration-related tasks to avoid any unnecessary code duplication across applications while it still allows the core to stay without any external dependency. Elektra abstracts from cross-platform-related issues with a consistent API, and allows applications to be aware of other applications' configurations, leveraging easy application integration.

See the website for more information <http://www.libelektra.org>

1.2 API docu

This document occupies with the API implementation, documentation, internals and plugins. On the one hand it gives an overview and an introduction for developers using Elektra, on the other hand it gives an informal description what methods must and may provide to allow an alternative implementation of the API.

The current version (for stable releases) of this document can be found at <http://doc.libelektra.org/api/current/html>

The latest version (from git master) of this document can be found at <http://doc.libelektra.org/api/latest/html>

1.3 Using the Elektra Library

A C or C++ source file that wants to use Elektra should include:

```
#include <kdb.h>
```

To link an executable with the Elektra library, the correct way is to use the `pkg-config` tool:

```
bash$ cc `pkg-config --libs elektra` -o myapp myapp.c
```

1.4 Elektra API

The API was written in pure C because Elektra was designed to be useful even for the most basic system programs, which are all made in C. Also, being C, bindings to other languages can appear easily.

The API follows an Object Oriented design, and there are 3 main classes as shown by the figure:

Some general things you can do with each class are:

KDB

- [The four lowlevel functions](#)
- [Open](#) and [Close](#) the Database
- [Get](#) and [Set](#) [KeySet](#) in the Database
- See [class documentation](#) for more

Key

- Get and Set key properties like [name](#) , [string](#) or [binary](#) values, [permissions](#) , [changed time](#) and [comment](#)
- Test if it is a [user/](#) or [system/](#) key, etc
- See [class documentation](#) for more

KeySet

- Linked list of Key objects
- Append [a single key](#) or an entire [KeySet](#)
- [Work with](#) its [internal cursor](#)
- See [class documentation](#) for more

1.5 Key Names and Namespaces

There are 2 trees of keys: `system` and `user`

- The "system" Subtree It is provided to store system-wide configuration keys, that is, configurations that daemons and system services will use. But all other programs will also try to fetch system keys to have a fallback managed by the distributor or admin when the user does not have configuration for its own.
- The "user" Subtree Used to store user-specific configurations, like the personal settings of a user to certain programs. The user subtree will always be favoured if present (except for security concerns the user subtree may not be considered). See [Cascading](#) in the documentation of [ksLookupByName\(\)](#) how the selection of user and system keys works.

1.6 Rules for Key Names

When using Elektra to store your application's configuration and state, please keep in mind the following rules:

- You are not allowed to create keys right under `system` or `user`. They are reserved for more generic purposes.
- The keys for your application, called say *MyApp*, should be created under `system/sw/MyApp/current` and `user/sw/MyApp/current`
- `current` is the default configuration profile, users may symlink to the profile they want.
- That means you just need to [kdbGet\(\)](#) `system/sw/MyApp/profile` and `user/sw/MyApp/profile` and then [ksLookupByName\(\)](#) in `/sw/MyApp/profile` while profile defaults to `current`, but may be changed by the user or admin. See [Cascading](#) to learn more about that feature.

1.7 Backend Overview

The core of elektra does not store configuration itself to the harddisk. Instead this work is delegated to backends.

If you want to develop a backend, you should already have some experience with Elektra from the user point of view. You should be familiar with the data structures: [Key](#) and [KeySet](#). Then you can start reading about Backends, which are composed out of [Plugins](#).

1.8 Glossary

- **pop**, used in [ksPop\(\)](#) and [KDB_O_POP](#) means to remove a key from a keyset.
- **delete**, or abbr. **del**, used in [keyDel\(\)](#), [ksDel\(\)](#) and [KDB_O_DEL](#) means to free a key or keyset. The memory can be used for something else afterwards.
- **remove** means that the key/value information in the physical database will be removed permanently.

Chapter 2

Deprecated List

Global `KDB_O_SORT`

dont use

Global `keyGetATime` (const Key *key)

This API is obsolete.

Global `keyGetCTime` (const Key *key)

This API is obsolete.

Global `keyGetGID` (const Key *key)

This API is obsolete.

Global `keyGetMode` (const Key *key)

This API is obsolete.

Global `keyGetMTime` (const Key *key)

This API is obsolete.

Global `keyGetUID` (const Key *key)

This API is obsolete.

Global `keySetATime` (Key *key, time_t atime)

This API is obsolete.

Global `keySetCTime` (Key *key, time_t ctime)

This API is obsolete.

Global `keySetDir` (Key *key)

This API is obsolete.

Global `keySetGID` (Key *key, gid_t gid)

This API is obsolete.

Global `keySetMode` (Key *key, mode_t mode)

This API is obsolete. It is only a mapping to `keySetMeta(key, "mode", str)` which should be preferred.

Global `keySetMTime` (Key *key, time_t mtime)

This API is obsolete.

Global `keySetUID` (Key *key, uid_t uid)

This API is obsolete.

Chapter 3

Module Index

3.1 Modules

Here is a list of all modules:

KDB	13
Key	20
Basic Methods	29
Meta Info Manipulation Methods	30
Methods for Making Tests	66
Name Manipulation Methods	43
Value Manipulation Methods	75
KeySet	52
Plugins	84

Chapter 4

Namespace Index

4.1 Namespace List

Here is a list of all documented namespaces with brief descriptions:

[kdb](#)

See `examples/cpp_example_userexception.cpp` for how to use `USER_DEFINED_EXEPTIONS`

[93](#)

Chapter 5

Data Structure Index

5.1 Data Structures

Here are the data structures with brief descriptions:

kdb::KDB	Opens the session with the Key database.	95
kdb::Key	A Key is the essential class that encapsulates key name , value and metainfo	102
kdb::KeySet	A keyset holds together a set of keys	149

Chapter 6

Module Documentation

6.1 KDB

General methods to access the Key database.

Enumerations

- enum `option_t` {
 [KDB_O_NONE](#) = 0, [KDB_O_DEL](#) = 1, [KDB_O_POP](#) = 1<<1, [KDB_O_NODIR](#) = 1<<2,
 [KDB_O_DIRONLY](#) = 1<<3, [KDB_O_NOREMOVE](#) = 1<<6, [KDB_O_REMOVEONLY](#) = 1<<7, [KDB_O_I-](#)
 [NACTIVE](#) = 1<<8,
 [KDB_O_SYNC](#) = 1<<9, [KDB_O_SORT](#) = 1<<10, [KDB_O_NORECURSIVE](#) = 1<<11, [KDB_O_NOCASE](#)
 = 1<<12,
 [KDB_O_WITHOWNER](#) = 1<<13, [KDB_O_NOALL](#) = 1<<14 }

Functions

- KDB * [kdbOpen](#) (Key *errorKey)
 Opens the session with the Key database.
- int [kdbClose](#) (KDB *handle, Key *errorKey)
- int [kdbGet](#) (KDB *handle, KeySet *ks, Key *parentKey)
 Retrieve keys in an atomic and universal way, all other [kdbGet\(\)](#) Functions rely on that one.
- int [kdbSet](#) (KDB *handle, KeySet *ks, Key *parentKey)

6.1.1 Detailed Description

General methods to access the Key database. To use them:

```
#include <kdb.h>
```

The `kdb*`() class of methods are used to access the storage, to get and set [Keys](#) or [KeySets](#) .

The most important functions are:

- [kdbOpen\(\)](#)
- [kdbClose\(\)](#)
- [kdbGet\(\)](#)
- [kdbSet\(\)](#)

The two essential functions for dynamic information about backends are:

- `kdbGetMountpoint()`

They use some backend implementation to know the details about how to access the storage. Currently we have this backends:

- `berkeleydb`: the keys are stored in a Berkeley DB database, providing very small footprint, speed, and other advantages.
- `filesystem`: the key hierarchy and data are saved as plain text files in the filesystem.
- `ini`: the key hierarchy are saved into configuration files.

See Also

<http://www.libelektra.org/Ini>

- `fstab`: a reference backend used to interpret the `/etc/fstab` file as a set of keys under `system/filesystems`.
- `gconf`: makes Elektra use the GConf daemon to access keys. Only the `user/` tree is available since GConf is not system wide.

Backends are physically a library named `/lib/libelektra-{NAME}.so`.

See [writing a new plugin](#) for information about how to write a plugin.

Language binding writers should follow the same rules:

- You must relay completely on the backend-dependent methods.
- You may use or reimplement the second set of methods.
- You should completely reimplement in your language the higher lever methods.
- Many methods are just for comfort in C. These methods are marked and need not to be implemented if the binding language has e.g. string operators which can do the operation easily.

6.1.2 Enumeration Type Documentation

6.1.2.1 `enum option_t`

Options to change the default behavior of `kdbGet()`, `kdbSet()` and `ksLookup()` functions.

These options can be ORed. That is the `|`-Operator in C.

See Also

[kdbGet\(\)](#), [kdbSet\(\)](#)

Enumerator:

`KDB_O_NONE` No Option set. Will be recursive with no inactive keys.

See Also

[kdbGet\(\)](#), [kdbSet\(\)](#), [ksLookup\(\)](#)

`KDB_O_DEL` Delete parentKey key in [kdbGet\(\)](#), [kdbSet\(\)](#) or [ksLookup\(\)](#).

See Also

[kdbGet\(\)](#), [kdbSet\(\)](#)

`KDB_O_POP` Pop Parent out of keyset key in [kdbGet\(\)](#).

See Also

[ksPop\(\)](#).

KDB_O_NODIR Exclude keys containing other keys in result.

Only return leaves.

See Also

[keyIsDir\(\)](#)

KDB_O_DIRONLY Retrieve only directory keys (keys containing other keys). This will give you an skeleton without leaves. This must not be used together with KDB_O_NODIR.

See Also

[keyIsDir\(\)](#)

KDB_O_NOREMOVE Don't remove any keys. This must not be used together with KDB_O_REMOVEONLY.

KDB_O_REMOVEONLY Only remove keys. This must not be used together with KDB_O_NOREMOVE.

KDB_O_INACTIVE Do not ignore inactive keys (that name begins with .).

See Also

[keyIsInactive\(\)](#)

KDB_O_SYNC Set keys independent of sync status.

See Also

[keyNeedSync\(\)](#)

KDB_O_SORT This option has no effect. KeySets are always sorted.

Deprecated dont use

KDB_O_NORECURSIVE Do not call [kdbGet\(\)](#) for every key containing other keys ([keyIsDir\(\)](#)).

KDB_O_NOCASE Ignore case.

KDB_O_WITHOWNER Search with owner.

KDB_O_NOALL Only search from start -> cursor to cursor -> end.

6.1.3 Function Documentation

6.1.3.1 int kdbClose (KDB * *handle*, Key * *errorKey*)

Closes the session with the Key database.

You should call this method when you finished your affairs with the key database. You can manipulate Key and KeySet objects also after [kdbClose\(\)](#). You must not use any kdb* call afterwards. You can implement [kdbClose\(\)](#) in the atexit() handler.

This is the counterpart of [kdbOpen\(\)](#).

The *handle* parameter will be finalized and all resources associated to it will be freed. After a [kdbClose\(\)](#), this *handle* can't be used anymore, unless it gets initialized again with another call to [kdbOpen\(\)](#).

See Also

[kdbOpen\(\)](#)

Parameters

<i>handle</i>	contains internal information of opened key database
<i>errorKey</i>	the key which holds error information

Returns

- 0 on success
- 1 on NULL pointer

6.1.3.2 int kdbGet (KDB * *handle*, KeySet * *ks*, Key * *parentKey*)

Retrieve keys in an atomic and universal way, all other [kdbGet\(\)](#) Functions rely on that one.

The returned KeySet must be initialized. The returned KeySet may already contain some keys. The new retrieved keys will be appended using [ksAppendKey\(\)](#).

It will fully retrieve all keys under the `parentKey` folder, with all subfolders and their children.

6.1.4 Example

This example demonstrates the typical usecase within an application without updating.

Example:

```

KeySet *myConfig = ksNew(0);
Key *key = keyNew("system/sw/MyApp", KEY_END);
KDB *handle = kdbOpen(key);

kdbGet(handle, myConfig, key);

keySetName(key, "user/sw/MyApp");
kdbGet(handle, myConfig, key);

// check for errors in key
keyDel(key);

key = ksLookupByName(myConfig, "/sw/MyApp/key", 0);
// check if key is not 0 and work with it...

ksDel(myConfig); // delete the in-memory configuration

// maybe you want kdbSet() myConfig here

kdbClose(handle, 0); // no more affairs with the key database.

```

6.1.5 Details

When no backend could be found (e.g. no backend mounted) the default backend will be used.

If you pass NULL on any parameter [kdbGet\(\)](#) will fail immediately without doing anything.

When a backend fails [kdbGet\(\)](#) will return -1 without any changes to one of the parameter.

6.1.6 Updating

In the first run of `kdbGet` all keys are retrieved. On subsequent calls only the keys are retrieved where something was changed inside the key database. The other keys stay unchanged in the keyset, even when they were manipulated.

It is your responsibility to save the original keyset if you need it afterwards.

If you must get the same keyset again, e.g. in another thread you need to open a second handle to the key database using [kdbOpen\(\)](#).

Parameters

<i>handle</i>	contains internal information of opened key database
<i>parentKey</i>	parent key holds the information which keys should be get - invalid name gets all keys
<i>ks</i>	the (pre-initialized) KeySet returned with all keys found will not be changed on error or if no update is required

See Also

[ksLookupByName\(\)](#) for powerful lookups after the KeySet was retrieved

Returns

- 1 if the keys were retrieved successfully
- 0 if there was no update - no changes are made to the keyset then
- 1 on failure - no changes are made to the keyset then

6.1.6.1 KDB* kdbOpen (Key * errorKey)

Opens the session with the Key database.

The first step is to open the default backend. With it system/elektra/mountpoints will be loaded and all needed libraries and mountpoints will be determined. These libraries for backends will be loaded and with it the KDB datastructure will be initialized.

You must always call this method before retrieving or committing any keys to the database. In the end of the program, after using the key database, you must not forget to [kdbClose\(\)](#). You can use the `atexit ()` handler for it.

The pointer to the KDB structure returned will be initialized like described above, and it must be passed along on any `kdb*()` method your application calls.

Get a KDB handle for every thread using `elektra`. Don't share the handle across threads, and also not the pointer accessing it:

```
thread1
{
    KDB * h;
    h = kdbOpen(0);
    // fetch keys and work with them
    kdbClose(h, 0);
}
thread2
{
    KDB * h;
    h = kdbOpen(0);
    // fetch keys and work with them
    kdbClose(h, 0);
}
```

You don't need to use the [kdbOpen\(\)](#) if you only want to manipulate plain in-memory Key or KeySet objects without any affairs with the backend key database or when your application loads plugins directly.

Parameters

<i>errorKey</i>	the key which holds errors and warnings which were issued must be given
-----------------	---

See Also

[kdbClose\(\)](#) to end all affairs to the [Key](#) database.

Returns

- a KDB pointer on success
- NULL on failure

6.1.6.2 int kdbSet (KDB * handle, KeySet * ks, Key * parentKey)

Set keys in an atomic and universal way.

All other `kdbSet` Functions rely on that one.

6.1.7 parentKey

With `parentKey` you can only store a part of the given keyset.

```
KeySet *ks = ksNew(0);
Key *parentKey = keyNew("user/app/myapp/default", KEY_END);
kdbGet (h, ks, parentKey);

//now only set everything below user
if (kdbSet (h, ks, parentKey) == -1)
{
    // in parentKey you can check the error cause
    // ksCurrent(ks) is the faulty key
}

ksDel (ks);
```

If you pass a `parentKey` without a name the whole keyset will be set in an atomic way.

6.1.8 Update

Each key is checked with `keyNeedSync()` before being actually committed. So only changed keys are updated. If no key of a backend needs to be synced any affairs to backends omitted and 0 is returned.

6.1.9 Error Situations

If some error occurs, `kdbSet()` will stop. In this situation the `KeySet` internal cursor will be set on the key that generated the error.

None of the keys are actually committed.

You should present the error message to the user and let the user decide what to do. Possible solutions are:

- repeat the same `kdbSet` (for temporary errors)
- remove the key and set it again (for validation or type errors)
- change the value and try it again (for validation errors)
- do a `kdbGet` and then (for conflicts ...)
- set the same keyset again (in favour of what was set by this user)
- drop the old keyset (in favour of what was set elsewhere)
- export the configuration into a file (for unresolvable errors)

Example of how this method can be used:

```
int i;
KeySet *ks; // the KeySet I want to set
// fill ks with some keys
for (i=0; i< NR_OF_TRIES; i++) // limit to NR_OF_TRIES tries
{
    ret=kdbSet(handle, ks, parentKey);
    if (ret == -1)
    {
        // We got an error. Warn user.
        Key *problemKey = ksCurrent(ks);
        // parentKey has the errorInformation
        // problemKey is the faulty key (may be null)
        int userInput = showElektraErrorDialog (parentKey, problemKey);
        switch (userInput)
        {
            case INPUT_REPEAT: continue;
            case INPUT_REMOVE: ksLookup (ks, parentKey, KDB_O_POP
        ); break;
        }
        ...
    }
}
```

Parameters

<i>handle</i>	contains internal information of opened key database
<i>ks</i>	a KeySet which should contain changed keys, otherwise nothing is done
<i>parentKey</i>	holds the information below which key keys should be set, see above

Returns

1 on success
0 if nothing had to be done
-1 on failure


See Also

[keyNeedSync\(\)](#), [ksNext\(\)](#), [ksCurrent\(\)](#)

6.2 Key

A Key is the essential class that encapsulates key [name](#) , [value](#) and [metainfo](#) .

Collaboration diagram for Key:



Modules

- [Basic Methods](#)
Key construction and initialization methods.
- [Meta Info Manipulation Methods](#)
Methods to do various operations on Key metainfo.
- [Methods for Making Tests](#)
Methods to do various tests on Keys.
- [Name Manipulation Methods](#)
Methods to do various operations on Key names.
- [Value Manipulation Methods](#)
Methods to do various operations on Key values.

Enumerations

- `enum keyswitch_t {`
`KEY_NAME = 1, KEY_VALUE = 1<<1, KEY_OWNER = 1<<2, KEY_COMMENT = 1<<3,`
`KEY_BINARY = 1<<4, KEY_UID = 1<<5, KEY_GID = 1<<6, KEY_MODE = 1<<7,`
`KEY_ETIME = 1<<8, KEY_MTIME = 1<<9, KEY_CTIME = 1<<10, KEY_SIZE = 1<<11,`
`KEY_DIR = 1<<14, KEY_END = 0 }`

Functions

- Key * [keyNew](#) (const char *[keyName](#),...)
- Key * [keyDup](#) (const Key *source)
- int [keyCopy](#) (Key *dest, const Key *source)
- int [keyDel](#) (Key *key)
- int [keyClear](#) (Key *key)
- ssize_t [keyIncRef](#) (Key *key)
- ssize_t [keyDecRef](#) (Key *key)
- ssize_t [keyGetRef](#) (const Key *key)

6.2.1 Detailed Description

A Key is the essential class that encapsulates key [name](#) , [value](#) and [metainfo](#) .

6.2.2 Enumeration Type Documentation

6.2.2.1 enum keyswitch_t

Switches to denote the various Key attributes in methods throughout this library.

This enum switch provide a flag for every metadata in a key.

In case of [keyNew\(\)](#) they give Information what Parameter comes next.

See Also

[keyNew\(\)](#)
[ksToStream\(\)](#), [keyToStream\(\)](#)

Enumerator:

KEY_NAME Flag for the key name
KEY_VALUE Flag for the key data
KEY_OWNER Flag for the key user domain
KEY_COMMENT Flag for the key comment
KEY_BINARY Flag if the key is binary
KEY_UID Flag for the key UID
KEY_GID Flag for the key GID
KEY_MODE Flag for the key permissions
KEY_ETIME Flag for the key access time
KEY_MTIME Flag for the key change time
KEY_CTIME Flag for the key status change time
KEY_SIZE Flag for maximum size to limit value
KEY_DIR Flag for the key directories
KEY_END Used as a parameter terminator to [keyNew\(\)](#)

6.2.3 Function Documentation

6.2.3.1 int keyClear (Key * key)

Key Object Cleaner.

Will reset all internal data.

After this call you will receive a fresh key.

The reference counter will stay unmodified.

Note

that you might also `clear()` all aliases with this operation.

```
int f (Key *k)
{
    keyClear (k);
    // you have a fresh key k here
    keySetString (k, "value");
    // the caller will get an empty key k with an value
}
```

Returns

returns 0 on success

-1 on null pointer

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

6.2.3.2 int keyCopy (Key * *dest*, const Key * *source*)

Copy or Clear a key.

Most often you may prefer `keyDup()` which allocates a new key and returns a duplication of another key.

But when you need to copy into an existing key, e.g. because it was passed by a pointer in a function you can do so:

```
void h (Key *k)
{
    // receive key c
    keyCopy (k, c);
    // the caller will see the changed key k
}
```

The reference counter will not be changed for both keys. Affiliation to keysets are also not affected.

When you pass a NULL-pointer as source the data of dest will be cleaned completely (except reference counter, see `keyClear()`) and you get a fresh dest key.

```
void g (Key *k)
{
    keyCopy (k, 0);
    // k is now an empty and fresh key
}
```

The meta data will be duplicated for the destination key. So it will not take much additional space, even with lots of metadata.

If you want to copy all metadata, but keep the old value you can use `keyCopy()` too.

```
void j (Key *k)
{
    size_t size = keyGetValueSize (k);
    char *value = malloc (size);
    int bstring = keyIsString (k);

    // receive key c
    memcpy (value, keyValue(k), size);
}
```



```

keyCopy (k, c);
if (bstring) keySetString (k, value);
else keySetBinary (k, value, size);
free (value);
// the caller will see the changed key k
// with the metadata from c
}

```

Note

Next to the value itself we also need to remember if the value was string or binary. So in fact the meta data of the resulting key *k* in that example is not a complete duplicate, because the meta data "binary" may differ. Similar considerations might be necessary for the type of the key and so on, depending on the concrete situation.

Parameters

<i>dest</i>	the key which will be written to
<i>source</i>	the key which should be copied or NULL to clean the destination key

Returns

-1 on failure when a NULL pointer was passed for *dest* or a dynamic property could not be written. Both name and value are empty then.
 0 when *dest* was cleaned
 1 when source was successfully copied

See Also

[keyDup\(\)](#) to get a duplication of a [Key](#)

6.2.3.3 ssize_t keyDecRef (Key * key)

Decrement the viability of a key object.

The references will be decremented for [ksPop\(\)](#) or successful calls of [ksLookup\(\)](#) with the option KDB_O_POP. It will also be decremented with an following [keyDel\(\)](#) in the case that an old key is replaced with another key with the same name.

The reference counter can't be decremented once it reached 0. In that situation nothing will happen and 0 will be returned.

Note

[keyDup\(\)](#) will reset the references for dopped key.

Returns

the value of the new reference counter
 -1 on null pointer
 0 when the key is ready to be freed

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

See Also

[keyGetRef\(\)](#), [keyDel\(\)](#), [keyIncRef\(\)](#)

6.2.3.4 `int keyDel (Key * key)`

A destructor for Key objects.

Every key created by [keyNew\(\)](#) must be deleted with [keyDel\(\)](#).

It is save to delete keys which are in a keyset, the number of references will be returned then.

It is save to delete a nullpointer, -1 will be returned then.

It is also save to delete a multiple referenced key, nothing will happen then and the reference counter will be returned.

Parameters

<code>key</code>	the key object to delete
------------------	--------------------------

See Also

[keyNew\(\)](#), [keyInc\(\)](#), [keyGetRef\(\)](#)

Returns

the value of the reference counter if the key is within keyset(s)
 0 when the key was freed
 -1 on null pointers

6.2.3.5 `Key* keyDup (const Key * source)`

Return a duplicate of a key.

Memory will be allocated as needed for dynamic properties.

The new key will not be member of any KeySet and will start with a new reference counter at 0. A subsequent [keyDel\(\)](#) will delete the key.

```
int f (const Key * source)
{
    Key * dup = keyDup (source);
    // work with duplicate
    keyDel (dup);
    // everything related to dup is freed
    // and source is unchanged
}
```

Like for a new key after [keyNew\(\)](#) a subsequent [ksAppend\(\)](#) makes a KeySet to take care of the lifecycle of the key.

```
int g (const Key * source, KeySet * ks)
{
    Key * dup = keyDup (source);
    // work with duplicate
    ksAppendKey (ks, dup);
    // ksDel(ks) will also free the duplicate
    // source remains unchanged.
}
```

Duplication of keys should be preferred to [keyNew\(\)](#), because data like owner can be filled with a copy of the key instead of asking the environment. It can also be optimized in the checks, because the keyname is known to be valid.

Parameters

<i>source</i>	has to be an initialised source Key
---------------	-------------------------------------

Returns

0 failure or on NULL pointer
a fully copy of source on success

See Also

[ksAppend\(\)](#), [keyDel\(\)](#), [keyNew\(\)](#)

6.2.3.6 ssize_t keyGetRef (const Key * key)

Return how many references the key has.

The references will be incremented on successful calls to [ksAppendKey\(\)](#) or [ksAppend\(\)](#).

Note

[keyDup\(\)](#) will reset the references for dupped key.

For your own applications you can use [keyIncRef\(\)](#) and [keyDecRef\(\)](#) for reference counting. Keys with zero references will be deleted when using [keyDel\(\)](#).

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

the number of references
-1 on null pointer

See Also

[keyIncRef\(\)](#) and [keyDecRef\(\)](#)

6.2.3.7 ssize_t keyIncRef (Key * key)

Increment the viability of a key object.

This function is intended for applications using their own reference counter for key objects. With it you can increment the reference and thus avoid destruction of the object in a subsequent [keyDel\(\)](#).

```
Key *k;
keyInc (k);
function_that_keyDec(k);
// work with k
keyDel (k); // now really free it
```

The reference counter can't be incremented once it reached SSIZE_MAX. In that situation nothing will happen and SSIZE_MAX will be returned.

Note

[keyDup\(\)](#) will reset the references for dupped key.

Returns

the value of the new reference counter
 -1 on null pointer
 SSIZE_MAX when maximum exceeded

Parameters

<code>key</code>	the key object to work with
------------------	-----------------------------

See Also

[keyGetRef\(\)](#), [keyDecRef\(\)](#), [keyDel\(\)](#)

6.2.3.8 Key* keyNew (const char * *keyName*, ...)

A practical way to fully create a Key object in one step.

This function tries to mimic the C++ way for constructors.

To just get a key object, simple do:

```
Key *k = keyNew(0);
// work with it
keyDel (k);
```

If you want the key object to contain a name, value, comment and other meta info read on.

Note

When you already have a key with similar properties its easier and cheaper to [keyDup\(\)](#) the key.

Due to ABI compatibility, the `Key` structure is not defined in `kdb.h`, only declared. So you can only declare pointers to Keys in your program, and allocate and free memory for them with [keyNew\(\)](#) and [keyDel\(\)](#) respectively. See <http://tldp.org/HOWTO/Program-Library-HOWTO/shared-libraries.html#A-EN135>

You can call it in many different ways depending on the attribute tags you pass as parameters. Tags are represented as the `keyswitch_t` values, and tell [keyNew\(\)](#) which Key attribute comes next.

The simplest and minimum way to use it is with no tags, only a key name:

```
Key *nullKey, *emptyNamedKey;

// Create a key that has no name, is completely empty, but is initialized
nullKey=keyNew(0);
keyDel (nullKey);

// Is the same as above
nullKey=keyNew("", KEY_END);
keyDel (nullKey);

// Create and initialize a key with a name and nothing else
emptyNamedKey=keyNew("user/some/example", KEY_END);
keyDel (emptyNamedKey);
```

[keyNew\(\)](#) allocates memory for a key object and cleans everything up. After that, it processes the given argument list.

The Key attribute tags are the following:

- `keyswitch_t::KEY_TYPE`

Next parameter is a type of the value. Default assumed is `KEY_TYPE_UNDEFINED`. Set this attribute so that a subsequent `KEY_VALUE` can toggle to [keySetString\(\)](#) or [keySetBinary\(\)](#) regarding to [keyIsString\(\)](#) or [keyIsBinary\(\)](#). If you don't use `KEY_TYPE` but a `KEY_VALUE` follows afterwards, `KEY_TYPE_STRING` will be used.

- `keyswitch_t::KEY_SIZE`

Define a maximum length of the value. This is especially useful for setting a binary key. So make sure you use that before you `KEY_VALUE` for binary keys.

- `keyswitch_t::KEY_VALUE`

Next parameter is a pointer to the value that will be set to the key. If no `keyswitch_t::KEY_TYPE` was used before, `keyswitch_t::KEY_TYPE_STRING` is assumed. If `KEY_TYPE` was previously passed with a `KEY_TYPE_BINARY`, you should have passed `KEY_SIZE` before! Otherwise it will be cut off with first `\0` in string!

- `keyswitch_t::KEY_UID`, `keyswitch_t::KEY_GID`

Next parameter is taken as the UID (`uid_t`) or GID (`gid_t`) that will be defined on the key. See [keySetUID\(\)](#) and [keySetGID\(\)](#).

- `keyswitch_t::KEY_MODE`

Next parameter is taken as mode permissions (`mode_t`) to the key. See [keySetMode\(\)](#).

- `keyswitch_t::KEY_DIR`

Define that the key is a directory rather than a ordinary key. This means its executable bits in its mode are set. This option allows the key to have subkeys. See [keySetDir\(\)](#).

- `keyswitch_t::KEY_OWNER`

Next parameter is the owner. See [keySetOwner\(\)](#).

- `keyswitch_t::KEY_COMMENT`

Next parameter is a comment. See [keySetComment\(\)](#).

- `keyswitch_t::KEY_END`

Must be the last parameter passed to [keyNew\(\)](#). It is always required, unless the `keyName` is 0.

Example:

```
KeySet *ks=keyNew(0);

ksAppendKey(ks,keyNew(0));           // an empty key

ksAppendKey(ks,keyNew("user/sw",      // the name of
    KEY_END));                       // no more args

ksAppendKey(ks,keyNew("user/tmp/ex1",
    KEY_VALUE,"some data",           // set a string value
    KEY_END));                       // end of args

ksAppendKey(ks,keyNew("user/tmp/ex2",
    KEY_VALUE,"some data",           // with a simple value
    KEY_MODE,0777,                   // permissions
    KEY_END));                       // end of args

ksAppendKey(ks,keyNew("user/tmp/ex4",
    KEY_TYPE,KEY_TYPE_BINARY,        // key type
    KEY_SIZE,7,                      // assume binary length 7
    KEY_VALUE,"some data",           // value that will be
    truncated in 7 bytes              //
    KEY_COMMENT,"value is truncated",
    KEY_OWNER,"root",                // owner (not uid) is root
    KEY_UID,0,                      // root uid
    KEY_END));                       // end of args

ksAppendKey(ks,keyNew("user/tmp/ex5",
    KEY_TYPE,KEY_TYPE_DIR | KEY_TYPE_BINARY, // dir key with a binary value
    KEY_SIZE,7,                      //
    KEY_VALUE,"some data",           // value that will be
    truncated in 7 bytes              //
    KEY_COMMENT,"value is truncated",
    KEY_OWNER,"root",                // owner (not uid) is root
    KEY_UID,0,                      // root uid
    KEY_END));                       // end of args

ksDel(ks);
```

The reference counter (see [keyGetRef\(\)](#)) will be initialized with 0, that means a subsequent call of [keyDel\(\)](#) will delete the key. If you append the key to a keyset the reference counter will be incremented by one (see [keyInc\(\)](#)) and the key can't be deleted by a [keyDel\(\)](#).

```
Key *k = keyNew(0); // ref counter 0
ksAppendKey(ks, k); // ref counter of key 1
ksDel(ks); // key will be deleted with keyset
*
```

If you increment only by one with [keyInc\(\)](#) the same as said above is valid:

```
Key *k = keyNew(0); // ref counter 0
keyIncRef(k); // ref counter of key 1
keyDel(k); // has no effect
keyDecRef(k); // ref counter back to 0
keyDel(k); // key is now deleted
*
```

If you add the key to more keySets:

```
Key *k = keyNew(0); // ref counter 0
ksAppendKey(ks1, k); // ref counter of key 1
ksAppendKey(ks2, k); // ref counter of key 2
ksDel(ks1); // ref counter of key 1
ksDel(ks2); // k is now deleted
*
```

or use [keyInc\(\)](#) more than once:

```
Key *k = keyNew(0); // ref counter 0
keyIncRef(k); // ref counter of key 1
keyDel(k); // has no effect
keyIncRef(k); // ref counter of key 2
keyDel(k); // has no effect
keyDecRef(k); // ref counter of key 1
keyDel(k); // has no effect
keyDecRef(k); // ref counter is now 0
keyDel(k); // k is now deleted
*
```

they key won't be deleted by a [keyDel\(\)](#) as long refcounter is not 0.

The key's sync bit will always be set for any call, except:

```
Key *k = keyNew(0);
// keyNeedSync() will be false
```

Parameters

<i>keyName</i>	a valid name to the key, or NULL to get a simple initialized, but really empty, object
----------------	--

See Also

[keyDel\(\)](#)

Returns

a pointer to a new allocated and initialized Key object, or NULL if an invalid `keyName` was passed (see [keySetName\(\)](#)).

6.3 Basic Methods

Key construction and initialization methods.

Collaboration diagram for Basic Methods:



Key construction and initialization methods. To use them:

```
#include <kdb.h>
```

Key properties are:


- [Key name](#)
- [Key value](#)
- [Key comment](#)
- [Key owner](#)
- [UID, GID and filesystem-like mode permissions](#)
- [Mode, change and modification times](#)

Described here the methods to allocate and free the key.

6.4 Meta Info Manipulation Methods

Methods to do various operations on Key metainfo.

Collaboration diagram for Meta Info Manipulation Methods:



Functions

- int [keyRewindMeta](#) (Key *key)
- const Key * [keyNextMeta](#) (Key *key)
- const Key * [keyCurrentMeta](#) (const Key *key)
- int [keyCopyMeta](#) (Key *dest, const Key *source, const char *metaName)
- int [keyCopyAllMeta](#) (Key *dest, const Key *source)
- const Key * [keyGetMeta](#) (const Key *key, const char *metaName)
- ssize_t [keySetMeta](#) (Key *key, const char *metaName, const char *newMetaString)
- uid_t [keyGetUID](#) (const Key *key)
- int [keySetUID](#) (Key *key, uid_t uid)
- gid_t [keyGetGID](#) (const Key *key)
- int [keySetGID](#) (Key *key, gid_t gid)
- int [keySetDir](#) (Key *key)
- mode_t [keyGetMode](#) (const Key *key)
- int [keySetMode](#) (Key *key, mode_t mode)
- time_t [keyGetATime](#) (const Key *key)
- int [keySetATime](#) (Key *key, time_t atime)
- time_t [keyGetMTime](#) (const Key *key)
- int [keySetMTime](#) (Key *key, time_t mtime)
- time_t [keyGetCTime](#) (const Key *key)
- int [keySetCTime](#) (Key *key, time_t ctime)

6.4.1 Detailed Description

Methods to do various operations on Key metainfo. To use them:

```
#include <kdb.h>
```

Next to [Name \(key and owner\)](#) and [Value \(data and comment\)](#) there is the so called meta information inside every key.

Key meta information are an unlimited number of key/value pairs strongly related to a key. It main purpose is to give keys special semantics, so that plugins can treat them differently.

File system information (see stat(2) for more information):

- uid: the user id (positive number)

- gid: the group id (positive number)
- mode: filesystem-like mode permissions (positive octal number)
- atime: When was the key accessed the last time.
- mtime: When was the key modified the last time.
- ctime: When the uid, gid or mode of a key changes. (times are represented through a positive number as unix timestamp)

The comment can contain userdata which directly belong to that key. The name of the meta information is "comment" for a general purpose comment about the key. Multi-Language comments are also supported by appending [LANG] to the name.

Validators are regular expressions which are tested against the key value. The metakey "validator" can hold a regular expression which will be matched against.

Types can be expressed with the meta information "type".

The relevance of the key can be tagged with a value from -20 to 20. Negative numbers are the more important and must be present in order to start the program.

A version of a key may be stored with "version". Its format is full.major.minor where all of these are integers.

The order inside a persistent storage can be described with the tag "order" which contains a positive number.

The meta key "app" describes to which application a key belongs. It can be used to remove keys from an application no longer installed.

The meta key "path" describes where the key is physically stored.

The "owner" is the user that owns the key. It only works for the user/ hierarchy. It rather says where the key is stored and says nothing about the filesystem properties.

6.4.2 Function Documentation

6.4.2.1 `int keyCopyAllMeta (Key * dest, const Key * source)`

Do a shallow copy of all meta data from source to dest.

The key dest will additionally have all meta data source had. Meta data not present in source will not be changed. Meta data which was present in source and dest will be overwritten.

For example the meta data type is copied into the Key k.

```
void l(Key *k)
{
    // receive c
    keyCopyMeta(k, c);
    // the caller will see the changed key k
    // with all the metadata from c
}
```

The main purpose of this function is for plugins or applications which want to add the same meta data to n keys. When you do that with `keySetMeta()` it will take n times the memory for the key. This can be considerable amount of memory for many keys with some meta data for each.

To avoid that problem you can use `keyCopyAllMeta()` or `keyCopyMeta()`.

```
void o(KeySet *ks)
{
    Key *current;
    Key *shared = keyNew (0);
    keySetMeta(shared, "shared1", "this meta data should be
shared among many keys");
    keySetMeta(shared, "shared2", "this meta data should be
shared among many keys also");
    keySetMeta(shared, "shared3", "this meta data should be
```

```

        shared among many keys too");

    ksRewind(ks);
    while ((current = ksNext(ks)) != 0)
    {
        if (needs_shared_data(current)) keyCopyAllMeta(
current, shared);
    }
}

```

Postcondition

for every metaName present in source: `keyGetMeta(source, metaName) == keyGetMeta(dest, metaName)`

Returns

1 if was successfully copied
0 if source did not have any meta data
-1 on null pointers (source or dest)
-1 on memory problems

Parameters

<i>dest</i>	the destination where the meta data should be copied too
<i>source</i>	the key where the meta data should be copied from

6.4.2.2 int keyCopyMeta (Key * dest, const Key * source, const char * metaName)

Do a shallow copy of meta data from source to dest.

The key dest will have the same meta data referred with metaName afterwards then source.

For example the meta data type is copied into the Key k.

```

void l(Key *k)
{
    // receive c
    keyCopyMeta(k, c, "type");
    // the caller will see the changed key k
    // with the metadata "type" from c
}

```

The main purpose of this function is for plugins or applications which want to add the same meta data to n keys. When you do that with `keySetMeta()` it will take n times the memory for the key. This can be considerable amount of memory for many keys with some meta data for each.

To avoid that problem you can use `keyCopyAllMeta()` or `keyCopyMeta()`.

```

void o(KeySet *ks)
{
    Key *current;
    Key *shared = keyNew (0);
    keySetMeta(shared, "shared", "this meta data should be shared
among many keys");

    ksRewind(ks);
    while ((current = ksNext(ks)) != 0)
    {
        if (needs_shared_data(current)) keyCopyMeta(current,
shared, "shared");
    }
}

```

Postcondition

`keyGetMeta(source, metaName) == keyGetMeta(dest, metaName)`

Returns

- 1 if was successfully copied
- 0 if the meta data in dest was removed too
- 1 on null pointers (source or dest)
- 1 on memory problems

Parameters

<i>dest</i>	the destination where the meta data should be copied too
<i>source</i>	the key where the meta data should be copied from
<i>metaName</i>	the name of the meta data which should be copied

6.4.2.3 `const Key* keyCurrentMeta (const Key * key)`

Returns the Value of a Meta-Information which is current.

The pointer is NULL if you reached the end or after [ksRewind\(\)](#).

Note

You must not delete or change the returned key, use [keySetMeta\(\)](#) if you want to delete or change it.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

- a buffer to the value pointed by *key*'s cursor
- 0 on NULL pointer

See Also

[keyNextMeta\(\)](#), [keyRewindMeta\(\)](#)
[ksCurrent\(\)](#) for pedant in iterator interface of KeySet

6.4.2.4 `time_t keyGetATime (const Key * key)`

Get last time the key data was read from disk.

Deprecated This API is obsolete.

Every [kdbGet\(\)](#) might update the access time of a key. You get information when the key was read the last time from the database.

You will get 0 when the key was not read already.

Beware that multiple copies of keys with [keyDup\(\)](#) might have different atimes because you [kdbGet\(\)](#) one, but not the other. You can use this information to decide which key is the latest.

Parameters

<i>key</i>	Key to get information from.
------------	------------------------------

Returns

the time you got the key with `kdbGet()`
 0 on key that was never `kdbGet()`
 (time_t)-1 on NULL pointer

See Also

`keySetATime()`
`kdbGet()`

6.4.2.5 time_t keyGetCTime (const Key * key)

Get last time the key metadata was changed from disk.

Deprecated This API is obsolete.

You will get 0 when the key was not read already.

Any changed field in metadata will influence the ctime of a key.

This time is not updated if only value or comment are changed.

Not changed keys will not update this time, even after `kdbSet()`.

It is possible that other keys written to disc influence this time if the backend is not grained enough.

Parameters

<i>key</i>	Key to get information from.
------------	------------------------------

See Also

`keySetCTime()`

Returns

(time_t)-1 on NULL pointer
 the metadata change time

6.4.2.6 gid_t keyGetGID (const Key * key)

Get the group ID of a key.

Deprecated This API is obsolete.

6.4.3 GID

The group ID is a unique identification for every group present on a system. Keys will belong to root (0) as long as you did not get their real GID with `kdbGet()`.

Unlike UID users might change their group. This makes it possible to share configuration between some users.

A fresh key will have (gid_t)-1 also known as the group nogroup. It means that the key is not related to a group ID at the moment.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

the system's GID of the key
(gid_t)-1 on NULL key or currently unknown ID

See Also

[keySetGID\(\)](#), [keyGetUID\(\)](#)

6.4.3.1 const Key* keyGetMeta (const Key * key, const char * metaName)

Returns the Value of a Meta-Information given by name.

This is a much more efficient version of [keyGetMeta\(\)](#). But unlike with [keyGetMeta](#) you are not allowed to modify the resulting string.

```
int f(Key *k)
{
    if (!strcmp(keyValue(keyGetMeta(k, "type")), "boolean"))
    {
        // the type of the key is boolean
    }
}
```

Note

You must not delete or change the returned key, use [keySetMeta\(\)](#) if you want to delete or change it.

Parameters

<i>key</i>	the key object to work with
<i>metaName</i>	the name of the meta information you want the value from

Returns

0 if the key or metaName is 0
0 if no such metaName is found
value of Meta-Information if Meta-Information is found

See Also

[keyGetMeta\(\)](#), [keySetMeta\(\)](#)

6.4.3.2 mode_t keyGetMode (const Key * key)

Return the key mode permissions.

Deprecated This API is obsolete.

Default is 0664 (octal) for keys and 0775 for directory keys which used [keySetDir\(\)](#).

The defaults are defined with the macros KDB_FILE_MODE and KDB_DIR_MODE.

For more information about the mode permissions see [Modes](#).

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

mode permissions of the key
 KDB_FILE_MODE as defaults
 (mode_t)-1 on NULL pointer

See Also

[keySetMode\(\)](#)

6.4.3.3 `time_t keyGetMTime (const Key * key)`

Get last modification time of the key on disk.

Deprecated This API is obsolete.

You will get 0 when the key was not read already.

Everytime you change value or comment and [kdbSet\(\)](#) the key the mtime will be updated. When you [kdbGet\(\)](#) the key, the atime is set appropriate.

Not changed keys may not even passed to [kdbSet_backend\(\)](#) so it will not update this time, even after [kdbSet\(\)](#).

It is possible that other keys written to disc influence this time if the backend is not grained enough.

If you add or remove a key the key thereunder in the hierarchy will update the mtime if written with [kdbSet\(\)](#) to disc.

Parameters

<i>key</i>	Key to get information from.
------------	------------------------------

See Also

[keySetMTime\(\)](#)

Returns

the last modification time
 (time_t)-1 on NULL pointer

6.4.3.4 `uid_t keyGetUID (const Key * key)`

Get the user ID of a key.

Deprecated This API is obsolete.

6.4.4 UID

The user ID is a unique identification for every user present on a system. Keys will belong to root (0) as long as you did not get their real UID with [kdbGet\(\)](#).

Although usually the same, the UID of a key is not related to its owner.

A fresh key will have no UID.

Parameters

<code>key</code>	the key object to work with
------------------	-----------------------------

Returns

the system's UID of the key
(uid_t)-1 on NULL key

See Also

[keyGetGID\(\)](#), [keySetUID\(\)](#), [keyGetOwner\(\)](#)

6.4.4.1 `const Key* keyNextMeta (Key * key)`

Iterate to the next meta information.

Keys have an internal cursor that can be reset with [keyRewindMeta\(\)](#). Every time [keyNextMeta\(\)](#) is called the cursor is incremented and the new current Name of Meta Information is returned.

You'll get a NULL pointer if the meta information after the end of the Key was reached. On subsequent calls of [keyNextMeta\(\)](#) it will still return the NULL pointer.

The `key` internal cursor will be changed, so it is not const.

Note

That the resulting key is guaranteed to have a value, because meta information has no binary or null pointer semantics.

You must not delete or change the returned key, use [keySetMeta\(\)](#) if you want to delete or change it.

Parameters

<code>key</code>	the key object to work with
------------------	-----------------------------

Returns

a key representing meta information
0 when the end is reached
0 on NULL pointer

See Also

[ksNext\(\)](#) for pedant in iterator interface of KeySet

6.4.4.2 `int keyRewindMeta (Key * key)`

Rewind the internal iterator to first meta data.

Use it to set the cursor to the beginning of the Key Meta Infos. [keyCurrentMeta\(\)](#) will then always return NULL afterwards. So you want to [keyNextMeta\(\)](#) first.

```
Key *key;
const Key *meta;

keyRewindMeta (key);
while ((meta = keyNextMeta (key)) != 0)
{
    printf ("name: %s, value: %s", keyName(meta), (const char*)
        keyValue(meta));
}
```

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

- 0 on success
- 0 if there is no meta information for that key ([keyNextMeta\(\)](#) will always return 0 in that case)
- 1 on NULL pointer

See Also

[keyNextMeta\(\)](#), [keyCurrentMeta\(\)](#)
[ksRewind\(\)](#) for pedant in iterator interface of KeySet

6.4.4.3 `int keySetATime (Key * key, time_t atime)`

Update the atime information for a key.

Deprecated This API is obsolete.

When you do manual sync of keys you might also update the atime to make them indistinguishable.

It can also be useful if you work with keys not using a keydatabase.

Parameters

<i>key</i>	The Key object to work with
<i>atime</i>	The new access time for the key

Returns

- 0 on success
- 1 on NULL pointer

See Also

[keyGetATime\(\)](#)

6.4.4.4 `int keySetCTime (Key * key, time_t ctime)`

Update the ctime information for a key.

Deprecated This API is obsolete.

Parameters

<i>key</i>	The Key object to work with
<i>ctime</i>	The new change metadata time for the key

Returns

- 0 on success
- 1 on NULL pointer

See Also

[keyGetCTime\(\)](#)

6.4.4.5 int keySetDir (Key * key)

Set mode so that key will be recognized as directory.

Deprecated This API is obsolete.

The function will add all executable bits.

- Mode 0200 will be translated to 0311
- Mode 0400 will be translated to 0711
- Mode 0664 will be translated to 0775

The macro KDB_DIR_MODE (defined to 0111) will be used for that.

The executable bits show that child keys are allowed and listable. There is no way to have child keys which are not listable for anyone, but it is possible to restrict listing the keys to the owner only.

- Mode 0000 means that it is a key not read or writable to anyone.
- Mode 0111 means that it is a directory not read or writable to anyone. But it is recognized as directory to anyone.

For more about mode see [keySetMode\(\)](#).

It is not possible to access keys below a not executable key. If a key is not writeable and executable [kdbSet\(\)](#) will fail to access the keys below. If a key is not readable and executable [kdbGet\(\)](#) will fail to access the keys below.

Parameters

<i>key</i>	the key to set permissions to be recognized as directory.
------------	---

Returns

- 0 on success
- 1 on NULL pointer

See Also

[keySetMode\(\)](#)

6.4.4.6 int keySetGID (Key * key, gid_t gid)

Set the group ID of a key.

Deprecated This API is obsolete.

See [GID](#) for more information about group IDs.

Parameters

<i>key</i>	the key object to work with
<i>gid</i>	is the group ID

Returns

0 on success
-1 on NULL key

See Also

[keyGetGID\(\)](#), [keySetUID\(\)](#)

6.4.4.7 ssize_t keySetMeta (Key * key, const char * metaName, const char * newMetaString)

Set a new Meta-Information.

Will set a new Meta-Information pair consisting of metaName and newMetaString.

Will add a new Pair for Meta-Information if metaName was not added up to now.

It will modify a existing Pair of Meta-Information if the the metaName was inserted already.

It will remove a meta information if newMetaString is 0.

Parameters

<i>key</i>	the key object to work with
<i>metaName</i>	the name of the meta information where you want to change the value
<i>newMetaString</i>	the new value for the meta information

Returns

-1 on error if key or metaName is 0, out of memory or names are not valid
0 if the Meta-Information for metaName was removed
size (>0) of newMetaString if Meta-Information was successfully added

See Also

[keyGetMeta\(\)](#)

6.4.4.8 int keySetMode (Key * key, mode_t mode)

Set the key mode permissions.

Deprecated This API is obsolete. It is only a mapping to `keySetMeta(key, "mode", str)` which should be preferred.

The mode consists of 9 individual bits for mode permissions. In the following explanation the octal notation with leading zero will be used.

Default is 0664 (octal) for keys and 0775 for directory keys which used [keySetDir\(\)](#).

The defaults are defined with the macros `KDB_FILE_MODE` and `KDB_DIR_MODE`.

Note

libelektra 0.7.0 only allows 0775 (directory keys) and 0664 (other keys). More will be added later in a sense of the description below.

6.4.5 Modes

0000 is the most restrictive mode. No user might read, write or execute the key.

Reading the key means to get the value by [kdbGet\(\)](#).

Writing the key means to set the value by [kdbSet\(\)](#).

Execute the key means to make a step deeper in the hierarchy. But you must be able to read the key to be able to list the keys below. See also [keySetDir\(\)](#) in that context. But you must be able to write the key to be able to add or remove keys below.

0777 is the most relaxing mode. Every user is allowed to read, write and execute the key, if he is allowed to execute and read all keys below.

0700 allows every action for the current user, identified by the uid. See [keyGetUID\(\)](#) and [keySetUID\(\)](#).

To be more specific for the user the single bits can elect the mode for read, write and execute. 0100 only allows executing which gives the information that it is a directory for that user, but not accessible. 0200 only allows reading. This information may be combined to 0300, which allows execute and reading of the directory. Last 0400 decides about the writing permissions.

The same as above is also valid for the 2 other octal digits. 0070 decides about the group permissions, in that case full access. Groups are identified by the gid. See [keyGetGID\(\)](#) and [keySetGID\(\)](#). In that example everyone with a different uid, but the gid of the the key, has full access.

0007 decides about the world permissions. This is taken into account when neither the uid nor the gid matches. So that example would allow everyone with a different uid and gid of that key gains full access.

Parameters

<i>key</i>	the key to set mode permissions
<i>mode</i>	the mode permissions

Returns

0 on success
-1 on NULL key

See Also

[keyGetMode\(\)](#)

6.4.5.1 int keySetMTime (Key * key, time_t mtime)

Update the mtime information for a key.

Deprecated This API is obsolete.

Parameters

<i>key</i>	The Key object to work with
<i>mtime</i>	The new modification time for the key

Returns

0 on success

See Also

[keyGetMTime\(\)](#)

6.4.5.2 `int keySetUID (Key * key, uid_t uid)`

Set the user ID of a key.

Deprecated This API is obsolete.

See [UID](#) for more information about user IDs.

Parameters

<i>key</i>	the key object to work with
<i>uid</i>	the user ID to set

Returns

0 on success

-1 on NULL key or conversion error


See Also

[keySetGID\(\)](#), [keyGetUID\(\)](#), [keyGetOwner\(\)](#)

6.5 Name Manipulation Methods

Methods to do various operations on Key names.

Collaboration diagram for Name Manipulation Methods:



Functions

- `const char * keyName (const Key *key)`
- `ssize_t keyGetNameSize (const Key *key)`
- `ssize_t keyGetName (const Key *key, char *returnedName, size_t maxSize)`
- `ssize_t keySetName (Key *key, const char *newName)`
- `ssize_t keyGetFullNameSize (const Key *key)`
- `ssize_t keyGetFullName (const Key *key, char *returnedName, size_t maxSize)`
- `const char * keyBaseName (const Key *key)`
- `ssize_t keyGetBaseNameSize (const Key *key)`
- `ssize_t keyGetBaseName (const Key *key, char *returned, size_t maxSize)`
- `ssize_t keyAddBaseName (Key *key, const char *baseName)`
- `ssize_t keySetBaseName (Key *key, const char *baseName)`
- `const char * keyOwner (const Key *key)`
- `ssize_t keyGetOwnerSize (const Key *key)`
- `ssize_t keyGetOwner (const Key *key, char *returnedOwner, size_t maxSize)`
- `ssize_t keySetOwner (Key *key, const char *newOwner)`

6.5.1 Detailed Description

Methods to do various operations on Key names. To use them:

```
#include <kdb.h>
```

These functions make it easier for c programmers to work with key names. Everything here can also be done with `keySetName`, described in `key`.

Rules for Key Names

When using Elektra to store your application's configuration and state, please keep in mind the following rules:

- You are not allowed to create keys right under `system` or `user`.
- You are not allowed to create folder keys right under `system` or `user`. They are reserved for very essential OS subsystems.
- The keys for your application, called say *MyApp*, should be created under `system/sw/MyApp` and/or `user/sw/MyApp`.

- It is suggested to make your application look for default keys under `system/sw/MyApp/current` and/or `user/sw/MyApp/current`. This way, from a sysadmin perspective, it will be possible to copy the `system/sw/MyApp/current` tree to something like `system/sw/MyApp/old`, and keep system clean and organized.
- `\0` must not occur in names.
- `/` is the separator.

6.5.2 Function Documentation

6.5.2.1 `ssize_t keyAddBaseName (Key * key, const char * baseName)`

Adds `baseName` to the current key name.

Assumes that `key` is a directory. `baseName` is appended to it. The function adds `'/'` if needed while concatenating.

So if `key` has name `"system/dir1/dir2"` and this method is called with `baseName "mykey"`, the resulting key will have name `"system/dir1/dir2/mykey"`.

When `baseName` is 0 or "" nothing will happen and the size of the name is returned.

Warning

You should not change a keys name once it belongs to a keyset. See `ksSort()` for more information.

TODO: does not recognise `..` and `.` in the string!

Parameters

<code>key</code>	the key object to work with
<code>baseName</code>	the string to append to the name

Returns

the size in bytes of the new key name including the ending NULL
 -1 if the key had no name
 -1 on NULL pointers

See Also

[keySetBaseName\(\)](#)
[keySetName\(\)](#) to set a new name.

6.5.2.2 `const char* keyBaseName (const Key * key)`

Returns a pointer to the real internal key name where the `basename` starts.

This is a much more efficient version of [keyGetBaseName\(\)](#) and you should use it if you are responsible enough to not mess up things. The name might change or even point to a wrong place after a [keySetName\(\)](#). If you need a copy of the `basename` consider to use [keyGetBaseName\(\)](#).

[keyBaseName\(\)](#) returns "" when there is no `keyBaseName`. The reason is

```
key=keyNew(0);
keySetName(key, "");
keyBaseName(key); // you would expect "" here
keySetName(key, "user");
keyBaseName(key); // you would expect "" here
keyDel(key);
```

Note

Note that the Key structure keeps its own size field that is calculated by library internal calls, so to avoid inconsistencies, you must never use the pointer returned by [keyBaseName\(\)](#) method to set a new value. Use [keySetBaseName\(\)](#) instead.

Parameters

<i>key</i>	the object to obtain the basename from
------------	--

Returns

a pointer to the basename
 "" when the key has no (base)name
 0 on NULL pointer

See Also

[keyGetBaseName\(\)](#), [keyGetBaseNameSize\(\)](#)
[keyName\(\)](#) to get a pointer to the name
[keyOwner\(\)](#) to get a pointer to the owner

6.5.2.3 `ssize_t keyGetBaseName (const Key * key, char * returned, size_t maxSize)`

Calculate the basename of a key name and put it in `returned` finalizing the string with NULL.

Some examples:

- basename of `system/some/keyname` is `keyname`
- basename of `"user/tmp/some key"` is `"some key"`

Parameters

<i>key</i>	the key to extract basename from
<i>returned</i>	a pre-allocated buffer to store the basename
<i>maxSize</i>	size of the <code>returned</code> buffer

Returns

number of bytes copied to `returned`
 1 on empty name
 -1 on NULL pointers
 -1 when `maxSize` is 0 or larger than `SSIZE_MAX`

See Also

[keyBaseName\(\)](#), [keyGetBaseNameSize\(\)](#)
[keyName\(\)](#), [keyGetName\(\)](#), [keySetName\(\)](#)

6.5.2.4 `ssize_t keyGetBaseNameSize (const Key * key)`

Calculates number of bytes needed to store basename of `key`.

Key names that have only root names (e.g. `"system"` or `"user"` or `"user:domain"`) does not have base-names, thus the function will return 1 bytes to store "".

Basenames are denoted as:

- `system/some/thing/basename -> basename`
- `user:domain/some/thing/base\name > base\name`

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

size in bytes of *key*'s basename including ending NULL

See Also

[keyBaseName\(\)](#), [keyGetBaseName\(\)](#)
[keyName\(\)](#), [keyGetName\(\)](#), [keySetName\(\)](#)

6.5.2.5 `ssize_t keyGetFullName (const Key * key, char * returnedName, size_t maxSize)`

Get key full name, including the user domain name.

Returns

number of bytes written
 1 on empty name
 -1 on NULL pointers
 -1 if maxSize is 0 or larger than SSIZE_MAX

Parameters

<i>key</i>	the key object
<i>returnedName</i>	pre-allocated memory to write the key name
<i>maxSize</i>	maximum number of bytes that will fit in returnedName, including the final NULL

6.5.2.6 `ssize_t keyGetFullNameSize (const Key * key)`

Bytes needed to store the key name including user domain and ending NULL.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

number of bytes needed to store key name including user domain
 1 on empty name
 -1 on NULL pointer

See Also

[keyGetFullName\(\)](#), [keyGetNameSize\(\)](#)

6.5.2.7 `ssize_t keyGetName (const Key * key, char * returnedName, size_t maxSize)`

Get abbreviated key name (without owner name).

When there is not enough space to write the name, nothing will be written and -1 will be returned.

maxSize is limited to SSIZE_MAX. When this value is exceeded -1 will be returned. The reason for that is that any value higher is just a negative return value passed by accident. Of course malloc is not as failure tolerant and will try to allocate.

```
char *getBack = malloc (keyGetNameSize(key));
keyGetName(key, getBack, keyGetNameSize(key));
```

Returns

- number of bytes written to returnedName
- 1 when only a null was written
- 1 when keyname is longer then maxSize or 0 or any NULL pointer

Parameters

<i>key</i>	the key object to work with
<i>returnedName</i>	pre-allocated memory to write the key name
<i>maxSize</i>	maximum number of bytes that will fit in returnedName, including the final NULL

See Also

[keyGetNameSize\(\)](#), [keyGetFullName\(\)](#), [keyGetFullNameSize\(\)](#)

6.5.2.8 `ssize_t keyGetNameSize (const Key * key)`

Bytes needed to store the key name without owner.

For an empty key name you need one byte to store the ending NULL. For that reason 1 is returned.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

- number of bytes needed, including ending NULL, to store key name without owner
- 1 if there is is no key Name
- 1 on NULL pointer

See Also

[keyGetName\(\)](#), [keyGetFullNameSize\(\)](#)

6.5.2.9 `ssize_t keyGetOwner (const Key * key, char * returnedOwner, size_t maxSize)`

Return the owner of the key.

- Given `user:someuser/.....` return `someuser`
- Given `user:some.user/....` return `some.user`

- Given `user/...` return the current user

Only `user/...` keys have a owner. For `system/...` keys (that doesn't have a key owner) an empty string ("") is returned.

Although usually the same, the owner of a key is not related to its UID. Owner are related to WHERE the key is stored on disk, while UIDs are related to mode controls of a key.

Parameters

<i>key</i>	the object to work with
<i>returnedOwner</i>	a pre-allocated space to store the owner
<i>maxSize</i>	maximum number of bytes that fit returned

Returns

- number of bytes written to buffer
- 1 if there is no owner
- 1 on NULL pointers
- 1 when `maxSize` is 0, larger than `SSIZE_MAX` or too small for `ownername`

See Also

[keySetName\(\)](#), [keySetOwner\(\)](#), [keyOwner\(\)](#), [keyGetFullName\(\)](#)

6.5.2.10 `ssize_t keyGetOwnerSize (const Key * key)`

Return the size of the owner of the Key with concluding 0.

The returned number can be used to allocate a string. 1 will returned on an empty owner to store the concluding 0 on using [keyGetOwner\(\)](#).

```
char * buffer;
buffer = malloc (keyGetOwnerSize (key));
// use buffer and keyGetOwnerSize (key) for maxSize
```

Note

that -1 might be returned on null pointer, so when you directly allocate afterwards its best to check if you will pass a null pointer before.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

- number of bytes
- 1 if there is no owner
- 1 on NULL pointer

See Also

[keyGetOwner\(\)](#)

6.5.2.11 `const char* keyName (const Key * key)`

Returns a pointer to the abbreviated real internal `key` name.

This is a much more efficient version of `keyGetName()` and can use it if you are responsible enough to not mess up things. You are not allowed to change anything in the returned array. The content of that string may change after `keySetName()` and similar functions. If you need a copy of the name, consider using `keyGetName()`.

The name will be without owner, see `keyGetFullName()` if you need the name with its owner.

`keyName()` returns "" when there is no `keyName`. The reason is

```
key=keyNew(0);
keySetName(key, "");
keyName(key); // you would expect "" here
keyDel(key);
```

Note

Note that the `Key` structure keeps its own size field that is calculated by library internal calls, so to avoid inconsistencies, you must never use the pointer returned by `keyName()` method to set a new value. Use `keySetName()` instead.

Parameters

<code>key</code>	the key object to work with
------------------	-----------------------------

Returns

a pointer to the keyname which must not be changed.

"" when there is no (a empty) keyname

0 on NULL pointer

See Also

`keyGetNameSize()` for the string length
`keyGetFullName()`, `keyGetFullNameSize()` to get the full name
`keyGetName()` as alternative to get a copy
`keyOwner()` to get a pointer to owner

6.5.2.12 `const char* keyOwner (const Key * key)`

Return a pointer to the real internal `key` owner.

This is a much more efficient version of `keyGetOwner()` and you should use it if you are responsible enough to not mess up things. You are not allowed to modify the returned string in any way. If you need a copy of the string, consider to use `keyGetOwner()` instead.

`keyOwner()` returns "" when there is no `keyOwner`. The reason is

```
key=keyNew(0);
keySetOwner(key, "");
keyOwner(key); // you would expect "" here
keySetOwner(key, "system");
keyOwner(key); // you would expect "" here
```

Note

Note that the `Key` structure keeps its own size field that is calculated by library internal calls, so to avoid inconsistencies, you must never use the pointer returned by `keyOwner()` method to set a new value. Use `keySetOwner()` instead.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

a pointer to internal owner
 "" when there is no (a empty) owner
 0 on NULL pointer

See Also

[keyGetOwnerSize\(\)](#) for the size of the string with concluding 0
[keyGetOwner\(\)](#), [keySetOwner\(\)](#)
[keyName\(\)](#) for name without owner
[keyGetFullName\(\)](#) for name with owner

6.5.2.13 `ssize_t keySetBaseName (Key * key, const char * baseName)`

Sets *baseName* as the new basename for *key*.

All text after the last ' / ' in the *key* keyname is erased and *baseName* is appended.

So lets suppose *key* has name "system/dir1/dir2/key1". If *baseName* is "key2", the resulting key name will be "system/dir1/dir2/key2". If *baseName* is empty or NULL, the resulting key name will be "system/dir1/dir2".

Warning

You should not change a keys name once it belongs to a keyset. See [ksSort\(\)](#) for more information.

TODO: does not work with .. and .

Parameters

<i>key</i>	the key object to work with
<i>baseName</i>	the string used to overwrite the basename of the key

Returns

the size in bytes of the new key name
 -1 on NULL pointers

See Also

[keyAddBaseName\(\)](#)
[keySetName\(\)](#) to set a new name

6.5.2.14 `ssize_t keySetName (Key * key, const char * newName)`

Set a new name to a key.

A valid name is of the forms:

- system/something
- user/something

- `user:username/something`

The last form has explicitly set the owner, to let the library know in which user folder to save the key. A owner is a user name. If it is not defined (the second form) current user is used.

You should always follow the guidelines for key tree structure creation.

A private copy of the key name will be stored, and the `newName` parameter can be freed after this call.

.., . and / will be handled correctly. A valid name will be build out of the (valid) name what you pass, e.g. `user///sw/./sw/././MyApp -> user/sw/MyApp`

On invalid names, NULL or "" the name will be "" afterwards.

Warning

You shall not change a key name once it belongs to a keyset.

Return values

<i>size</i>	in bytes of this new key name including ending NULL
<i>0</i>	if <code>newName</code> is an empty string or a NULL pointer (name will be empty afterwards)
<i>-1</i>	if <code>newName</code> is invalid (name will be empty afterwards)

Parameters

<i>key</i>	the key object to work with
<i>newName</i>	the new key name

See Also

[keyNew\(\)](#), [keySetOwner\(\)](#)
[keyGetName\(\)](#), [keyGetFullName\(\)](#), [keyName\(\)](#)
[keySetBaseName\(\)](#), [keyAddBaseName\(\)](#) to manipulate a name

6.5.2.15 `ssize_t keySetOwner (Key * key, const char * newOwner)`

Set the owner of a key.

A owner is a name of a system user related to a UID. The owner decides on which location on the disc the key goes.

A private copy is stored, so the passed parameter can be freed after the call.

Parameters

<i>key</i>	the key object to work with
<i>newOwner</i>	the string which describes the owner of the key

Returns

the number of bytes actually saved including final NULL
 1 when owner is freed (by setting 0 or "")
 -1 on null pointer or memory problems

See Also

[keySetName\(\)](#), [keyGetOwner\(\)](#), [keyGetFullName\(\)](#)

6.6 KeySet

Methods to manipulate KeySets.

Functions

- KeySet * [ksNew](#) (size_t alloc,...)
- KeySet * [ksDup](#) (const KeySet *source)
- int [ksCopy](#) (KeySet *dest, const KeySet *source)

- int [ksDel](#) (KeySet *ks)
- ssize_t [ksGetSize](#) (const KeySet *ks)
- ssize_t [ksAppendKey](#) (KeySet *ks, Key *toAppend)

- ssize_t [ksAppend](#) (KeySet *ks, const KeySet *toAppend)

- KeySet * [ksCut](#) (KeySet *ks, const Key *cutpoint)
- Key * [ksPop](#) (KeySet *ks)
- int [ksRewind](#) (KeySet *ks)
- Key * [ksNext](#) (KeySet *ks)
- Key * [ksCurrent](#) (const KeySet *ks)
- Key * [ksHead](#) (const KeySet *ks)
- Key * [ksTail](#) (const KeySet *ks)
- cursor_t [ksGetCursor](#) (const KeySet *ks)
- int [ksSetCursor](#) (KeySet *ks, cursor_t cursor)
- Key * [ksLookup](#) (KeySet *ks, Key *key, option_t options)
- Key * [ksLookupByName](#) (KeySet *ks, const char *name, option_t options)

6.6.1 Detailed Description

Methods to manipulate KeySets. A KeySet is a sorted set of keys. So the correct name actually would be KeyMap.

With [ksNew\(\)](#) you can create a new KeySet.

You can add keys with [ksAppendKey\(\)](#) in the keyset. Using [ksAppend\(\)](#) you can append a whole keyset.

Note

Because the key is not copied, also the pointer to the current metadata [keyNextMeta\(\)](#) will be shared.

[ksGetSize\(\)](#) tells you the current size of the keyset.

With [ksRewind\(\)](#) and [ksNext\(\)](#) you can navigate through the keyset. Don't expect any particular order, but it is assured that you will get every key of the set.

KeySets have an [internal cursor](#) . This is used for [ksLookup\(\)](#) and [kdbSet\(\)](#).

KeySet has a fundamental meaning inside elektra. It makes it possible to get and store many keys at once inside the database. In addition to that the class can be used as high level datastructure in applications. With [ksLookup-ByName\(\)](#) it is possible to fetch easily specific keys out of the list of keys.

You can easily create and iterate keys:

```
#include <kdb.h>

// create a new keyset with 3 keys
// with a hint that about 20 keys will be inside
KeySet *myConfig = ksNew(20,
    keyNew ("user/name1", 0),
    keyNew ("user/name2", 0),
```

```

        keyNew ("user/name3", 0),
        KS_END);
// append a key in the keyset
ksAppendKey(myConfig, keyNew("user/name4", 0));

Key *current;
ksRewind(myConfig);
while ((current=ksNext(myConfig))!=0)
{
    printf("Key name is %s.\n", keyName (current));
}
ksDel (myConfig); // delete keyset and all keys appended

```

6.6.2 Function Documentation

6.6.2.1 ssize_t ksAppend (KeySet * ks, const KeySet * toAppend)

Append all toAppend contained keys to the end of the ks.

toAppend KeySet will be left unchanged.

If a key is both in toAppend and ks, the Key in ks will be overridden.

Note

Because the key is not copied, also the pointer to the current metadata [keyNextMeta\(\)](#) will be shared.

Postcondition

Sorted KeySet ks with all keys it had before and additionally the keys from toAppend

Returns

the size of the KeySet after transfer
-1 on NULL pointers

Parameters

<i>ks</i>	the KeySet that will receive the keys
<i>toAppend</i>	the KeySet that provides the keys that will be transfered

See Also

[ksAppendKey\(\)](#)

6.6.2.2 ssize_t ksAppendKey (KeySet * ks, Key * toAppend)

Appends a Key to the end of ks.

A pointer to the key will be stored, and not a private copy. So a future [ksDel\(\)](#) on ks may [keyDel\(\)](#) the toAppend object, see [keyGetRef\(\)](#).

The reference counter of the key will be incremented, and thus toAppend is not const.

Note

Because the key is not copied, also the pointer to the current metadata [keyNextMeta\(\)](#) will be shared.

If the keyname already existed, it will be replaced with the new key.

The KeySet internal cursor will be set to the new key.

Returns

- the size of the KeySet after insertion
- 1 on NULL pointers
- 1 if insertion failed, the key will be deleted then.

Parameters

<i>ks</i>	KeySet that will receive the key
<i>toAppend</i>	Key that will be appended to ks

See Also

[ksAppend\(\)](#), [keyNew\(\)](#), [ksDel\(\)](#)
[keyIncRef\(\)](#)

6.6.2.3 int ksCopy (KeySet * *dest*, const KeySet * *source*)

Copy a keyset.

Most often you may want a duplicate of a keyset, see [ksDup\(\)](#) or append keys, see [ksAppend\(\)](#). But in some situations you need to copy a keyset to a existing keyset, for that this function exists.

You can also use it to clear a keyset when you pass a NULL pointer as *source*.

Note that all keys in *dest* will be deleted. Afterwards the content of the source will be added to the destination and the [ksCurrent\(\)](#) is set properly in *dest*.

A flat copy is made, so the keys will not be duplicated, but there reference counter is updated, so both keysets need to be [ksDel\(\)](#).

Note

Because the key is not copied, also the pointer to the current metadata [keyNextMeta\(\)](#) will be shared.

```
int f (KeySet *ks)
{
    KeySet *c = ksNew (20, ..., KS_END);
    // c receives keys
    ksCopy (ks, c); // pass the keyset to the caller

    ksDel (c);
    // caller needs to ksDel (ks)
}
```

Parameters

<i>source</i>	has to be an initialized source KeySet or NULL
<i>dest</i>	has to be an initialized KeySet where to write the keys

Returns

- 1 on success
- 0 if dest was cleared successfully (source is NULL)
- 1 on NULL pointer

See Also

[ksNew\(\)](#), [ksDel\(\)](#), [ksDup\(\)](#)
[keyCopy\(\)](#) for copying keys

6.6.2.4 Key* ksCurrent (const KeySet * ks)

Return the current Key.

The pointer is NULL if you reached the end or after [ksRewind\(\)](#).

Note

You must not delete the key or change the key, use [ksPop\(\)](#) if you want to delete it.

Parameters

<i>ks</i>	the keyset object to work with
-----------	--------------------------------

Returns

pointer to the Key pointed by *ks*'s cursor
0 on NULL pointer

See Also

[ksNext\(\)](#), [ksRewind\(\)](#)

6.6.2.5 KeySet* ksCut (KeySet * ks, const Key * cutpoint)

Cuts out a keyset at the cutpoint.

Searches for the cutpoint inside the KeySet *ks*. If found it cuts out everything which is below (see [keysBelow\(\)](#)) this key. If not found an empty keyset is returned.

The cursor will stay at the same key as it was before. If the cursor was inside the region of cutted (moved) keys, the cursor will be set to the key before the cutpoint.

Returns

a new allocated KeySet which needs to be deleted with [ksDel\(\)](#). The keyset consists of all keys (of the original keyset *ks*) below the cutpoint. If the key cutpoint exists, it will also be appended.

Return values

0	on null pointers, no key name or allocation problems
---	--

Parameters

<i>ks</i>	the keyset to cut. It will be modified by removing all keys below the cutpoint. The cutpoint itself will also be removed.
<i>cutpoint</i>	the point where to cut out the keyset

6.6.2.6 int ksDel (KeySet * ks)

A destructor for KeySet objects.

Cleans all internal dynamic attributes, decrement all reference pointers to all keys and then [keyDel\(\)](#) all contained Keys, and free()s the release the KeySet object memory (that was previously allocated by [ksNew\(\)](#)).

Parameters

<i>ks</i>	the keyset object to work with
-----------	--------------------------------

Returns

0 when the keyset was freed
-1 on null pointer

See Also

[ksNew\(\)](#)

6.6.2.7 KeySet* ksDup (const KeySet * *source*)

Return a duplicate of a keyset.

Objects created with [ksDup\(\)](#) must be destroyed with [ksDel\(\)](#).

Memory will be allocated as needed for dynamic properties, so you need to [ksDel\(\)](#) the returned pointer.

A flat copy is made, so the keys will not be duplicated, but there reference counter is updated, so both keysets need [ksDel\(\)](#).

Parameters

<i>source</i>	has to be an initialised source KeySet
---------------	--

Returns

a flat copy of source on success
0 on NULL pointer

See Also

[ksNew\(\)](#), [ksDel\(\)](#)
[keyDup\(\)](#) for [Key](#) duplication

6.6.2.8 cursor_t ksGetCursor (const KeySet * *ks*)

Get the KeySet internal cursor.

Use it to get the cursor of the actual position.

Warning

Cursors are getting invalid when the key was [ksPop\(\)](#)ed or [ksLookup\(\)](#) with KDB_O_POP was used.

6.6.3 Read ahead

With the cursors it is possible to read ahead in a keyset:

```
cursor_t jump;
ksRewind (ks);
while ((key = keyNextMeta (ks)) != 0)
{
    // now mark this key
    jump = ksGetCursor(ks);
}
```

```

//code..
keyNextMeta (ks); // now browse on
// use ksCurrent(ks) to check the keys
//code..

// jump back to the position marked before
ksSetCursor (ks, jump);
}

```

6.6.4 Restoring state

It can also be used to restore the state of a keyset in a function

```

int f (KeySet *ks)
{
    cursor_t state = ksGetCursor(ks);

    // work with keyset

    // now bring the keyset to the state before
    ksSetCursor (ks, state);
}

```

It is of course possible to make the KeySet const and cast its const away to set the cursor. Another way to achieve the same is to [ksDup\(\)](#) the keyset, but it is not as efficient.

An invalid cursor will be returned directly after [ksRewind\(\)](#). When you set an invalid cursor [ksCurrent\(\)](#) is 0 and [ksNext\(\) == ksHead\(\)](#).

Note

Only use a cursor for the same keyset which it was made for.

Parameters

<i>ks</i>	the keyset object to work with
-----------	--------------------------------

Returns

a valid cursor on success
 an invalid cursor on NULL pointer or after [ksRewind\(\)](#)

See Also

[ksNext\(\)](#), [ksSetCursor\(\)](#)

6.6.4.1 ssize_t ksGetSize (const KeySet * ks)

Return the number of keys that *ks* contains.

Parameters

<i>ks</i>	the keyset object to work with
-----------	--------------------------------

Returns

the number of keys that *ks* contains.
 -1 on NULL pointer

See Also

`ksNew(0)`, `ksDel()`

6.6.4.2 `Key* ksHead (const KeySet * ks)`

Return the first key in the KeySet.

The KeySets cursor will not be effected.

If `ksCurrent()==ksHead()` you know you are on the first key.

Parameters

<code>ks</code>	the keyset object to work with
-----------------	--------------------------------

Returns

the first Key of a keyset

0 on NULL pointer or empty keyset

See Also

`ksTail()` for the last `Key`

`ksRewind()`, `ksCurrent()` and `ksNext()` for iterating over the `KeySet`

6.6.4.3 `Key* ksLookup (KeySet * ks, Key * key, option_t options)`

Look for a Key contained in `ks` that matches the name of the `key`.

6.6.5 Introduction

`ksLookup()` is designed to let you work with entirely pre-loaded KeySets, so instead of `kdbGetKey()`, key by key, the idea is to fully `kdbGet()` for your application root key and process it all at once with `ksLookup()`.

This function is very efficient by using binary search. Together with `kdbGet()` which can you load the whole configuration with only some communication to backends you can write very effective but short code for configuration.

6.6.6 Usage

If found, `ks` internal cursor will be positioned in the matched key (also accessible by `ksCurrent()`), and a pointer to the Key is returned. If not found, `ks` internal cursor will not move, and a NULL pointer is returned.

Cascading is done if the first character is a `/`. This leads to ignoring the prefix like `system/` and `user/`.

```
if (kdbGet(handle, "user/myapp", myConfig, 0) == -1)
    errorHandler ("Could not get Keys");

if (kdbGet(handle, "system/myapp", myConfig, 0) == -1)
    errorHandler ("Could not get Keys");

if ((myKey = ksLookup(myConfig, key, 0)) == NULL)
    errorHandler ("Could not Lookup Key");
```

This is the way multi user Programs should get there configuration and search after the values. It is guaranteed that more namespaces can be added easily and that all values can be set by admin and user.

6.6.6.1 KDB_O_NOALL

When KDB_O_NOALL is set the keyset will be only searched from [ksCurrent\(\)](#) to [ksTail\(\)](#). You need to [ksRewind\(\)](#) the keyset yourself. [ksCurrent\(\)](#) is always set properly after searching a key, so you can go on searching another key after the found key.

When KDB_O_NOALL is not set the cursor will stay untouched and all keys are considered. A much more efficient binary search will be used then.

6.6.6.2 KDB_O_POP

When KDB_O_POP is set the key which was found will be [ksPop\(\)](#)ed. [ksCurrent\(\)](#) will not be changed, only iff [ksCurrent\(\)](#) is the searched key, then the keyset will be [ksRewind\(\)](#)ed.

Note

Like in [ksPop\(\)](#) the popped key always needs to be [keyDel\(\)](#) afterwards, even if it is appended to another keyset.

Warning

All cursors on the keyset will be invalid iff you use KDB_O_POP, so don't use this if you rely on a cursor, see [ksGetCursor\(\)](#).

You can solve this problem by using KDB_O_NOALL, risking you have to iterate n^2 instead of n .

The more elegant way is to separate the keyset you use for [ksLookup\(\)](#) and [ksAppendKey\(\)](#):

```
int f(KeySet *iterator, KeySet *lookup)
{
    KeySet *append = ksNew (ksGetSize(lookup), KS_END);
    Key *key;
    Key *current;

    ksRewind(iterator);
    while (current=ksNext(iterator))
    {
        key = ksLookup (lookup, current, KDB_O_POP);
        // do something...
        ksAppendKey(append, key); // now append it to
append, not lookup!
        keyDel (key); // make sure to ALWAYS delete popped keys.
    }
    ksAppend(lookup, append);
    // now lookup needs to be sorted only once, append never
    ksDel (append);
}
```

Parameters

<i>ks</i>	where to look for
<i>key</i>	the key object you are looking for
<i>options</i>	<p>some KDB_O_* option bits:</p> <ul style="list-style-type: none"> • KDB_O_NOCASE Lookup ignoring case. • KDB_O_WITHOWNER Also consider correct owner. • KDB_O_NOALL Only search from ksCurrent() to end of keyset, see above text. • KDB_O_POP Pop the key which was found. • KDB_O_DEL Delete the passed key.

Returns

pointer to the Key found, 0 otherwise
0 on NULL pointers

See Also

[ksLookupByName\(\)](#) to search by a name given by a string
[ksCurrent\(\)](#), [ksRewind\(\)](#), [ksNext\(\)](#) for iterating over a [KeySet](#)

6.6.6.3 Key* ksLookupByName (KeySet * ks, const char * name, option_t options)

Look for a Key contained in `ks` that matches `name`.

[ksLookupByName\(\)](#) is designed to let you work with entirely pre-loaded KeySets, so instead of `kdbGetKey()`, key by key, the idea is to fully `kdbGetByName()` for your application root key and process it all at once with [ksLookupByName\(\)](#).

This function is very efficient by using binary search. Together with `kdbGetByName()` which can you load the whole configuration with only some communication to backends you can write very effective but short code for configuration.

If found, `ks` internal cursor will be positioned in the matched key (also accessible by [ksCurrent\(\)](#)), and a pointer to the Key is returned. If not found, `ks` internal cursor will not move, and a NULL pointer is returned. If requested to pop the key, the cursor will be rewinded.

6.6.7 Cascading

Cascading is done if the first character is a `/`. This leads to ignoring the prefix like `system/` and `user/`.

```
if (kdbGet(handle, "user/sw/myapp/current", myConfig, parentKey) == -1)
    errorHandler ("Could not get Keys", parentKey);

if (kdbGet(handle, "system/sw/myapp/current", myConfig, parentKey) == -1)
    errorHandler ("Could not get Keys", parentKey);

if ((myKey = ksLookupByName (myConfig, "/myapp/current/key", 0))
    == NULL)
    errorHandler ("Could not Lookup Key");
```

This is the way multi user Programs should get there configuration and search after the values. It is guaranteed that more namespaces can be added easily and that all values can be set by admin and user.

It is up to the application to implement a sophisticated cascading algorithm, for e.g. a list of profiles (specific, group and fallback):

```
if ((myKey = ksLookupByName (myConfig, "
/myapp/current/specific/key", 0)) == NULL)
    if ((myKey = ksLookupByName (myConfig, "
/myapp/current/group/key", 0)) == NULL)
        if ((myKey = ksLookupByName (myConfig, "
/myapp/current/fallback/key", 0)) == NULL)
            errorHandler ("All fallbacks failed to lookup key");
```

Note that for every profile both the user and the system key are searched. The first key found will be used.

6.6.8 Full Search

When `KDB_O_NOALL` is set the keyset will be only searched from [ksCurrent\(\)](#) to [ksTail\(\)](#). You need to [ksRewind\(\)](#) the keyset yourself. [ksCurrent\(\)](#) is always set properly after searching a key, so you can go on searching another key after the found key.

When `KDB_O_NOALL` is not set the cursor will stay untouched and all keys are considered. A much more efficient binary search will be used then.

Parameters

<i>ks</i>	where to look for
<i>name</i>	key name you are looking for
<i>options</i>	some KDB_O_* option bits: <ul style="list-style-type: none"> • KDB_O_NOCASE Lookup ignoring case. • KDB_O_WITHOWNER Also consider correct owner. • KDB_O_NOALL Only search from ksCurrent() to end of keyset, see above text. • KDB_O_POP Pop the key which was found.

Currently no options supported.

Returns

pointer to the Key found, 0 otherwise
0 on NULL pointers

See Also

[keyCompare\(\)](#) for very powerfull Key lookups in KeySets
[ksCurrent\(\)](#), [ksRewind\(\)](#), [ksNext\(\)](#)

6.6.8.1 KeySet* ksNew (size_t alloc, ...)

Allocate, initialize and return a new KeySet object.

Objects created with [ksNew\(\)](#) must be destroyed with [ksDel\(\)](#).

You can use a various long list of parameters to preload the keyset with a list of keys. Either your first and only parameter is 0 or your last parameter must be KEY_END.

So, terminate with [ksNew\(0\)](#) or [ksNew\(20, ..., KS_END\)](#)

For most uses

```
KeySet *keys = ksNew(0);
// work with it
ksDel (keys);
```

goes ok, the alloc size will be 16, defined in kdbprivate.h. The alloc size will be doubled whenever size reaches alloc size, so it also performs out large keysets.

But if you have any clue how large your keyset may be you should read the next statements.

If you want a keyset with length 15 (because you know of your application that you normally need about 12 up to 15 keys), use:

```
KeySet * keys = ksNew (15,
    keyNew ("user/sw/app/fixedConfiguration/key01", KEY_SWITCH_VALUE,
"value01", 0),
    keyNew ("user/sw/app/fixedConfiguration/key02", KEY_SWITCH_VALUE,
"value02", 0),
    keyNew ("user/sw/app/fixedConfiguration/key03", KEY_SWITCH_VALUE,
"value03", 0),
    // ...
    keyNew ("user/sw/app/fixedConfiguration/key15", KEY_SWITCH_VALUE,
"value15", 0),
    KS_END);
// work with it
ksDel (keys);
```

If you start having 3 keys, and your application needs approximately 200-500 keys, you can use:

```
KeySet * config = ksNew (500,
    keyNew ("user/sw/app/fixedConfiguration/key1", KEY_SWITCH_VALUE,
        "value1", 0),
    keyNew ("user/sw/app/fixedConfiguration/key2", KEY_SWITCH_VALUE,
        "value2", 0),
    keyNew ("user/sw/app/fixedConfiguration/key3", KEY_SWITCH_VALUE,
        "value3", 0),
    KS_END); // don't forget the KS_END at the end!
// work with it
ksDel (config);
```

Alloc size is 500, the size of the keyset will be 3 after ksNew. This means the keyset will reallocate when appending more than 497 keys.

The main benefit of taking a list of variant length parameters is to be able to have one C-Statement for any possible KeySet.

Due to ABI compatibility, the KeySet structure is only declared in kdb.h, and not defined. So you can only declare pointers to KeySets in your program. See <http://tldp.org/HOWTO/Program-Library-HOWTO/shared-libraries.html#AEN135>

See Also

[ksDel\(\)](#) to free the [KeySet](#) afterwards
[ksDup\(\)](#) to duplicate an existing [KeySet](#)

Parameters

<i>alloc</i>	gives a hint for the size how many Keys may be stored initially
--------------	---

Returns

a ready to use KeySet object
0 on memory error

6.6.8.2 Key* ksNext (KeySet * ks)

Returns the next Key in a KeySet.

KeySets have an internal cursor that can be reset with [ksRewind\(\)](#). Every time [ksNext\(\)](#) is called the cursor is incremented and the new current Key is returned.

You'll get a NULL pointer if the key after the end of the KeySet was reached. On subsequent calls of [ksNext\(\)](#) it will still return the NULL pointer.

The *ks* internal cursor will be changed, so it is not const.

Note

You must not delete or change the key, use [ksPop\(\)](#) if you want to delete it.

Parameters

<i>ks</i>	the keyset object to work with
-----------	--------------------------------

Returns

the new current Key
0 when the end is reached
0 on NULL pointer

See Also

[ksRewind\(\)](#), [ksCurrent\(\)](#)

6.6.8.3 Key* ksPop (KeySet * ks)

Remove and return the last key of `ks`.

The reference counter will be decremented by one.

The KeySets cursor will not be effected if it did not point to the popped key.

Note

You need to [keyDel\(\)](#) the key afterwards, if you don't append it to another keyset. It has the same semantics like a key allocated with [keyNew\(\)](#) or [keyDup\(\)](#).

```
ks1=ksNew(0);
ks2=ksNew(0);

k1=keyNew("user/name", KEY_END); // ref counter 0
ksAppendKey(ks1, k1); // ref counter 1
ksAppendKey(ks2, k1); // ref counter 2

k1=ksPop (ks1); // ref counter 1
k1=ksPop (ks2); // ref counter 0, like after keyNew()

ksAppendKey(ks1, k1); // ref counter 1

ksDel (ks1); // key is deleted too
ksDel (ks2);
*
```

Returns

the last key of `ks`

NULL if `ks` is empty or on NULL pointer

Parameters

<code>ks</code>	KeySet to work with
-----------------	---------------------

See Also

[ksAppendKey\(\)](#), [ksAppend\(\)](#)
[commandList\(\)](#) for an example

6.6.8.4 int ksRewind (KeySet * ks)

Rewinds the KeySet internal cursor.

Use it to set the cursor to the beginning of the KeySet. [ksCurrent\(\)](#) will then always return NULL afterwards. So you want to [ksNext\(\)](#) first.

```
ksRewind (ks);
while ((key = ksNext (ks)) !=0) {}
```

Parameters

<code>ks</code>	the keyset object to work with
-----------------	--------------------------------

Returns

0 on success
-1 on NULL pointer

See Also

[ksNext\(\)](#), [ksCurrent\(\)](#)

6.6.8.5 int ksSetCursor (KeySet * ks, cursor_t cursor)

Set the KeySet internal cursor.

Use it to set the cursor to a stored position. [ksCurrent\(\)](#) will then be the position which you got with.

Warning

Cursors may get invalid when the key was [ksPop\(\)](#)ed or [ksLookup\(\)](#) was used together with KDB_O_POP.

```
cursor_t cursor;
..
// key now in any position here
cursor = ksGetCursor (ks);
while ((key = keyNextMeta (ks)) != 0) {}
ksSetCursor (ks, cursor); // reset state
ksCurrent(ks); // in same position as before
```

An invalid cursor will set the keyset to its beginning like [ksRewind\(\)](#). When you set an invalid cursor [ksCurrent\(\)](#) is 0 and [ksNext\(\)](#) == [ksHead\(\)](#).

Parameters

<i>cursor</i>	the cursor to use
<i>ks</i>	the keyset object to work with

Returns

0 when the keyset is [ksRewind\(\)](#)ed
1 otherwise
-1 on NULL pointer

See Also

[ksNext\(\)](#), [ksGetCursor\(\)](#)

6.6.8.6 Key* ksTail (const KeySet * ks)

Return the last key in the KeySet.

The KeySets cursor will not be effected.

If [ksCurrent\(\)](#)==[ksTail\(\)](#) you know you are on the last key. [ksNext\(\)](#) will return a NULL pointer afterwards.

Parameters

<i>ks</i>	the keyset object to work with
-----------	--------------------------------

Returns

the last Key of a keyset
0 on NULL pointer or empty keyset


See Also

[ksHead\(\)](#) for the first [Key](#)
[ksRewind\(\)](#), [ksCurrent\(\)](#) and [ksNext\(\)](#) for iterating over the [KeySet](#)

6.7 Methods for Making Tests

Methods to do various tests on Keys.

Collaboration diagram for Methods for Making Tests:



Functions

- `int keyCmp (const Key *k1, const Key *k2)`
- `int keyNeedSync (const Key *key)`
- `int keyIsSystem (const Key *key)`
- `int keyIsUser (const Key *key)`
- `int keyIsBelow (const Key *key, const Key *check)`
- `int keyIsDirectBelow (const Key *key, const Key *check)`
- `int keyRel (const Key *key, const Key *check)`
- `int keyIsInactive (const Key *key)`
- `int keyIsDir (const Key *key)`
- `int keyIsBinary (const Key *key)`
- `int keyIsString (const Key *key)`
- `keyswitch_t keyCompare (const Key *key1, const Key *key2)`

6.7.1 Detailed Description

Methods to do various tests on Keys. To use them:

```
#include <kdb.h>
```

6.7.2 Function Documentation

6.7.2.1 `int keyCmp (const Key * k1, const Key * k2)`

Compare the name of two keys.

Returns

a number less than, equal to or greater than zero if k1 is found, respectively, to be less than, to match, or be greater than k2.

The comparison is based on a `strcmp` of the keynames, and iff they match a `strcmp` of the owner will be used to distinguish. If even this matches the keys are found to be exactly the same and 0 is returned. These two keys can't be used in the same KeySet.

`keyCmp()` defines the sorting order for a KeySet.

The following 3 points are the rules for null values. They only take account when none of the preceding rules matched.

- A null pointer will be found to be smaller than every other key. If both are null pointers, 0 is returned.
- A null name will be found to be smaller than every other name. If both are null names, 0 is returned.
- No owner will be found to be smaller than every other owner. If both don't have a owner, 0 is returned.

Note

the owner will only be used if the names are equal.

Often is enough to know if the other key is less then or greater then the other one. But Sometimes you need more precise information, see [keyRel\(\)](#).

Given any Keys k1 and k2 constructed with [keyNew\(\)](#), following equation hold true:

```
// keyCmp(0,0) == 0
// keyCmp(k1,0) == 1
// keyCmp(0,k2) == -1
```

You can write similar equation for the other rules.

Here are some more examples with equation:

```
Key *k1 = keyNew("user/a", KEY_END);
Key *k2 = keyNew("user/b", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0

Key *k1 = keyNew("user/a", KEY_OWNER, "markus", KEY_END);
Key *k2 = keyNew("user/a", KEY_OWNER, "max", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0
```

Warning

Do not try to strcmp the [keyName\(\)](#) yourself because the used strcmp implementation is allowed to differ from simple ascii comparison.

Parameters

<i>k1</i>	the first key object to compare with
<i>k2</i>	the second key object to compare with

See Also

[ksAppendKey\(\)](#), [ksAppend\(\)](#) will compare keys when appending
[ksLookup\(\)](#) will compare keys during searching

6.7.2.2 keyswitch_t keyCompare (const Key * key1, const Key * key2)

Compare 2 keys.

The returned flags bit array has 1s (differ) or 0s (equal) for each key meta info compared, that can be logically ORed using [keyswitch_t](#) flags. [KEY_NAME](#) , [KEY_VALUE](#) , [KEY_OWNER](#) , [KEY_COMMENT](#) , [KEY_UID](#) , [KEY_GID](#) , [KEY_MODE](#) and

A very simple example would be

```
Key *key1, *key;
```

```

uint32_t changes;

// omitted key1 and key2 initialization and manipulation

changes=keyCompare(key1,key2);

if (changes == 0) printf("key1 and key2 are identicall\n");

if (changes & KEY_VALUE)
    printf("key1 and key2 have different values\n");

if (changes & KEY_UID)
    printf("key1 and key2 have different UID\n");

```

Example of very powerful specific Key lookup in a KeySet:

```

KDB *handle = kdbOpen();
KeySet *ks=ksNew(0);
Key *base = keyNew ("user/sw/MyApp/something", KEY_END);
Key *current;
uint32_t match;
uint32_t interests;

kdbGetByName(handle, ks, "user/sw/MyApp", 0);

// we are interested only in key type and access permissions
interests=(KEY_TYPE | KEY_MODE);

ksRewind(ks); // put cursor in the begining
while ((current=ksNext(ks)) {
    match=keyCompare(current,base);

    if ((~match & interests) == interests)
        printf("Key %s has same type and permissions of base key",
            keyName(current));

    // continue walking in the KeySet....
}

// now we want same name and/or value
interests=(KEY_NAME | KEY_VALUE);

// we don't really need ksRewind(), since previous loop achieved end of KeySet
ksRewind(ks);
while ((current=ksNext(ks)) {
    match=keyCompare(current,base);

    if ((~match & interests) == interests) {
        printf("Key %s has same name, value, and sync status
            of base key",keyName(current));
    }

    // continue walking in the KeySet....
}

keyDel(base);
ksDel(ks);
kdbClose (handle);

```

Returns

a bit array pointing the differences

Parameters

<i>key1</i>	first key
<i>key2</i>	second key

See Also

[keyswitch_t](#)

6.7.2.3 int keyIsBelow (const Key * key, const Key * check)

Check if the key check is below the key key or not.

Example:

```
key user/sw/app
check user/sw/app/key
```

returns true because check is below key

Example:

```
key user/sw/app
check user/sw/app/folder/key
```

returns also true because check is indirect below key

Parameters

<i>key</i>	the key object to work with
<i>check</i>	the key to find the relative position of

Returns

- 1 if check is below key
- 0 if it is not below or if it is the same key

See Also

[keySetName\(\)](#), [keyGetName\(\)](#), [keyIsDirectBelow\(\)](#)

6.7.2.4 int keyIsBinary (const Key * key)

Check if a key is binary type.

The function checks if the key is a binary. Opposed to string values binary values can have '\0' inside the value and may not be terminated by a null character. Their disadvantage is that you need to pass their size.

Make sure to use this function and don't test the binary type another way to ensure compatibility and to write less error prone programs.

Returns

- 1 if it is binary
- 0 if it is not
- 1 on NULL pointer

See Also

[keyGetBinary\(\)](#), [keySetBinary\(\)](#)

Parameters

<i>key</i>	the key to check
------------	------------------

6.7.2.5 int keyIsDir (const Key * key)

Check if the mode for the key has access privileges.

In the filesystem backend a key represented through a file has the mode 664, but a key represented through a folder

775. `keyIsDir()` checks if all 3 executeable bits are set.

If any executable bit is set it will be recognized as a directory.

Note

`keyIsDir` may return true even though you can't access the directory.

To know if you can access the directory, you need to check, if your

- user ID is equal the key's user ID and the mode & 100 is true
- group ID is equal the key's group ID and the mode & 010 is true
- mode & 001 is true

Accessing does not mean that you can get any value or comments below, see [Modes](#) for more information.

Note

currently mountpoints can only where `keyIsDir()` is true (0.7.0) but this is likely to change.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

- 1 if key is a directory, 0 otherwise
- 1 on NULL pointer

See Also

[keySetDir\(\)](#), [keySetMode\(\)](#)

6.7.2.6 `int keyIsDirectBelow (const Key * key, const Key * check)`

Check if the key check is direct below the key key or not.

Example:

```
key user/sw/app
check user/sw/app/key
```

returns true because check is below key

Example:

```
key user/sw/app
check user/sw/app/folder/key
```

does not return true, because there is only a indirect relation

Parameters

<i>key</i>	the key object to work with
<i>check</i>	the key to find the relative position of

Returns

1 if check is below key
0 if it is not below or if it is the same key
-1 on null pointer

See Also

[keyIsBelow\(\)](#), [keySetName\(\)](#), [keyGetName\(\)](#)

6.7.2.7 int keyIsInactive (const Key * key)

Check whether a key is inactive or not.

In elektra terminology any key is inactive if the it's basename starts with '.'. Inactive keys must not have any meaning to applications, they are reserved for users and administrators.

To remove a whole hierarchy in elektra, don't forget to pass option_t::KDB_O_INACTIVE to [kdbGet\(\)](#) to receive the inactive keys in order to remove them.

Otherwise you should not fetch these keys.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

1 if the key is inactive, 0 otherwise
-1 on NULL pointer or when key has no name

6.7.2.8 int keyIsString (const Key * key)

Check if a key is string type.

String values are null terminated and are not allowed to have any '\0' characters inside the string.

Make sure to use this function and don't test the string type another way to ensure compatibility and to write less error prone programs.

Returns

1 if it is string
0 if it is not
-1 on NULL pointer

See Also

[keyGetString\(\)](#), [keySetString\(\)](#)

Parameters

<i>key</i>	the key to check
------------	------------------

6.7.2.9 int `keyIsSystem` (const Key * *key*)

Check whether a key is under the `system` namespace or not

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

1 if key name begins with `system`, 0 otherwise
-1 on NULL pointer

See Also

[keyIsUser\(\)](#), [keySetName\(\)](#), [keyName\(\)](#)

6.7.2.10 int `keyIsUser` (const Key * *key*)

Check whether a key is under the `user` namespace or not.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

1 if key name begins with `user`, 0 otherwise
-1 on NULL pointer

See Also

[keyIsSystem\(\)](#), [keySetName\(\)](#), [keyName\(\)](#)

6.7.2.11 int `keyNeedSync` (const Key * *key*)

Test if a key needs to be synced to backend storage.

If any key modification took place the key will be flagged with `KEY_FLAG_SYNC` so that [kdbSet\(\)](#) knows which keys were modified and which not.

After [keyNew\(\)](#) the flag will normally be set, but after [kdbGet\(\)](#) and [kdbSet\(\)](#) the flag will be removed. When you modify the key the flag will be set again.

In your application you can make use of that flag to know if you changed something in a key after a [kdbGet\(\)](#) or [kdbSet\(\)](#).

Note

Note that also changes in the meta data will set that flag.

See Also

[keyNew\(\)](#)

Parameters

<code>key</code>	the key object to work with
------------------	-----------------------------

Returns

- 1 if `key` was changed in memory, 0 otherwise
- 1 on NULL pointer

6.7.2.12 int keyRel (const Key * *key*, const Key * *check*)

Information about the relation in the hierarchy between two keys.

Unlike `keyCmp()` the number gives information about hierarchical information.

- If the keys are the same 0 is returned. So it is the key itself.

```
user/key
user/key
```

```
keySetName (key, "user/key/folder");
keySetName (check, "user/key/folder");
succeed_if (keyRel (key, check) == 0, "should be same");
*
```

Note

this relation can be checked with `keyCmp()` too.

- If the key is direct below the other one 1 is returned. That means that, in terms of hierarchy, no other key is between them - it is a direct child.

```
user/key/folder
user/key/folder/child
```

```
keySetName (key, "user/key/folder");
keySetName (check, "user/key/folder/child");
succeed_if (keyRel (key, check) == 1, "should be direct below");
*
```

- If the key is below the other one, but not directly 2 is returned. This is also called grand-child.

```
user/key/folder
user/key/folder/any/depth/deeper/grand-child
```

```
keySetName (key, "user/key/folder");
keySetName (check, "user/key/folder/any/depth/deeper/grand-child");
succeed_if (keyRel (key, check) >= 2, "should be below (but not direct)");
;
succeed_if (keyRel (key, check) > 0, "should be below");
succeed_if (keyRel (key, check) >= 0, "should be the same or below");
*
```

- If a invalid or null ptr key is passed, -1 is returned
- If the keys have no relations, but are not invalid, -2 is returned.
- If the keys are in the same hierarchy, a value smaller then -2 is returned. It means that the key is not below.

```
user/key/myself
user/key/sibling
```

```
keySetName (key, "user/key/folder");
keySetName (check, "user/notsame/folder");
succeed_if (keyRel (key, check) < -2, "key is not below, but same
namespace");
```

TODO Below is an idea how it could be extended: It could continue the search into the other direction if any (grand-)parents are equal.

- If the keys are direct below a key which is next to the key, -2 is returned. This is also called nephew. (TODO not implemented)

```
user/key/myself
user/key/sibling
```

- If the keys are direct below a key which is next to the key, -2 is returned. This is also called nephew. (TODO not implemented)

```
user/key/myself
user/key/sibling/nephew
```

- If the keys are below a key which is next to the key, -3 is returned. This is also called grand-nephew. (TODO not implemented)

```
user/key/myself
user/key/sibling/any/deeper/deeper/grand-nephew
```

The same holds true for the other direction, but with negative values. For no relation INT_MIN is returned.

Note

to check if the keys are the same, you must use `keyCmp() == 0`! `keyRel()` does not give you the information if it did not find a relation or if it is the same key.

Returns

dependend on the relation: 2.. if below 1.. if direct below 0.. if the same -1.. on null or invalid keys -2.. if none of any other relation -3.. if same hierarchy (none of those below) -4.. if sibling (in same hierarchy) -5.. if nephew (in same hierarchy)


Parameters

<i>key</i>	the key object to work with
<i>check</i>	the second key object to check the relation with

6.8 Value Manipulation Methods

Methods to do various operations on Key values.

Collaboration diagram for Value Manipulation Methods:



Functions

- `const void * keyValue (const Key *key)`
- `const char * keyString (const Key *key)`
- `ssize_t keyGetValueSize (const Key *key)`
- `ssize_t keyGetString (const Key *key, char *returnedString, size_t maxSize)`
- `ssize_t keySetString (Key *key, const char *newStringValue)`
- `ssize_t keyGetBinary (const Key *key, void *returnedBinary, size_t maxSize)`
- `ssize_t keySetBinary (Key *key, const void *newBinary, size_t dataSize)`
- `const char * keyComment (const Key *key)`
- `ssize_t keyGetCommentSize (const Key *key)`
- `ssize_t keyGetComment (const Key *key, char *returnedComment, size_t maxSize)`
- `ssize_t keySetComment (Key *key, const char *newComment)`

6.8.1 Detailed Description

Methods to do various operations on Key values. A key can contain a value in different format. The most likely situation is, that the value is interpreted as text. Use `keyGetString()` for that. You can save any Unicode Symbols and Elektra will take care that you get the same back, independent of your current environment.

In some situations this idea fails. When you need exactly the same value back without any interpretation of the characters, there is `keySetBinary()`. If you use that, its very likely that your Configuration is not according to the standard. Also for Numbers, Booleans and Date you should use `keyGetString()`. To do so, you might use `strtod()` `strtoll()` and then `atol()` or `atof()` to convert back.

To use them:

```
#include <kdb.h>
```

6.8.2 Function Documentation

6.8.2.1 `const char* keyComment (const Key * key)`

Return a pointer to the real internal `key` comment.

This is a much more efficient version of `keyGetComment()` and you should use it if you are responsible enough to not mess up things. You are not allowed to change anything in the memory region the returned pointer points to.

`keyComment()` returns "" when there is no `keyComment`. The reason is

```
key=keyNew(0);
keySetComment(key, "");
keyComment(key); // you would expect "" here
keyDel(key);
```

See [keySetComment\(\)](#) for more information on comments.

Note

Note that the Key structure keeps its own size field that is calculated by library internal calls, so to avoid inconsistencies, you must never use the pointer returned by [keyComment\(\)](#) method to set a new value. Use [keySetComment\(\)](#) instead.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

a pointer to the internal managed comment
 "" when there is no comment
 0 on NULL pointer

See Also

[keyGetCommentSize\(\)](#) for size and [keyGetComment\(\)](#) as alternative

6.8.2.2 ssize_t keyGetBinary (const Key * key, void * returnedBinary, size_t maxSize)

Get the value of a key as a binary.

If the type is not binary -1 will be returned.

When the binary data is empty (this is not the same as "") 0 will be returned and the returnedBinary will not be changed.

For string values see [keyGetString\(\)](#) and [keyIsString\(\)](#).

When the returnedBinary is too small to hold the data (its maximum size is given by maxSize), the returnedBinary will not be changed and -1 is returned.

Example:

```
Key *key = keyNew ("user/keyname", KEY_TYPE, KEY_TYPE_BINARY, KEY_END);
char buffer[300];

if (keyGetBinary(key,buffer,sizeof(buffer)) == -1)
{
    // handle error
}
```

Parameters

<i>key</i>	the object to gather the value from
<i>returnedBinary</i>	pre-allocated memory to store a copy of the key value
<i>maxSize</i>	number of bytes of pre-allocated memory in returnedBinary

Returns

the number of bytes actually copied to `returnedBinary`

Return values

0	if the binary is empty
-1	on NULL pointers
-1	if maxSize is 0
-1	if maxSize is too small for string
-1	if maxSize is larger than SSIZE_MAX
-1	on type mismatch: binary expected, but found string

See Also

[keyValue\(\)](#), [keyGetValueSize\(\)](#), [keySetBinary\(\)](#)
[keyGetString\(\)](#) and [keySetString\(\)](#) as preferred alternative to binary
[keyIsBinary\(\)](#) to see how to check for binary type

6.8.2.3 ssize_t keyGetComment (const Key * key, char * returnedComment, size_t maxSize)

Get the key comment.

6.8.3 Comments

A Key comment is description for humans what this key is for. It may be a textual explanation of valid values, when and why a user or administrator changed the key or any other text that helps the user or administrator related to that key.

Don't depend on a comment in your program. A user is always allowed to remove or change it in any way he wants to. But you are allowed or even encouraged to always show the content of the comment to the user and allow him to change it.

Parameters

<i>key</i>	the key object to work with
<i>returned-Comment</i>	pre-allocated memory to copy the comments to
<i>maxSize</i>	number of bytes that will fit returnedComment

Returns

the number of bytes actually copied to `returnedString`, including final NULL
 1 if the string is empty
 -1 on NULL pointer
 -1 if maxSize is 0, not enough to store the comment or when larger then SSIZE_MAX

See Also

[keyGetCommentSize\(\)](#), [keySetComment\(\)](#)

6.8.3.1 ssize_t keyGetCommentSize (const Key * key)

Calculates number of bytes needed to store a key comment, including final NULL.

Use this method to know to size for allocated memory to retrieve a key comment.

See [keySetComment\(\)](#) for more information on comments.

For an empty key name you need one byte to store the ending NULL. For that reason 1 is returned.

```
char *buffer;
buffer = malloc (keyGetCommentSize (key));
// use this buffer to store the comment
// pass keyGetCommentSize (key) for maxSize
```

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

number of bytes needed
 1 if there is no comment
 -1 on NULL pointer

See Also

[keyGetComment\(\)](#), [keySetComment\(\)](#)

6.8.3.2 ssize_t keyGetString (const Key * *key*, char * *returnedString*, size_t *maxSize*)

Get the value of a key as a string.

When there is no value inside the string, 1 will be returned and the returnedString will be empty "" to avoid programming errors that old strings are shown to the user.

For binary values see [keyGetBinary\(\)](#) and [keyIsBinary\(\)](#).

Example:

```
Key *key = keyNew ("user/keyname", KEY_END);
char buffer[300];

if (keyGetString(key,buffer,sizeof(buffer)) == -1)
{
    // handle error
} else {
    printf ("buffer: %s\n", buffer);
}
```

Parameters

<i>key</i>	the object to gather the value from
<i>returnedString</i>	pre-allocated memory to store a copy of the key value
<i>maxSize</i>	number of bytes of allocated memory in returnedString

Returns

the number of bytes actually copied to returnedString, including final NULL

Return values

1	if the string is empty
-1	on any NULL pointers
-1	on type mismatch: string expected, but found binary

-1	maxSize is 0
-1	if maxSize is too small for string
-1	if maxSize is larger than SSIZE_MAX

See Also

[keyValue\(\)](#), [keyGetValueSize\(\)](#), [keySetString\(\)](#), [keyString\(\)](#)
[keyGetBinary\(\)](#) for working with binary data

6.8.3.3 ssize_t keyGetValueSize (const Key * key)

Returns the number of bytes needed to store the key value, including the NULL terminator.

It returns the correct size, independent of the Key Type. If it is a binary there might be '\0' values in it.

For an empty string you need one byte to store the ending NULL. For that reason 1 is returned. This is not true for binary data, so there might be returned 0 too.

A binary key has no '\0' termination. String types have it, so to there length will be added 1 to have enough space to store it.

This method can be used with malloc() before [keyGetString\(\)](#) or [keyGetBinary\(\)](#) is called.

```
char *buffer;
buffer = malloc (keyGetValueSize (key));
// use this buffer to store the value (binary or string)
// pass keyGetValueSize (key) for maxSize
```

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

the number of bytes needed to store the key value
 1 when there is no data and type is not binary
 0 when there is no data and type is binary
 -1 on null pointer

See Also

[keyGetString\(\)](#), [keyGetBinary\(\)](#), [keyValue\(\)](#)

6.8.3.4 ssize_t keySetBinary (Key * key, const void * newBinary, size_t dataSize)

Set the value of a key as a binary.

A private copy of `newBinary` will allocated and saved inside `key`, so the parameter can be deallocated after the call.

Binary values might be encoded in another way then string values depending on the plugin. Typically character encodings should not take place on binary data. Consider using a string key instead.

When `newBinary` is a NULL pointer the binary will be freed and 0 will be returned.

Note

The meta data "binary" will be set to mark that the key is binary from now on. When the key is already binary the meta data won't be changed. This will only happen in the successful case, but not when -1 is returned.

Parameters

<i>key</i>	the object on which to set the value
<i>newBinary</i>	is a pointer to any binary data or NULL to free the previous set data
<i>dataSize</i>	number of bytes to copy from <i>newBinary</i>

Returns

the number of bytes actually copied to internal struct storage
 0 when the internal binary was freed and is now a null pointer
 -1 if key is a NULL pointer
 -1 when dataSize is 0 (but newBinary not NULL) or larger than SSIZE_MAX

See Also

[keyGetBinary\(\)](#)
[keyIsBinary\(\)](#) to check if the type is binary
[keyGetString\(\)](#) and [keySetString\(\)](#) as preferred alternative to binary

6.8.3.5 ssize_t keySetComment (Key * key, const char * newComment)

Set a comment for a key.

A key comment is like a configuration file comment. See [keySetComment\(\)](#) for more information.

Parameters

<i>key</i>	the key object to work with
<i>newComment</i>	the comment, that can be freed after this call.

Returns

the number of bytes actually saved including final NULL
 0 when the comment was freed (newComment NULL or empty string)
 -1 on NULL pointer or memory problems

See Also

[keyGetComment\(\)](#)

6.8.3.6 ssize_t keySetString (Key * key, const char * newStringValue)

Set the value for key as newStringValue.

The function will allocate and save a private copy of newStringValue, so the parameter can be freed after the call.

String values will be saved in backend storage, when kdbSetKey() will be called, in UTF-8 universal encoding, regardless of the program's current encoding, when iconv plugin is present.

Note

The type will be set to KEY_TYPE_STRING. When the type of the key is already a string type it won't be changed.

Parameters

<i>key</i>	the key to set the string value
<i>newStringValue</i>	NULL-terminated text string to be set as <i>key</i> 's value

Returns

the number of bytes actually saved in private struct including final NULL

Return values

1	if newStringValue is a NULL pointer, this will make the string empty (string only containing null termination)
-1	if key is a NULL pointer

See Also

[keyGetString\(\)](#), [keyValue\(\)](#), [keyString\(\)](#)

6.8.3.7 const char* keyString (const Key * key)

Get the c-string representing the value.

Will return (null) on null pointers. Will return (binary) on binary data not ended with a null byte.

It is not checked if it is actually a string, only if it terminates for security reasons.

Returns

the c-string of the value

Return values

(null)	on null keys
""	if no data found
(binary)	on binary keys

Parameters

<i>key</i>	the key object to get the string from
------------	---------------------------------------

6.8.3.8 const void* keyValue (const Key * key)

Return a pointer to the real internal *key* value.

This is a much more efficient version of [keyGetString\(\)](#) [keyGetBinary\(\)](#), and you should use it if you are responsible enough to not mess up things. You are not allowed to modify anything in the returned string. If you need a copy of the Value, consider to use [keyGetString\(\)](#) or [keyGetBinary\(\)](#) instead.

6.8.4 String Handling

If `key` is string (`keyIsString()`), you may cast the returned as a `"char *"` because you'll get a NULL terminated regular string.

`keyValue()` returns "" in string mode when there is no value. The reason is

```
key=keyNew(0);
keySetString(key, "");
keyValue(key); // you would expect "" here
keyDel(key);
```

6.8.5 Binary Data Handling

If the data is binary, the size of the value must be determined by `keyGetValueSize()`, any `strlen()` operations are not suitable to determine the size.

`keyValue()` returns 0 in binary mode when there is no value. The reason is

```
key=keyNew(0);
keySetBinary(key, 0, 0);
keyValue(key); // you would expect 0 here

keySetBinary(key, "", 1);
keyValue(key); // you would expect "" (a pointer to '\0') here

int i=23;
keySetBinary(key, (void*)&i, 4);
(int*)keyValue(key); // you would expect a pointer to (int)23 here
keyDel(key);
```

Note

Note that the Key structure keeps its own size field that is calculated by library internal calls, so to avoid inconsistencies, you must never use the pointer returned by `keyValue()` method to set a new value. Use `keySetString()` or `keySetBinary()` instead.

Warning

Binary keys will return a NULL pointer when there is no data in contrast to `keyName()`, `keyBaseName()`, `keyOwner()` and `keyComment()`. For string value the behaviour is the same.

Example:

```
KDB *handle = kdbOpen();
KeySet *ks=ksNew(0);
Key *current=0;

kdbGetByName(handle, ks, "system/sw/my", KDB_O_SORT|KDB_O_RECURSIVE);

ksRewind(ks);
while(current=ksNext(ks)) {
    size_t size=0;

    if (keyIsBin(current)) {
        size=keyGetValueSize(current);
        printf("Key %s has a value of size %d bytes. Value: <BINARY>\n",
            Comment: %s",
                keyName(current),
                size,
                keyComment(current));
    } else {
        size=elektraStrLen((char *)keyValue(current));
        printf("Key %s has a value of size %d bytes. Value: %s\n",
            Comment: %s",
                keyName(current),
                size,
                (char *)keyValue(current),
                keyComment(current));
    }
}

ksDel(ks);
kdbClose(handle);
```

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

a pointer to internal value
"" when there is no data and key is not binary
0 where there is no data and key is binary
0 on NULL pointer

See Also

[keyGetValueSize\(\)](#), [keyGetString\(\)](#), [keyGetBinary\(\)](#)

6.9 Plugins

Elektra plugin framework.

Functions

- int [docOpen](#) (Plugin *handle, Key *errorKey)
- int [docClose](#) (Plugin *handle, Key *errorKey)
- int [docGet](#) (Plugin *handle, KeySet *returned, Key *parentKey)
- int [docSet](#) (Plugin *handle, KeySet *returned, Key *parentKey)
- int [docError](#) (Plugin *handle, KeySet *returned, Key *parentKey)
- Plugin * [ELEKTRA_PLUGIN_EXPORT](#) (doc)

6.9.1 Detailed Description

Elektra plugin framework.

Since

version 0.4.9, Elektra can dynamically load different key storage plugins.
 version 0.7.0 Elektra can have multiple backends, mounted at any place in the key database.
 version 0.8.0 Elektra backends are composed out of multiple plugins.

Overview

There are different types of plugins for different concerns. The types of plugins handled in this document:

- file storage plugins (also called just storage plugins here)
- filter plugins
 See <http://www.libelektra.org/ftp/elektra/thesis.pdf> for an detailed explanation and description of other types of plugins.
 A plugin can implement anything related to configuration. There are 5 possible entry points, as described in this document:
 - elektraDocOpen()
 - elektraDocClose()
 - elektraDocGet()
 - elektraDocSet()
 - elektraDocError() (not needed by storage or filter plugins)
 Depending of the type of plugin you need not to implement all of them.

Note

that the Doc within the name is just because the plugin described here is called doc (see `src/plugins/doc/doc.c`). Always replace Doc with the name of the plugin you are going to implement.

See the descriptions below what each of them is supposed to do.

Storage Plugins

A filter plugin is a plugin which already receives some keys. It may process or change the keyset. Or it may reject specific keysets which do not meet some criteria.

Filter Plugins

A storage plugin gets an empty keyset and constructs the information out from a file.
 Other persistent storage then a file is not handled within this document because it involves many other issues.
 For files the resolver plugin already takes care for transactions and rollback.

Error and Warnings

In any case of trouble, use ELEKTRA_SET_ERROR and return with -1. You might add warnings with ELEKTRA_ADD_WARNING if you think it is appropriate.

Note

some docu in this section might be confusing or not updated, please refer to <http://www.libelektra.org/ftp/elektra/thesis.pdf> or ask at the mailinglist if something is unclear.

6.9.2 Function Documentation

6.9.2.1 int docClose (Plugin * *handle*, Key * *errorKey*)

Finalize the plugin.

Called prior to unloading the plugin dynamic module. Should ensure that no functions or static/global variables from the module will ever be accessed again.

Make sure to free all memory that your plugin requested at runtime.

After this call, libelektra.so will unload the plugin library, so this is the point to shutdown any affairs with the storage.

Parameters

<i>handle</i>	contains internal information of the plugin
<i>errorKey</i>	is needed to add warnings using ELEKTRA_ADD_WARNING

Return values

0	on success
---	------------

See Also

[kdbClose\(\)](#)

[elektraPluginGetData\(\)](#), [elektraPluginSetData\(\)](#) and [elektraPluginGetConfig\(\)](#)

6.9.2.2 int docError (Plugin * *handle*, KeySet * *returned*, Key * *parentKey*)

Rollback in case of errors.

Parameters

<i>handle</i>	contains internal information of the plugin
<i>returned</i>	contains a keyset with relevant keys
<i>parentKey</i>	contains the information where to set the keys

Return values

1	on success
0	on success with no action
-1	on failure

6.9.2.3 int docGet (Plugin * *handle*, KeySet * *returned*, Key * *parentKey*)

Retrieve information from a permanent storage to construct a keyset.

6.9.3 Introduction

The `elektraDocGet()` function handle everything related to receiving keys.

6.9.3.1 Storage Plugins

For storage plugins the filename is written in the value of the `parentKey`. So the first task of the plugin is to open that file. Then it should parse its content and construct a keyset with all information of that file.

You need to be able to reconstruct the same file with the information of the keyset. So be sure to copy all comments, whitespaces and so on into some metadata of the keys. Otherwise the information is lost after writing the file the next time.

Now lets look at an example how the typical `elektraDocGet()` might be implemented. To explain we introduce some pseudo functions which do all the work with the storage (which is of course 90% of the work for a real plugin):

- `parse_key` will parse a key and a value from an open file handle

The typical loop for a storage plugin will be like:

```
int elektraDocGet(Plugin *handle, KeySet *returned, const Key *parentKey)
{
    // contract handling, see below

    FILE *fp = fopen (keyString(parentKey), "r");
    char *key;
    char *value;

    while ((n = parse_key(fp, &key, &value)) >= 1)
    {
        Key *read = keyNew(0);
        if (keySetName(read, key) == -1)
        {
            fclose (fp);
            keyDel (read);
            ELEKTRA_SET_ERROR(59, parentKey, key);
            return -1;
        }
        keySetString(read, value);

        ksAppendKey (returned, read);
        free (key);
        free (value);
    }

    if (feof(fp) == 0)
    {
        fclose (fp);
        ELEKTRA_SET_ERROR(60, parentKey, "not at the end of file");
        return -1;
    }

    fclose (fp);

    return 1; // success
}
```

6.9.3.2 Filter Plugins

For filter plugins the actual task is rather unspecified. You basically can do anything with the keyset. To get roundtrip properties you might want to undo any changes you did in `elektraDocSet()`.

The pseudo functions (which do the real work) are:

- `do_action()` which processes every key in this filter

```
int elektraDocGet(Plugin *handle, KeySet *returned, Key *parentKey)
{
    // contract handling

    Key *k;
```



```

    ksRewind (returned);
    while ((k = ksNext (returned)) != 0)
    {
        do_action(k);
    }

    return 1; // success
}

```

6.9.4 Conditions

Precondition

The caller `kdbGet()` will make sure before you are called that the `parentKey`:

- is a valid key (means that it is a system or user key).
- is below (see `keysBelow()`) your mountpoint and that your plugin is responsible for it. and that the returned:
- is a valid keyset.
- has all keys with the flag `KEY_FLAG_SYNC` set.
- contains only valid keys direct below (see `keysDirectBelow()`) your `parentKey`. That also means, that the `parentKey` will not be in that keyset.
- is in a sorted order, see `ksSort()`. and that the handle:
 - is a valid KDB for your plugin.
 - that `elektraPluginGetBackendHandle()` contains the same handle for lifetime `kdbOpen()` until `elektraPluginClose()` was called.

The caller `kdbGet()` will make sure that afterwards you were called, whenever the user requested it with the options, that:

- hidden keys they will be thrown away.
- dirs or only dirs `kdbGet()` will remove the other.
- you will be called again recursively with all subdirectories.
- the keyset will be sorted when needed.
- the keys in returned having `KEY_FLAG_SYNC` will be sorted out.

Invariant

There are no global variables and `elektraPluginSetData()` stores all information. The handle is to be guaranteed to be the same if it is the same plugin.

Postcondition

The keyset returned has the `parentKey` and all keys direct below (`keysDirectBelow()`) with all information from the storage. Make sure to return all keys, all directories and also all hidden keys. If some of them are not wished, the caller `kdbGet()` will drop these keys, see above.

6.9.5 Updating

To get all keys out of the storage over and over again can be very inefficient. You might know a more efficient method to know if the key needs update or not, e.g. by stating it or by an external time stamp info. For file storage plugins this is automatically done for you. For other types (e.g. databases) you need to implement your own resolver doing this.

Note

If any calls you use `change_errno`, make sure to restore the old `errno`.

See Also

[kdbGet\(\)](#) for caller.

Parameters

<i>handle</i>	contains internal information of opened key database
<i>returned</i>	contains a keyset where the function need to append the keys got from the storage. There might be also some keys inside it, see conditions. You may use them to support efficient updating of keys, see Updating .
<i>parentKey</i>	contains the information below which key the keys should be gotten.

Returns

1 on success
 0 when nothing was to do
 -1 on failure, the current key in returned shows the position. use ELEKTRA_SET_ERROR of kdberrors.h to define the error code

6.9.5.1 int docOpen (Plugin * *handle*, Key * *errorKey*)

Initialize the plugin.

This is the first method called after dynamically loading this plugin.

This method is responsible for:

- plugin's specific configuration gathering
- all plugin's internal structs initialization
- if unavoidable initial setup of all I/O details such as opening a file, connecting to a database, setup connection to a server, etc.

You may also read the configuration you can get with `elektraPluginGetConfig()` and transform it into other structures used by your plugin.

Note

The plugin must not have any global variables. If you do Elektra will not be threadsafe. It is not a good assumption that your plugin will be opened only once.

Instead you can use `elektraPluginGetData()` and `elektraPluginSetData()` to store and get any information related to your plugin.

The correct substitute for global variables will be:

```
struct _GlobalData{ int global; };
typedef struct _GlobalData GlobalData;
int elektraDocOpen(Plugin *handle, Key *errorKey)
{
    GlobalData *data;
    data=malloc(sizeof(GlobalData));
    data.global = 20;
    elektraPluginSetData(handle,data);
}
```

Note

Make sure to free everything you allocate here within `elektraDocClose()`.

Returns

0 on success

Parameters

<i>handle</i>	contains internal information of the plugin
<i>errorKey</i>	defines an errorKey

See Also

[kdbOpen\(\)](#) which will call [elektraDocOpen\(\)](#)
[elektraPluginGetData\(\)](#), [elektraPluginSetData\(\)](#) and [elektraPluginGetConfig\(\)](#)

6.9.5.2 int docSet (Plugin * handle, KeySet * returned, Key * parentKey)

Store a keyset permanently.

This function does everything related to set and remove keys in a plugin. There is only one function for that purpose to make implementation and locking much easier.

The keyset `returned` was filled in with information from the application using `elektra` and the task of this function is to store it in a permanent way so that a subsequent call of `elektraPluginGet()` can rebuild the keyset as it was before. See the live cycle of a comment to understand:

```
void usercode (Key *key)
{
    keySetComment (key, "mycomment"); // the usercode stores a
    comment for the key
    ksAppendKey (keyset, key); // append the key to the keyset
    kdbSet (handle, keyset, 0, 0);
}

// so now kdbSet is called
int kdbSet (KDB *handle, KeySet *keyset, Key *parentKey, options)
{
    // find appropriate plugin
    elektraPluginSet (handle, keyset, 0); // the keyset with the key will
    be passed to this function
}

// so now elektraPluginSet(), which is the function described here, is called
elektraPluginSet (KDB *handle, KeySet *keyset, Key *parentKey)
{
    // the task of elektraPluginSet is now to store the comment
    Key *key = ksCurrent (keyset); // get out the key where the
    user set the comment before
    char *comment = allocate (size);
    keyGetComment (key, comment, size);
    savetodisc (comment);
}
```

Of course not only the comment, but all information of every key in the keyset `returned` need to be stored permanently. So this specification needs to give an exhaustive list of information present in a key.

Precondition

The keyset `returned` holds all keys which must be saved permanently for this keyset. The keyset is sorted and rewinded. All keys having children must be true for [keyIsDir\(\)](#).

The `parentKey` is the key which is the ancestor for all other keys in the keyset. The first key of the keyset `returned` has the same keyname. The `parentKey` is below the mountpoint, see `kdbhGetMountpoint()`.

The caller `kdbSet` will fulfill following parts:

- If the user does not want hidden keys they will be thrown away. All keys in `returned` need to be stored permanently.
- If the user does not want dirs or only dirs [kdbGet\(\)](#) will remove the other.

- Sorting of the keyset. It is not important in which order the keys are appended. So make sure to set all keys, all directories and also all hidden keys. If some of them are not wished, the caller `kdbSet()` will sort them out.

Invariant

There are no global variables and `kdbhGetBackendData()` only stores information which can be regenerated any time. The handle is the same when it is the same plugin.

Postcondition

The information of the keyset `returned` is stored permanently.

Lock your permanent storage in an exclusive way, no access of a concurrent `elektraPluginSet_plugin()` or `kdbGet()` is possible and these methods block until the function has finished. Otherwise declare `kdbcGetnoLock()`.

See Also

`kdbSet()` for caller.

Parameters

<i>handle</i>	contains internal information of the plugin
<i>returned</i>	contains a keyset with relevant keys
<i>parentKey</i>	contains the information where to set the keys

Returns

When everything works gracefully return the number of keys you set. The cursor position and the keys remaining in the keyset are not important.

Note

If any calls you use `change_errno`, make sure to restore the old `errno`.

Return values

<i>1</i>	on success
<i>0</i>	on success with no changed key in database
<i>-1</i>	on failure. The cause of the error needs to be added in <code>parentKey</code>

You also have to make sure that `ksGetCursor()` shows to the position where the error appeared.

6.9.5.3 Plugin* ELEKTRA_PLUGIN_EXPORT (doc)

All KDB methods implemented by the plugin can have random names, except `ELEKTRA_PLUGIN_EXPORT`. This is the single symbol that will be looked up when loading the plugin, and the first method of the backend implementation that will be called.

You need to use a macro so that both dynamic and static loading of the plugin works.

The first parameter is the name of the plugin. Then every plugin should have: `ELEKTRA_PLUGIN_OPEN`, `ELEKTRA_PLUGIN_CLOSE`, `ELEKTRA_PLUGIN_GET`, `ELEKTRA_PLUGIN_SET` and optionally `ELEKTRA_PLUGIN_ERROR`.

The list is terminated with `ELEKTRA_PLUGIN_END`.

You must use static "char arrays" in a read only segment. Don't allocate storage, it won't be freed.

Returns

Plugin

See Also

elektraPluginExport()

Chapter 7

Namespace Documentation

7.1 kdb Namespace Reference

See examples/cpp_example_userexception.cpp for how to use USER_DEFINED_EXEPTIONS.

Data Structures

- class [KDB](#)
Opens the session with the [Key](#) database.
- class [Key](#)
A Key is the essential class that encapsulates key [name](#) , [value](#) and [metainfo](#) .
- class [KeySet](#)
A keyset holds together a set of keys.

Functions

- `std::ostream & operator<< (std::ostream &os, kdb::Key const &k)`
Stream a the name of a key.
- `std::istream & operator>> (std::istream &is, kdb::Key &k)`
Reads a line with a keys name.
- `std::ostream & operator<< (std::ostream &os, kdb::KeySet const &ks)`
Outputs line per line the keynames.
- `std::istream & operator>> (std::istream &is, kdb::KeySet &ks)`
Reads line per line key names and appends those keys to ks.

7.1.1 Detailed Description

See examples/cpp_example_userexception.cpp for how to use USER_DEFINED_EXEPTIONS. See examples/cpp-example_userio.cpp for how to use USER_DEFINED_IO.

7.1.2 Function Documentation

7.1.2.1 `std::ostream& kdb::operator<< (std::ostream & os, kdb::Key const & k)` `[inline]`

Stream a the name of a key.

If you also want to stream the value, use the plugin framework.

Parameters

<i>os</i>	the stream to write to
<i>k</i>	the key which name should be streamed

Returns

the stream

7.1.2.2 `std::ostream& kdb::operator<< (std::ostream & os, kdb::KeySet const & cks)` `[inline]`

Outputs line per line the keynames.

To output values you should use the plugin framework.

Parameters

<i>os</i>	the stream to write to
<i>cks</i>	the keyset which should be streamed

Returns

the stream

7.1.2.3 `std::istream& kdb::operator>> (std::istream & is, kdb::Key & k)` `[inline]`

Reads a line with a keys name.

Parameters

<i>is</i>	the stream to read from
<i>k</i>	the key whose name will be set

Returns

the stream

7.1.2.4 `std::istream& kdb::operator>> (std::istream & is, kdb::KeySet & ks)` `[inline]`

Reads line per line key names and appends those keys to ks.

To input values you need to use the plugin framework.

Parameters

<i>is</i>	the stream to read from
<i>ks</i>	the keyset to append to

Returns

the stream

Chapter 8

Data Structure Documentation

8.1 kdb::KDB Class Reference

Opens the session with the [Key](#) database.

```
#include <kdb.hpp>
```

Public Member Functions

- [KDB](#) ()
Opens the session with the [Key](#) database.
- [KDB](#) ([Key](#) &errorKey)
Opens the session with the [Key](#) database.
- [~KDB](#) () throw ()
- void [open](#) ([Key](#) &errorKey)
Opens the session with the [Key](#) database.
- void [close](#) ([Key](#) &errorKey) throw ()
- int [get](#) ([KeySet](#) &returned, std::string const &keyname)
Retrieve keys in an atomic and universal way, all other [kdbGet\(\)](#) Functions rely on that one.
- int [get](#) ([KeySet](#) &returned, [Key](#) &parentKey)
Retrieve keys in an atomic and universal way, all other [kdbGet\(\)](#) Functions rely on that one.
- int [set](#) ([KeySet](#) &returned, std::string const &keyname)
- int [set](#) ([KeySet](#) &returned, [Key](#) &parentKey)

8.1.1 Detailed Description

Opens the session with the [Key](#) database.

Constructs a class [KDB](#).

Exceptions

KDBException	if database could not be opened
------------------------------	---------------------------------

The first step is to open the default backend. With it system/elektra/mountpoints will be loaded and all needed

libraries and mountpoints will be determined. These libraries for backends will be loaded and with it the [KDB](#) datastructure will be initialized.

You must always call this method before retrieving or committing any keys to the database. In the end of the program, after using the key database, you must not forget to [kdbClose\(\)](#). You can use the `atexit()` handler for it.

The pointer to the [KDB](#) structure returned will be initialized like described above, and it must be passed along on any `kdb*()` method your application calls.

Get a [KDB](#) handle for every thread using `elektra`. Don't share the handle across threads, and also not the pointer accessing it:

```
thread1
{
    KDB * h;
    h = kdbOpen(0);
    // fetch keys and work with them
    kdbClose(h, 0);
}
thread2
{
    KDB * h;
    h = kdbOpen(0);
    // fetch keys and work with them
    kdbClose(h, 0);
}
```

You don't need to use the [kdbOpen\(\)](#) if you only want to manipulate plain in-memory [Key](#) or [KeySet](#) objects without any affairs with the backend key database or when your application loads plugins directly.

Parameters

<i>errorKey</i>	the key which holds errors and warnings which were issued must be given
-----------------	---

See Also

[kdbClose\(\)](#) to end all affairs to the [Key](#) database.

Returns

a [KDB](#) pointer on success
NULL on failure

Access to the key database.

Invariant

the object holds an valid connection to the key database or is empty

8.1.2 Constructor & Destructor Documentation

8.1.2.1 `kdb::KDB::KDB()` `[inline]`

Opens the session with the [Key](#) database.

Constructs a class [KDB](#).

Exceptions

<i>KDBException</i>	if database could not be opened
---------------------	---------------------------------

The first step is to open the default backend. With it system/elektra/mountpoints will be loaded and all needed libraries and mountpoints will be determined. These libraries for backends will be loaded and with it the [KDB](#) datastructure will be initialized.

You must always call this method before retrieving or committing any keys to the database. In the end of the program, after using the key database, you must not forget to `kdbClose()`. You can use the `atexit ()` handler for it.

The pointer to the `KDB` structure returned will be initialized like described above, and it must be passed along on any `kdb*()` method your application calls.

Get a `KDB` handle for every thread using `elektra`. Don't share the handle across threads, and also not the pointer accessing it:

```
thread1
{
    KDB * h;
    h = kdbOpen(0);
    // fetch keys and work with them
    kdbClose(h, 0);
}
thread2
{
    KDB * h;
    h = kdbOpen(0);
    // fetch keys and work with them
    kdbClose(h, 0);
}
```

You don't need to use the `kdbOpen()` if you only want to manipulate plain in-memory `Key` or `KeySet` objects without any affairs with the backend key database or when your application loads plugins directly.

Parameters

<code>errorKey</code>	the key which holds errors and warnings which were issued must be given
-----------------------	---

See Also

`kdbClose()` to end all affairs to the `Key` database.

Returns

a `KDB` pointer on success
 NULL on failure

8.1.2.2 kdb::KDB::KDB (Key & errorKey) [inline]

Opens the session with the `Key` database.

Constructs a class `KDB`.

Parameters

<code>errorKey</code>	is useful if you want to get the warnings in the successful case, when no exception is thrown.
-----------------------	--

Exceptions

<code>KDBException</code>	if database could not be opened
---------------------------	---------------------------------

The first step is to open the default backend. With it `system/elektra/mountpoints` will be loaded and all needed libraries and mountpoints will be determined. These libraries for backends will be loaded and with it the `KDB` datastructure will be initialized.

You must always call this method before retrieving or committing any keys to the database. In the end of the program, after using the key database, you must not forget to `kdbClose()`. You can use the `atexit ()` handler for it.

The pointer to the `KDB` structure returned will be initialized like described above, and it must be passed along on any `kdb*()` method your application calls.

Get a [KDB](#) handle for every thread using `elektra`. Don't share the handle across threads, and also not the pointer accessing it:

```
thread1
{
    KDB * h;
    h = kdbOpen(0);
    // fetch keys and work with them
    kdbClose(h, 0);
}
thread2
{
    KDB * h;
    h = kdbOpen(0);
    // fetch keys and work with them
    kdbClose(h, 0);
}
```

You don't need to use the `kdbOpen()` if you only want to manipulate plain in-memory [Key](#) or [KeySet](#) objects without any affairs with the backend key database or when your application loads plugins directly.

Parameters

<i>errorKey</i>	the key which holds errors and warnings which were issued must be given
-----------------	---

See Also

[kdbClose\(\)](#) to end all affairs to the [Key](#) database.

Returns

a [KDB](#) pointer on success
NULL on failure

8.1.2.3 `kdb::~KDB::~~KDB () throw ()` `[inline]`

The destructor closes the database.

Closes the session with the [Key](#) database.

You should call this method when you finished your affairs with the key database. You can manipulate [Key](#) and [KeySet](#) objects also after `kdbClose()`. You must not use any `kdb*` call afterwards. You can implement `kdbClose()` in the `atexit()` handler.

This is the counterpart of `kdbOpen()`.

The `handle` parameter will be finalized and all resources associated to it will be freed. After a `kdbClose()`, this `handle` can't be used anymore, unless it gets initialized again with another call to `kdbOpen()`.

See Also

[kdbOpen\(\)](#)

Parameters

<i>handle</i>	contains internal information of opened key database
<i>errorKey</i>	the key which holds error information

Returns

0 on success
-1 on NULL pointer

8.1.3 Member Function Documentation

8.1.3.1 void kdb::KDB::close (Key & errorKey) throw () [inline]

Open the database.

The return value does not matter because its only a null pointer check.

Parameters

<i>errorKey</i>	is useful if you want to get the warnings
-----------------	---

Closes the session with the [Key](#) database.

You should call this method when you finished your affairs with the key database. You can manipulate [Key](#) and [KeySet](#) objects also after [kdbClose\(\)](#). You must not use any kdb* call afterwards. You can implement [kdbClose\(\)](#) in the `atexit()` handler.

This is the counterpart of [kdbOpen\(\)](#).

The `handle` parameter will be finalized and all resources associated to it will be freed. After a [kdbClose\(\)](#), this `handle` can't be used anymore, unless it gets initialized again with another call to [kdbOpen\(\)](#).

See Also

[kdbOpen\(\)](#)

Parameters

<i>handle</i>	contains internal information of opened key database
<i>errorKey</i>	the key which holds error information

Returns

0 on success
-1 on NULL pointer

8.1.3.2 int kdb::KDB::get (KeySet & returned, std::string const & keyname) [inline]

Retrieve keys in an atomic and universal way, all other [kdbGet\(\)](#) Functions rely on that one.

Get all keys below keyname inside returned.

The `returned` [KeySet](#) must be initialized. The `returned` [KeySet](#) may already contain some keys. The new retrieved keys will be appended using [ksAppendKey\(\)](#).

It will fully retrieve all keys under the `parentKey` folder, with all subfolders and their children.

```
#include <kdb.hpp>
#include <keyio.hpp>
using namespace kdb;

int main()
{
    KeySet config;
    KDB kdb;
    kdb.get(config, "/sw/MyApp");

    Key k = config.lookup("/sw/MyApp/mykey");
    if (k)
    {
        std::cout << k << " is " << k.get<int>() << std::endl;
    }
    else
```

```

    {
        std::cerr << "No key found" << std::endl;
        return 1;
    }
}

```

Parameters

<i>returned</i>	the keyset where the keys will be in
<i>keyname</i>	the root keyname which should be used to get keys below it

Return values

0	if no key was updated
1	if user or system keys were updated
2	if user and system keys were updated

Exceptions

<i>KDBException</i>	if there were problems with the database
---------------------	--

See Also

[KDB::get](#) ([KeySet](#) & returned, [Key](#) & parentKey)

8.1.3.3 int kdb::KDB::get (KeySet & returned, Key & parentKey) [inline]

Retrieve keys in an atomic and universal way, all other [kdbGet\(\)](#) Functions rely on that one.

Get all keys below parentKey inside returned.

The returned [KeySet](#) must be initialized. The returned [KeySet](#) may already contain some keys. The new retrieved keys will be appended using [ksAppendKey\(\)](#).

It will fully retrieve all keys under the parentKey folder, with all subfolders and their children.

Parameters

<i>returned</i>	the keyset where the keys will be in
<i>parentKey</i>	the parentKey of returned

Return values

0	if no key was updated
1	if user or system keys were updated
2	if user and system keys were updated

Exceptions

<i>KDBException</i>	if there were problems with the database
---------------------	--

8.1.3.4 void kdb::KDB::open (Key & errorKey) [inline]

Opens the session with the [Key](#) database.

Open the database

Parameters

<i>errorKey</i>	is useful if you want to get the warnings in the successful case, when no exception is thrown.
-----------------	--

The first step is to open the default backend. With it system/elektra/mountpoints will be loaded and all needed libraries and mountpoints will be determined. These libraries for backends will be loaded and with it the [KDB](#) datastructure will be initialized.

You must always call this method before retrieving or committing any keys to the database. In the end of the program, after using the key database, you must not forget to [kdbClose\(\)](#). You can use the `atexit()` handler for it.

The pointer to the [KDB](#) structure returned will be initialized like described above, and it must be passed along on any `kdb*()` method your application calls.

Get a [KDB](#) handle for every thread using `elektra`. Don't share the handle across threads, and also not the pointer accessing it:

```
thread1
{
    KDB * h;
    h = kdbOpen(0);
    // fetch keys and work with them
    kdbClose(h, 0);
}
thread2
{
    KDB * h;
    h = kdbOpen(0);
    // fetch keys and work with them
    kdbClose(h, 0);
}
```

You don't need to use the [kdbOpen\(\)](#) if you only want to manipulate plain in-memory [Key](#) or [KeySet](#) objects without any affairs with the backend key database or when your application loads plugins directly.

Parameters

<i>errorKey</i>	the key which holds errors and warnings which were issued must be given
-----------------	---

See Also

[kdbClose\(\)](#) to end all affairs to the [Key](#) database.

Returns

a [KDB](#) pointer on success
NULL on failure

8.1.3.5 int kdb::KDB::set (KeySet & returned, std::string const & keyname) [inline]

Set all keys below keyname.

If the keyname of the parentKey is invalid (e.g. empty) all keys will be set.

Set keys in an atomic and universal way.

All other `kdbSet` Functions rely on that one.

Return values

0	if no key was updated
1	if user or system keys were updated
2	if user and system keys were updated

Parameters

<i>returned</i>	the keyset where the keys will be in
<i>keyname</i>	the keyname below the names should be set

Exceptions

<i>KDBException</i>	if there were problems with the database
---------------------	--

8.1.3.6 `int kdb::KDB::set (KeySet & returned, Key & parentKey) [inline]`

Set all keys below parentKey.

If the keyname of the parentKey is invalid (e.g. empty) all keys will be set.

Set keys in an atomic and universal way.

All other kdbSet Functions rely on that one.

Return values

<i>0</i>	if no key was updated
<i>1</i>	if user or system keys were updated
<i>2</i>	if user and system keys were updated

Parameters

<i>returned</i>	the keyset where the keys are passed to the user
<i>parentKey</i>	the parentKey of returned

Exceptions

<i>KDBException</i>	if there were problems with the database
---------------------	--

The documentation for this class was generated from the following file:

- `src/bindings/cpp/include/kdb.hpp`

8.2 kdb::Key Class Reference

A Key is the essential class that encapsulates key `name` , `value` and `metainfo` .

```
#include <key.hpp>
```

Public Member Functions

- `Key ()`
- `Key (ckdb::Key *k)`
- `Key (Key &k)`
- `Key (Key const &k)`
- `Key (const char *keyName,...)`
- `Key (const std::string keyName,...)`
- `Key (const char *keyName, va_list ap)`

- void `operator++` (int) const
- void `operator++` () const
- void `operator--` (int) const
- void `operator--` () const
- size_t `getReferenceCounter` () const
- `Key` & `operator=` (ckdb::Key *k)
- `Key` & `operator=` (const `Key` &k)
- void `copy` (const `Key` &other)
- void `clear` ()
- ckdb::Key * `getKey` () const
- ckdb::Key * `operator*` () const
- ckdb::Key * `release` ()
- ckdb::Key * `dup` () const
- `~Key` ()
- std::string `getName` () const
- size_t `getNameSize` () const
- std::string `getBaseName` () const
- size_t `getBaseNameSize` () const
- std::string `getDirName` () const
- void `setName` (const std::string &newName)
- void `setBaseName` (const std::string &baseName)
- void `addBaseName` (const std::string &baseName)
- size_t `getFullNameSize` () const
- std::string `getFullName` () const
- `Key` & `operator=` (const std::string &newName)
- `Key` & `operator+=` (const std::string &baseName)
- `Key` & `operator=` (const std::string &baseName)
- `Key` & `operator=` (const char *newName)
(const std::string &newName)
- `Key` & `operator+=` (const char *baseName)
(const std::string &)

- `Key & operator-=` (const char *baseName)
 (const std::string &)
- bool `operator==` (const `Key` &k) const
- bool `operator!=` (const `Key` &k) const
- bool `operator<` (const `Key` &other) const
- bool `operator<=` (const `Key` &other) const
- bool `operator>` (const `Key` &other) const
- bool `operator>=` (const `Key` &other) const
- `operator bool` () const
- template<class T >
 T `get` () const
- template<class T >
 void `set` (T x)
- std::string `getString` () const
- void `setString` (std::string newString)
- size_t `getStringSize` () const
- func_t `getFunc` () const
- const void * `getValue` () const
- std::string `getBinary` () const
- size_t `getBinarySize` () const
- size_t `setBinary` (const void *newBinary, size_t dataSize)
- template<class T >
 T `getMeta` (const std::string &metaName) const
- template<class T >
 void `setMeta` (const std::string &metaName, T x)
- void `copyMeta` (const `Key` &other, const std::string &metaName)
- void `copyAllMeta` (const `Key` &other)
- void `rewindMeta` () const
- const `Key` `nextMeta` ()
- const `Key` `currentMeta` () const

- bool [isValid](#) () const
- bool [isSystem](#) () const
- bool [isUser](#) () const
- bool [isString](#) () const
- bool [isBinary](#) () const
- bool [isInactive](#) () const
- bool [isBelow](#) (const [Key](#) &k) const
- bool [isBelowOrSame](#) (const [Key](#) &k) const
- bool [isDirectBelow](#) (const [Key](#) &k) const

8.2.1 Detailed Description

A Key is the essential class that encapsulates key [name](#) , [value](#) and [metainfo](#) .

This class is an wrapper for an optional, refcounted `ckdb::Key`. It is like an `shared_ptr<ckdb::Key>`, but the `shared_ptr` functionality is already within the [Key](#) and exposed with this wrapper.

optional

A key can be constructed with an null pointer, by using [Key](#) (`static_cast<ckdb::Key*>(0)`); or made empty afterwards by using [release\(\)](#) or assign a null key. To check if there is an associated managed object the user can use operator `bool()`.

references

Copies of keys are cheap because they are only flat. If you really need a deep copy, you can use [copy\(\)](#) or [dup\(\)](#). If you [release\(\)](#) an object, the reference counter will stay All other operations operate on references.

documentation

Note that the documentation is typically copied from the underlying function which is wrapped and sometimes extended with C++ specific details. So you might find C examples within the C++ documentation.

Invariant

[Key](#) either has a working underlying Elektra [Key](#) object or a null pointer. The [Key](#), however, might be invalid (see [isValid\(\)](#)) or null (see operator `bool()`).

Note

that the reference counting in the keys is mutable, so that const keys can be passed around by value.

8.2.2 Constructor & Destructor Documentation

8.2.2.1 `kdb::Key::Key ()` `[inline]`

Constructs an empty, invalid key.

Note

That this is not a null key, so the key will evaluate to true.

See Also

[isValid\(\)](#), operator bool()

8.2.2.2 `kdb::Key::Key (cldb::Key * k) [inline]`

Constructs a key out of a C key.

Note

If you pass a null pointer here, the key will evaluate to false.

Parameters

<code>k</code>	the key to work with
----------------	----------------------

See Also

[isValid\(\)](#), operator bool()

8.2.2.3 `kdb::Key::Key (Key & k) [inline]`

Takes a reference of another key.

The key will not be copied, but the reference counter will be increased.

Parameters

<code>k</code>	the key to work with
----------------	----------------------

8.2.2.4 `kdb::Key::Key (Key const & k) [inline]`

Takes a reference of another key.

The key will not be copied, but the reference counter will be increased.

Parameters

<code>k</code>	the key to work with
----------------	----------------------

8.2.2.5 `kdb::Key::Key (const char * keyName, ...) [inline],[explicit]`

A practical way to fully create a [Key](#) object in one step.

This function tries to mimic the C++ way for constructors.

To just get a key object, simple do:

```
Key *k = keyNew(0);
// work with it
keyDel (k);
```

If you want the key object to contain a name, value, comment and other meta info read on.

Note

When you already have a key with similar properties its easier and cheaper to [keyDup\(\)](#) the key.

Due to ABI compatibility, the [Key](#) structure is not defined in kdb.h, only declared. So you can only declare pointers to Keys in your program, and allocate and free memory for them with [keyNew\(\)](#) and [keyDel\(\)](#) respectively. See <http://tldp.org/HOWTO/Program-Library-HOWTO/shared-libraries.html#A-EN135>

You can call it in many different ways depending on the attribute tags you pass as parameters. Tags are represented as the [keyswitch_t](#) values, and tell [keyNew\(\)](#) which [Key](#) attribute comes next.

The simplest and minimum way to use it is with no tags, only a key name:

```
Key *nullKey, *emptyNamedKey;

// Create a key that has no name, is completely empty, but is initialized
nullKey=keyNew(0);
keyDel (nullKey);

// Is the same as above
nullKey=keyNew("", KEY_END);
keyDel (nullKey);

// Create and initialize a key with a name and nothing else
emptyNamedKey=keyNew("user/some/example", KEY_END);
keyDel (emptyNamedKey);
```

[keyNew\(\)](#) allocates memory for a key object and cleans everything up. After that, it processes the given argument list.

The [Key](#) attribute tags are the following:

- [keyswitch_t::KEY_TYPE](#)

Next parameter is a type of the value. Default assumed is [KEY_TYPE_UNDEFINED](#). Set this attribute so that a subsequent [KEY_VALUE](#) can toggle to [keySetString\(\)](#) or [keySetBinary\(\)](#) regarding to [keyIsString\(\)](#) or [keyIsBinary\(\)](#). If you don't use [KEY_TYPE](#) but a [KEY_VALUE](#) follows afterwards, [KEY_TYPE_STRING](#) will be used.

- [keyswitch_t::KEY_SIZE](#)

Define a maximum length of the value. This is especially useful for setting a binary key. So make sure you use that before you [KEY_VALUE](#) for binary keys.

- [keyswitch_t::KEY_VALUE](#)

Next parameter is a pointer to the value that will be set to the key. If no [keyswitch_t::KEY_TYPE](#) was used before, [keyswitch_t::KEY_TYPE_STRING](#) is assumed. If [KEY_TYPE](#) was previously passed with a [KEY_TYPE_BINARY](#), you should have passed [KEY_SIZE](#) before! Otherwise it will be cut of with first `\0` in string!

- [keyswitch_t::KEY_UID](#), [keyswitch_t::KEY_GID](#)

Next parameter is taken as the UID (`uid_t`) or GID (`gid_t`) that will be defined on the key. See [keySetUID\(\)](#) and [keySetGID\(\)](#).

- [keyswitch_t::KEY_MODE](#)

Next parameter is taken as mode permissions (`mode_t`) to the key. See [keySetMode\(\)](#).

- [keyswitch_t::KEY_DIR](#)

Define that the key is a directory rather than a ordinary key. This means its executable bits in its mode are set. This option allows the key to have subkeys. See [keySetDir\(\)](#).

- [keyswitch_t::KEY_OWNER](#)

Next parameter is the owner. See [keySetOwner\(\)](#).

- `keyswitch_t::KEY_COMMENT`

Next parameter is a comment. See [keySetComment\(\)](#).

- `keyswitch_t::KEY_END`

Must be the last parameter passed to [keyNew\(\)](#). It is always required, unless the `keyName` is 0.

Example:

```
KeySet *ks=keyNew(0);

ksAppendKey(ks,keyNew(0));           // an empty key

ksAppendKey(ks,keyNew("user/sw",      // the name of
    KEY_END));                       // no more args

ksAppendKey(ks,keyNew("user/tmp/ex1",
    KEY_VALUE,"some data",           // set a string value
    KEY_END));                       // end of args

ksAppendKey(ks,keyNew("user/tmp/ex2",
    KEY_VALUE,"some data",           // with a simple value
    KEY_MODE,0777,                   // permissions
    KEY_END));                       // end of args

ksAppendKey(ks,keyNew("user/tmp/ex4",
    KEY_TYPE,KEY_TYPE_BINARY,        // key type
    KEY_SIZE,7,                      // assume binary length 7
    KEY_VALUE,"some data",           // value that will be truncated in 7
    bytes
    KEY_COMMENT,"value is truncated",
    KEY_OWNER,"root",               // owner (not uid) is root
    KEY_UID,0,                      // root uid
    KEY_END));                       // end of args

ksAppendKey(ks,keyNew("user/tmp/ex5",
    KEY_TYPE,
    KEY_TYPE_DIR | KEY_TYPE_BINARY, // dir key with a binary value
    KEY_SIZE,7,
    KEY_VALUE,"some data",           // value that will be truncated in 7
    bytes
    KEY_COMMENT,"value is truncated",
    KEY_OWNER,"root",               // owner (not uid) is root
    KEY_UID,0,                      // root uid
    KEY_END));                       // end of args

ksDel(ks);
```

The reference counter (see [keyGetRef\(\)](#)) will be initialized with 0, that means a subsequent call of [keyDel\(\)](#) will delete the key. If you append the key to a keyset the reference counter will be incremented by one (see [keyInc\(\)](#)) and the key can't be deleted by a [keyDel\(\)](#).

```
Key *k = keyNew(0); // ref counter 0
ksAppendKey(ks, k); // ref counter of key 1
ksDel(ks); // key will be deleted with keyset
*
```

If you increment only by one with [keyInc\(\)](#) the same as said above is valid:

```
Key *k = keyNew(0); // ref counter 0
keyIncRef(k); // ref counter of key 1
keyDel(k); // has no effect
keyDecRef(k); // ref counter back to 0
keyDel(k); // key is now deleted
*
```

If you add the key to more keySets:

```
Key *k = keyNew(0); // ref counter 0
ksAppendKey(ks1, k); // ref counter of key 1
ksAppendKey(ks2, k); // ref counter of key 2
ksDel(ks1); // ref counter of key 1
ksDel(ks2); // k is now deleted
*
```

or use `keyInc()` more than once:

```
Key *k = keyNew(0); // ref counter 0
keyIncRef(k); // ref counter of key 1
keyDel (k); // has no effect
keyIncRef(k); // ref counter of key 2
keyDel (k); // has no effect
keyDecRef(k); // ref counter of key 1
keyDel (k); // has no effect
keyDecRef(k); // ref counter is now 0
keyDel (k); // k is now deleted
*
```

they key won't be deleted by a `keyDel()` as long refcounter is not 0.

The key's sync bit will always be set for any call, except:

```
Key *k = keyNew(0);
// keyNeedSync() will be false
```

Parameters

<i>keyName</i>	a valid name to the key, or NULL to get a simple initialized, but really empty, object
----------------	--

See Also

[keyDel\(\)](#)

Returns

a pointer to a new allocated and initialized [Key](#) object, or NULL if an invalid `keyName` was passed (see [keySetName\(\)](#)).

Parameters

<i>keyName</i>	the name of the new key
----------------	-------------------------

8.2.2.6 kdb::Key::Key (const std::string keyName, ...) [inline],[explicit]

A practical way to fully create a [Key](#) object in one step.

This function tries to mimic the C++ way for constructors.

To just get a key object, simple do:

```
Key *k = keyNew(0);
// work with it
keyDel (k);
```

If you want the key object to contain a name, value, comment and other meta info read on.

Note

When you already have a key with similar properties its easier and cheaper to [keyDup\(\)](#) the key.

Due to ABI compatibility, the [Key](#) structure is not defined in `kdb.h`, only declared. So you can only declare pointers to Keys in your program, and allocate and free memory for them with [keyNew\(\)](#) and [keyDel\(\)](#) respectively. See <http://tldp.org/HOWTO/Program-Library-HOWTO/shared-libraries.html#A-EN135>

You can call it in many different ways depending on the attribute tags you pass as parameters. Tags are represented as the `keyswitch_t` values, and tell [keyNew\(\)](#) which [Key](#) attribute comes next.

The simplest and minimum way to use it is with no tags, only a key name:

```
Key *nullKey, *emptyNamedKey;

// Create a key that has no name, is completely empty, but is initialized
nullKey=keyNew(0);
keyDel (nullKey);

// Is the same as above
nullKey=keyNew("", KEY_END);
keyDel (nullKey);

// Create and initialize a key with a name and nothing else
emptyNamedKey=keyNew("user/some/example", KEY_END);
keyDel (emptyNamedKey);
```

[keyNew\(\)](#) allocates memory for a key object and cleans everything up. After that, it processes the given argument list.

The [Key](#) attribute tags are the following:

- [keyswitch_t::KEY_TYPE](#)

Next parameter is a type of the value. Default assumed is [KEY_TYPE_UNDEFINED](#). Set this attribute so that a subsequent [KEY_VALUE](#) can toggle to [keySetString\(\)](#) or [keySetBinary\(\)](#) regarding to [keyIsString\(\)](#) or [keyIsBinary\(\)](#). If you don't use [KEY_TYPE](#) but a [KEY_VALUE](#) follows afterwards, [KEY_TYPE_STRING](#) will be used.

- [keyswitch_t::KEY_SIZE](#)

Define a maximum length of the value. This is especially useful for setting a binary key. So make sure you use that before you [KEY_VALUE](#) for binary keys.

- [keyswitch_t::KEY_VALUE](#)

Next parameter is a pointer to the value that will be set to the key. If no [keyswitch_t::KEY_TYPE](#) was used before, [keyswitch_t::KEY_TYPE_STRING](#) is assumed. If [KEY_TYPE](#) was previously passed with a [KEY_TYPE_BINARY](#), you should have passed [KEY_SIZE](#) before! Otherwise it will be cut of with first \0 in string!

- [keyswitch_t::KEY_UID](#), [keyswitch_t::KEY_GID](#)

Next parameter is taken as the UID ([uid_t](#)) or GID ([gid_t](#)) that will be defined on the key. See [keySetUID\(\)](#) and [keySetGID\(\)](#).

- [keyswitch_t::KEY_MODE](#)

Next parameter is taken as mode permissions ([mode_t](#)) to the key. See [keySetMode\(\)](#).

- [keyswitch_t::KEY_DIR](#)

Define that the key is a directory rather than a ordinary key. This means its executable bits in its mode are set. This option allows the key to have subkeys. See [keySetDir\(\)](#).

- [keyswitch_t::KEY_OWNER](#)

Next parameter is the owner. See [keySetOwner\(\)](#).

- [keyswitch_t::KEY_COMMENT](#)

Next parameter is a comment. See [keySetComment\(\)](#).

- [keyswitch_t::KEY_END](#)

Must be the last parameter passed to [keyNew\(\)](#). It is always required, unless the [keyName](#) is 0.

Example:

```
KeySet *ks=keyNew(0);

keyAppendKey(ks, keyNew(0));           // an empty key

keyAppendKey(ks, keyNew("user/sw",    // the name of
    the key
    KEY_END));                         // no more args
```



```

ksAppendKey(ks, keyNew("user/tmp/ex1",
    KEY_VALUE, "some data",      // set a string value
    KEY_END));                  // end of args

ksAppendKey(ks, keyNew("user/tmp/ex2",
    KEY_VALUE, "some data",      // with a simple value
    KEY_MODE, 0777,             // permissions
    KEY_END));                  // end of args

ksAppendKey(ks, keyNew("user/tmp/ex4",
    KEY_TYPE, KEY_TYPE_BINARY,   // key type
    KEY_SIZE, 7,                 // assume binary length 7
    KEY_VALUE, "some data",      // value that will be truncated in 7
    bytes
    KEY_COMMENT, "value is truncated",
    KEY_OWNER, "root",          // owner (not uid) is root
    KEY_UID, 0,                 // root uid
    KEY_END));                  // end of args

ksAppendKey(ks, keyNew("user/tmp/ex5",
    KEY_TYPE,
    KEY_TYPE_DIR | KEY_TYPE_BINARY, // dir key with a binary value
    KEY_SIZE, 7,
    KEY_VALUE, "some data",      // value that will be truncated in 7
    bytes
    KEY_COMMENT, "value is truncated",
    KEY_OWNER, "root",          // owner (not uid) is root
    KEY_UID, 0,                 // root uid
    KEY_END));                  // end of args

ksDel(ks);

```

The reference counter (see [keyGetRef\(\)](#)) will be initialized with 0, that means a subsequent call of [keyDel\(\)](#) will delete the key. If you append the key to a keyset the reference counter will be incremented by one (see [keyInc\(\)](#)) and the key can't be deleted by a [keyDel\(\)](#).

```

Key *k = keyNew(0); // ref counter 0
ksAppendKey(ks, k); // ref counter of key 1
ksDel(ks); // key will be deleted with keyset
*

```

If you increment only by one with [keyInc\(\)](#) the same as said above is valid:

```

Key *k = keyNew(0); // ref counter 0
keyIncRef(k); // ref counter of key 1
keyDel(k);    // has no effect
keyDecRef(k); // ref counter back to 0
keyDel(k);    // key is now deleted
*

```

If you add the key to more keySets:

```

Key *k = keyNew(0); // ref counter 0
ksAppendKey(ks1, k); // ref counter of key 1
ksAppendKey(ks2, k); // ref counter of key 2
ksDel(ks1); // ref counter of key 1
ksDel(ks2); // k is now deleted
*

```

or use [keyInc\(\)](#) more than once:

```

Key *k = keyNew(0); // ref counter 0
keyIncRef(k); // ref counter of key 1
keyDel(k);    // has no effect
keyIncRef(k); // ref counter of key 2
keyDel(k);    // has no effect
keyDecRef(k); // ref counter of key 1
keyDel(k);    // has no effect
keyDecRef(k); // ref counter is now 0
keyDel(k);    // k is now deleted
*

```

the key won't be deleted by a [keyDel\(\)](#) as long refcounter is not 0.

The key's sync bit will always be set for any call, except:

```
Key *k = keyNew(0);
// keyNeedSync() will be false
```

Parameters

<i>keyName</i>	a valid name to the key, or NULL to get a simple initialized, but really empty, object
----------------	--

See Also

[keyDel\(\)](#)

Returns

a pointer to a new allocated and initialized [Key](#) object, or NULL if an invalid `keyName` was passed (see [keySetName\(\)](#)).

Warning

Not supported on some compilers, e.g. clang which require you to only pass non-POD in varg lists.

Parameters

<i>keyName</i>	the name of the new key
----------------	-------------------------

8.2.2.7 kdb::Key::Key (const char * *keyName*, va_list *ap*) [inline], [explicit]

A practical way to fully create a [Key](#) object in one step.

This function tries to mimic the C++ way for constructors.

To just get a key object, simple do:

```
Key *k = keyNew(0);
// work with it
keyDel (k);
```

If you want the key object to contain a name, value, comment and other meta info read on.

Note

When you already have a key with similar properties its easier and cheaper to [keyDup\(\)](#) the key.

Due to ABI compatibility, the [Key](#) structure is not defined in `kdb.h`, only declared. So you can only declare pointers to Keys in your program, and allocate and free memory for them with [keyNew\(\)](#) and [keyDel\(\)](#) respectively. See <http://tldp.org/HOWTO/Program-Library-HOWTO/shared-libraries.html#A-EN135>

You can call it in many different ways depending on the attribute tags you pass as parameters. Tags are represented as the [keyswitch_t](#) values, and tell [keyNew\(\)](#) which [Key](#) attribute comes next.

The simplest and minimum way to use it is with no tags, only a key name:

```
Key *nullKey, *emptyNamedKey;

// Create a key that has no name, is completely empty, but is initialized
nullKey=keyNew(0);
keyDel (nullKey);

// Is the same as above
nullKey=keyNew("", KEY_END);
keyDel (nullKey);
```

```
// Create and initialize a key with a name and nothing else
emptyNamedKey=keyNew("user/some/example",KEY_END);
keyDel (emptyNamedKey);
```

[keyNew\(\)](#) allocates memory for a key object and cleans everything up. After that, it processes the given argument list.

The [Key](#) attribute tags are the following:

- [keyswitch_t::KEY_TYPE](#)
Next parameter is a type of the value. Default assumed is [KEY_TYPE_UNDEFINED](#). Set this attribute so that a subsequent [KEY_VALUE](#) can toggle to [keySetString\(\)](#) or [keySetBinary\(\)](#) regarding to [keyIsString\(\)](#) or [keyIsBinary\(\)](#). If you don't use [KEY_TYPE](#) but a [KEY_VALUE](#) follows afterwards, [KEY_TYPE_STRING](#) will be used.
- [keyswitch_t::KEY_SIZE](#)
Define a maximum length of the value. This is especially useful for setting a binary key. So make sure you use that before you [KEY_VALUE](#) for binary keys.
- [keyswitch_t::KEY_VALUE](#)
Next parameter is a pointer to the value that will be set to the key. If no [keyswitch_t::KEY_TYPE](#) was used before, [keyswitch_t::KEY_TYPE_STRING](#) is assumed. If [KEY_TYPE](#) was previously passed with a [KEY_TYPE_BINARY](#), you should have passed [KEY_SIZE](#) before! Otherwise it will be cut off with first `\0` in string!
- [keyswitch_t::KEY_UID](#), [keyswitch_t::KEY_GID](#)
Next parameter is taken as the UID (`uid_t`) or GID (`gid_t`) that will be defined on the key. See [keySetUID\(\)](#) and [keySetGID\(\)](#).
- [keyswitch_t::KEY_MODE](#)
Next parameter is taken as mode permissions (`mode_t`) to the key. See [keySetMode\(\)](#).
- [keyswitch_t::KEY_DIR](#)
Define that the key is a directory rather than a ordinary key. This means its executable bits in its mode are set. This option allows the key to have subkeys. See [keySetDir\(\)](#).
- [keyswitch_t::KEY_OWNER](#)
Next parameter is the owner. See [keySetOwner\(\)](#).
- [keyswitch_t::KEY_COMMENT](#)
Next parameter is a comment. See [keySetComment\(\)](#).
- [keyswitch_t::KEY_END](#)
Must be the last parameter passed to [keyNew\(\)](#). It is always required, unless the `keyName` is 0.

Example:

```
KeySet *ks=keyNew(0);

ksAppendKey(ks,keyNew(0));           // an empty key

ksAppendKey(ks,keyNew("user/sw",      // the name of
    KEY_END));                       // no more args

ksAppendKey(ks,keyNew("user/tmp/ex1",
    KEY_VALUE,"some data",           // set a string value
    KEY_END));                       // end of args

ksAppendKey(ks,keyNew("user/tmp/ex2",
    KEY_VALUE,"some data",           // with a simple value
    KEY_MODE,0777,                   // permissions
    KEY_END));                       // end of args

ksAppendKey(ks,keyNew("user/tmp/ex4",
    KEY_TYPE,KEY_TYPE_BINARY,        // key type
```

```

        KEY_SIZE, 7, // assume binary length 7
        KEY_VALUE, "some data", // value that will be truncated in 7
        bytes
        KEY_COMMENT, "value is truncated",
        KEY_OWNER, "root", // owner (not uid) is root
        KEY_UID, 0, // root uid
        KEY_END)); // end of args

    ksAppendKey(ks, keyNew("user/tmp/ex5",
        KEY_TYPE,
        KEY_TYPE_DIR | KEY_TYPE_BINARY, // dir key with a binary value
        KEY_SIZE, 7,
        KEY_VALUE, "some data", // value that will be truncated in 7
        bytes
        KEY_COMMENT, "value is truncated",
        KEY_OWNER, "root", // owner (not uid) is root
        KEY_UID, 0, // root uid
        KEY_END)); // end of args

    ksDel(ks);

```

The reference counter (see [keyGetRef\(\)](#)) will be initialized with 0, that means a subsequent call of [keyDel\(\)](#) will delete the key. If you append the key to a keyset the reference counter will be incremented by one (see [keyInc\(\)](#)) and the key can't be deleted by a [keyDel\(\)](#).

```

Key *k = keyNew(0); // ref counter 0
ksAppendKey(ks, k); // ref counter of key 1
ksDel(ks); // key will be deleted with keyset
*

```

If you increment only by one with [keyInc\(\)](#) the same as said above is valid:

```

Key *k = keyNew(0); // ref counter 0
keyIncRef(k); // ref counter of key 1
keyDel(k); // has no effect
keyDecRef(k); // ref counter back to 0
keyDel(k); // key is now deleted
*

```

If you add the key to more keySets:

```

Key *k = keyNew(0); // ref counter 0
ksAppendKey(ks1, k); // ref counter of key 1
ksAppendKey(ks2, k); // ref counter of key 2
ksDel(ks1); // ref counter of key 1
ksDel(ks2); // k is now deleted
*

```

or use [keyInc\(\)](#) more than once:

```

Key *k = keyNew(0); // ref counter 0
keyIncRef(k); // ref counter of key 1
keyDel(k); // has no effect
keyIncRef(k); // ref counter of key 2
keyDel(k); // has no effect
keyDecRef(k); // ref counter of key 1
keyDel(k); // has no effect
keyDecRef(k); // ref counter is now 0
keyDel(k); // k is now deleted
*

```

they key won't be deleted by a [keyDel\(\)](#) as long refcounter is not 0.

The key's sync bit will always be set for any call, except:

```

Key *k = keyNew(0);
// keyNeedSync() will be false

```

Parameters

<i>keyName</i>	a valid name to the key, or NULL to get a simple initialized, but really empty, object
----------------	--

See Also

[keyDel\(\)](#)

Returns

a pointer to a new allocated and initialized [Key](#) object, or NULL if an invalid `keyName` was passed (see [keySetName\(\)](#)).

Parameters

<i>keyName</i>	the name of the new key
<i>ap</i>	the variable argument list pointer

8.2.2.8 kdb::Key::~~Key () [inline]

Destructs the key.

See Also

[del\(\)](#)

8.2.3 Member Function Documentation

8.2.3.1 void kdb::Key::addBaseName (const std::string & *baseName*) [inline]

Adds a base name for a key

Adds `baseName` to the current key name.

Assumes that `key` is a directory. `baseName` is appended to it. The function adds `'/'` if needed while concatenating.

So if `key` has name `"system/dir1/dir2"` and this method is called with `baseName "mykey"`, the resulting key will have name `"system/dir1/dir2/mykey"`.

When `baseName` is 0 or "" nothing will happen and the size of the name is returned.

Warning

You should not change a keys name once it belongs to a keyset. See `ksSort()` for more information.

TODO: does not recognise `..` and `.` in the string!

Parameters

<i>key</i>	the key object to work with
<i>baseName</i>	the string to append to the name

Returns

the size in bytes of the new key name including the ending NULL
 -1 if the key had no name
 -1 on NULL pointers

See Also

[keySetBaseName\(\)](#)
[keySetName\(\)](#) to [set](#) a new name.

Exceptions

<i>KeyInvalidName</i>	if the name is not valid
-----------------------	--------------------------

8.2.3.2 void kdb::Key::clear () [inline]

Clears/Invalidates a key.

Afterwards the object is empty again.

Note

This is not a null key, so it will evaluate to true. [isValid\(\)](#) will, however, be false.

See Also

[release\(\)](#)
[isValid\(\)](#), operator bool()

[Key](#) Object Cleaner.

Will reset all internal data.

After this call you will receive a fresh key.

The reference counter will stay unmodified.

Note

that you might also [clear\(\)](#) all aliases with this operation.

```
int f (Key *k)
{
    keyClear (k);
    // you have a fresh key k here
    keySetString (k, "value");
    // the caller will get an empty key k with an value
}
```

Returns

returns 0 on success
-1 on null pointer

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

8.2.3.3 void kdb::Key::copy (const Key & other) [inline]

Copy or Clear a key.

Most often you may prefer [keyDup\(\)](#) which allocates a new key and returns a duplication of another key.

But when you need to copy into an existing key, e.g. because it was passed by a pointer in a function you can do

so:

```
void h (Key *k)
{
    // receive key c
    keyCopy (k, c);
    // the caller will see the changed key k
}
```

The reference counter will not be changed for both keys. Affiliation to keysets are also not affected.

When you pass a NULL-pointer as source the data of dest will be cleaned completely (except reference counter, see [keyClear\(\)](#)) and you get a fresh dest key.

```
void g (Key *k)
{
    keyCopy (k, 0);
    // k is now an empty and fresh key
}
```

The meta data will be duplicated for the destination key. So it will not take much additional space, even with lots of metadata.

If you want to copy all metadata, but keep the old value you can use [keyCopy\(\)](#) too.

```
void j (Key *k)
{
    size_t size = keyGetValueSize (k);
    char *value = malloc (size);
    int bstring = keyIsString (k);

    // receive key c
    memcpy (value, keyValue(k), size);
    keyCopy (k, c);
    if (bstring) keySetString (k, value);
    else keySetBinary (k, value, size);
    free (value);
    // the caller will see the changed key k
    // with the metadata from c
}
```

Note

Next to the value itself we also need to remember if the value was string or binary. So in fact the meta data of the resulting key k in that example is not a complete duplicate, because the meta data "binary" may differ. Similar considerations might be necessary for the type of the key and so on, depending on the concrete situation.

Parameters

<i>dest</i>	the key which will be written to
<i>source</i>	the key which should be copied or NULL to clean the destination key

Returns

-1 on failure when a NULL pointer was passed for dest or a dynamic property could not be written. Both name and value are empty then.
0 when dest was cleaned
1 when source was successfully copied

See Also

[keyDup\(\)](#) to get a duplication of a [Key](#)

8.2.3.4 void kdb::Key::copyAllMeta (const Key & other) [inline]

Do a shallow copy of all meta data from source to dest.

The key dest will additionally have all meta data source had. Meta data not present in source will not be changed. Meta data which was present in source and dest will be overwritten.

For example the meta data type is copied into the [Key](#) k.

```
void l(Key *k)
{
    // receive c
    keyCopyMeta(k, c);
    // the caller will see the changed key k
    // with all the metadata from c
}
```

The main purpose of this function is for plugins or applications which want to add the same meta data to n keys. When you do that with [keySetMeta\(\)](#) it will take n times the memory for the key. This can be considerable amount of memory for many keys with some meta data for each.

To avoid that problem you can use [keyCopyAllMeta\(\)](#) or [keyCopyMeta\(\)](#).

```
void o(KeySet *ks)
{
    Key *current;
    Key *shared = keyNew (0);
    keySetMeta(shared, "shared1", "this meta data should be
shared among many keys");
    keySetMeta(shared, "shared2", "this meta data should be
shared among many keys also");
    keySetMeta(shared, "shared3", "this meta data should be
shared among many keys too");

    ksRewind(ks);
    while ((current = ksNext(ks)) != 0)
    {
        if (needs_shared_data(current)) keyCopyAllMeta(
current, shared);
    }
}
```

Postcondition

for every metaName present in source: keyGetMeta(source, metaName) == keyGetMeta(dest, metaName)

Returns

- 1 if was successfully copied
- 0 if source did not have any meta data
- 1 on null pointers (source or dest)
- 1 on memory problems

Parameters

<i>dest</i>	the destination where the meta data should be copied too
<i>source</i>	the key where the meta data should be copied from

See Also

[getMeta\(\)](#), [setMeta\(\)](#), [copyMeta\(\)](#)

8.2.3.5 void kdb::Key::copyMeta (const Key & other, const std::string & metaName) [inline]

Do a shallow copy of meta data from source to dest.

The key dest will have the same meta data referred with metaName afterwards then source.

For example the meta data type is copied into the [Key](#) k.

```
void l(Key *k)
{
    // receive c
    keyCopyMeta(k, c, "type");
    // the caller will see the changed key k
    // with the metadata "type" from c
}
```

The main purpose of this function is for plugins or applications which want to add the same meta data to n keys. When you do that with [keySetMeta\(\)](#) it will take n times the memory for the key. This can be considerable amount of memory for many keys with some meta data for each.

To avoid that problem you can use [keyCopyAllMeta\(\)](#) or [keyCopyMeta\(\)](#).

```
void o(KeySet *ks)
{
    Key *current;
    Key *shared = keyNew (0);
    keySetMeta(shared, "shared", "this meta data should be shared
among many keys");

    ksRewind(ks);
    while ((current = ksNext(ks)) != 0)
    {
        if (needs_shared_data(current)) keyCopyMeta(current,
shared, "shared");
    }
}
```

Postcondition

`keyGetMeta(source, metaName) == keyGetMeta(dest, metaName)`

Returns

- 1 if was successfully copied
- 0 if the meta data in dest was removed too
- 1 on null pointers (source or dest)
- 1 on memory problems

Parameters

<i>dest</i>	the destination where the meta data should be copied too
<i>source</i>	the key where the meta data should be copied from
<i>metaName</i>	the name of the meta data which should be copied

See Also

[getMeta\(\)](#), [setMeta\(\)](#), [copyAllMeta\(\)](#)

8.2.3.6 const Key kdb::Key::currentMeta () const [inline]

Returns the Value of a Meta-Information which is current.

The pointer is NULL if you reached the end or after [ksRewind\(\)](#).

Note

You must not delete or change the returned key, use [keySetMeta\(\)](#) if you want to delete or change it.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

a buffer to the value pointed by *key*'s cursor
0 on NULL pointer

See Also

[keyNextMeta\(\)](#), [keyRewindMeta\(\)](#)
[ksCurrent\(\)](#) for pedant in iterator interface of [KeySet](#)

Note

that the key will be null if last meta data is found.

```
k.rewindMeta();
while (meta = k.nextMeta())
{
    cout << meta.getName() << " " << meta.getString() << endl;
}
```

See Also

[rewindMeta\(\)](#), [nextMeta\(\)](#)

8.2.3.7 `ckdb::Key * kdb::Key::dup () const` `[inline]`

Return a duplicate of a key.

Memory will be allocated as needed for dynamic properties.

The new key will not be member of any [KeySet](#) and will start with a new reference counter at 0. A subsequent [keyDel\(\)](#) will delete the key.

```
int f (const Key * source)
{
    Key * dup = keyDup (source);
    // work with duplicate
    keyDel (dup);
    // everything related to dup is freed
    // and source is unchanged
}
```

Like for a new key after [keyNew\(\)](#) a subsequent [ksAppend\(\)](#) makes a [KeySet](#) to take care of the lifecycle of the key.

```
int g (const Key * source, KeySet * ks)
{
    Key * dup = keyDup (source);
    // work with duplicate
    ksAppendKey (ks, dup);
    // ksDel(ks) will also free the duplicate
    // source remains unchanged.
}
```

Duplication of keys should be preferred to [keyNew\(\)](#), because data like owner can be filled with a copy of the key instead of asking the environment. It can also be optimized in the checks, because the keyname is known to be valid.

Parameters

<i>source</i>	has to be an initialised source Key
---------------	---

Returns

0 failure or on NULL pointer
a fully copy of source on success

See Also

[ksAppend\(\)](#), [keyDel\(\)](#), [keyNew\(\)](#)

8.2.3.8 std::string kdb::Key::get () const [inline]

Get a key value.

You can write your own template specialication, e.g.:

Returns

the string directly from the key.

It should be the same as [get\(\)](#).

Returns

empty string on null pointers

Exceptions

<i>KeyException</i>	on null key or not a valid size
<i>KeyTypeMismatch</i>	if key holds binary data and not a string

Note

unlike in the C version, it is safe to change the returned string.

See Also

[isString\(\)](#), [getBinary\(\)](#)

This method tries to serialize the string to the given type.

8.2.3.9 std::string kdb::Key::getBaseName () const [inline]

Returns a pointer to the real internal key name where the `basename` starts.

This is a much more efficient version of [keyGetBaseName\(\)](#) and you should use it if you are responsible enough to not mess up things. The name might change or even point to a wrong place after a [keySetName\(\)](#). If you need a copy of the basename consider to use [keyGetBaseName\(\)](#).

[keyBaseName\(\)](#) returns "" when there is no keyBaseName. The reason is

```
key=keyNew(0);
keySetName(key, "");
keyBaseName(key); // you would expect "" here
keySetName(key, "user");
keyBaseName(key); // you would expect "" here
keyDel(key);
```

Note

Note that the [Key](#) structure keeps its own size field that is calculated by library internal calls, so to avoid inconsistencies, you must never use the pointer returned by [keyBaseName\(\)](#) method to set a new value. Use [keySetBaseName\(\)](#) instead.

Parameters

<i>key</i>	the object to obtain the basename from
------------	--

Returns

a pointer to the basename
"" when the key has no (base)name
0 on NULL pointer

See Also

[keyGetBaseName\(\)](#), [keyGetBaseNameSize\(\)](#)
[keyName\(\)](#) to [get](#) a pointer to the name
[keyOwner\(\)](#) to [get](#) a pointer to the owner

8.2.3.10 `size_t kdb::Key::getBaseNameSize () const [inline]`

Calculates number of bytes needed to store basename of *key*.

[Key](#) names that have only root names (e.g. "system" or "user" or "user:domain") does not have base-names, thus the function will return 1 bytes to store "".

Basenames are denoted as:

- system/some/thing/basename -> basename
- user:domain/some/thing/base\name > base\name

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

size in bytes of *key*'s basename including ending NULL

See Also

[keyBaseName\(\)](#), [keyGetBaseName\(\)](#)
[keyName\(\)](#), [keyGetName\(\)](#), [keySetName\(\)](#)

8.2.3.11 `std::string kdb::Key::getBinary () const [inline]`**Returns**

the binary Value of the key.

Return values

""	on null pointers (size == 0) and on data only containing \0
----	---

Note

if you need to distinguish between null pointers and data containing \0 you can use [getValue\(\)](#).

Exceptions

<i>KeyException</i>	on invalid binary size
<i>KeyTypeMismatch</i>	if key is string and not a binary

Get the value of a key as a binary.

If the type is not binary -1 will be returned.

When the binary data is empty (this is not the same as "") 0 will be returned and the returnedBinary will not be changed.

For string values see [keyGetString\(\)](#) and [keyIsString\(\)](#).

When the returnedBinary is too small to hold the data (its maximum size is given by maxSize), the returnedBinary will not be changed and -1 is returned.

Example:

```
Key *key = keyNew ("user/keyname", KEY_TYPE, KEY_TYPE_BINARY, KEY_END);
char buffer[300];

if (keyGetBinary(key,buffer,sizeof(buffer)) == -1)
{
    // handle error
}
```

Parameters

<i>key</i>	the object to gather the value from
<i>returnedBinary</i>	pre-allocated memory to store a copy of the key value
<i>maxSize</i>	number of bytes of pre-allocated memory in returnedBinary

Returns

the number of bytes actually copied to returnedBinary

Return values

0	if the binary is empty
-1	on NULL pointers
-1	if maxSize is 0
-1	if maxSize is too small for string
-1	if maxSize is larger than SSIZE_MAX
-1	on type mismatch: binary expected, but found string

See Also

[keyValue\(\)](#), [keyGetValueSize\(\)](#), [keySetBinary\(\)](#)
[keyGetString\(\)](#) and [keySetString\(\)](#) as preferred alternative to binary
[keyIsBinary\(\)](#) to see how to check for binary type

See Also

[isBinary\(\)](#), [getString\(\)](#), [getValue\(\)](#)

8.2.3.12 `size_t kdb::Key::getBinarySize () const` `[inline]`

Returns the number of bytes needed to store the key value, including the NULL terminator.

It returns the correct size, independent of the [Key](#) Type. If it is a binary there might be '\0' values in it.

For an empty string you need one byte to store the ending NULL. For that reason 1 is returned. This is not true for binary data, so there might be returned 0 too.

A binary key has no '\0' termination. String types have it, so to there length will be added 1 to have enough space to store it.

This method can be used with `malloc()` before [keyGetString\(\)](#) or [keyGetBinary\(\)](#) is called.

```
char *buffer;
buffer = malloc (keyGetValueSize (key));
// use this buffer to store the value (binary or string)
// pass keyGetValueSize (key) for maxSize
```

Parameters

<code>key</code>	the key object to work with
------------------	-----------------------------

Returns

the number of bytes needed to store the key value
 1 when there is no data and type is not binary
 0 when there is no data and type is binary
 -1 on null pointer

See Also

[keyGetString\(\)](#), [keyGetBinary\(\)](#), [keyValue\(\)](#)

8.2.3.13 `std::string kdb::Key::getDirName () const` `[inline]`

Returns

the dir name of the key

e.g. `system/sw/dir/key` will return `system/sw/dir`

8.2.3.14 `std::string kdb::Key::getFullName () const` `[inline]`

Get key full name, including the user domain name.

Returns

number of bytes written
 1 on empty name
 -1 on NULL pointers
 -1 if maxSize is 0 or larger than `SSIZE_MAX`

Parameters

<i>key</i>	the key object
<i>returnedName</i>	pre-allocated memory to write the key name
<i>maxSize</i>	maximum number of bytes that will fit in returnedName, including the final NULL

Exceptions

<i>KeyException</i>	if key is null
---------------------	----------------

8.2.3.15 `size_t kdb::Key::getFullNameSize () const [inline]`

Bytes needed to store the key name including user domain and ending NULL.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

number of bytes needed to store key name including user domain
 1 on empty name
 -1 on NULL pointer

See Also

[keyGetFullName\(\)](#), [keyGetNameSize\(\)](#)

8.2.3.16 `Key::func_t kdb::Key::getFunc () const [inline]`

Elektra can store function pointers as binary. This function returns such a function pointer.

Exceptions

<i>KeyTypeMismatch</i>	if no binary data found, or binary data has not correct length
------------------------	--

Returns

a function pointer stored with [setBinary\(\)](#)

8.2.3.17 `ckdb::Key * kdb::Key::getKey () const [inline]`

Passes out the raw key pointer.

This pointer can be used to directly change the underlying key object.

Note

that the ownership remains in the object

8.2.3.18 `std::string kdb::Key::getMeta (const std::string & metaName) const [inline]`

Returns the Value of a Meta-Information given by name.

This is a much more efficient version of [keyGetMeta\(\)](#). But unlike with `keyGetMeta` you are not allowed to modify the resulting string.

```
int f(Key *k)
{
    if (!strcmp(keyValue(keyGetMeta(k, "type")), "boolean"))
    {
        // the type of the key is boolean
    }
}
```

Note

You must not delete or change the returned key, use [keySetMeta\(\)](#) if you want to delete or change it.

Parameters

<i>key</i>	the key object to work with
<i>metaName</i>	the name of the meta information you want the value from

Returns

0 if the key or metaName is 0
 0 if no such metaName is found
 value of Meta-Information if Meta-Information is found

See Also

[keyGetMeta\(\)](#), [keySetMeta\(\)](#)

You can specify your own template specialisation:

```
template<>
inline mode_t Key::getMeta(const std::string &name) const
{
    mode_t x;
    std::string str;
    str = std::string(
        static_cast<const char*>(
            ckdb::keyValue(
                ckdb::keyGetMeta(key, name.
            c_str())
        )
    );
    std::istringstream ist(str);
    ist >> std::oct >> x; // convert string to type
    return x;
}
```

Note

Because `mode_t` is in fact an `int`, this would also change all other `int` types.

Exceptions

<i>KeyBadMeta</i>	if meta data could not be parsed
<i>KeyNoSuchMeta</i>	if a value was requested, but none is available.

Note

No exception will be thrown if a `const Key` or `char*` is requested, but don't forget the `const`: [getMeta<const ckdb::Key*>](#), otherwise you will get an obfuscated compiler error.

- `char*` is null if meta data is not available
- `const Key` is null (evaluate to false) if no meta data is available

See Also

[setMeta\(\)](#), [copyMeta\(\)](#), [copyAllMeta\(\)](#)

8.2.3.19 `std::string kdb::Key::getName () const [inline]`

Returns a pointer to the abbreviated real internal `key` name.

This is a much more efficient version of [keyGetName\(\)](#) and can use it if you are responsible enough to not mess up things. You are not allowed to change anything in the returned array. The content of that string may change after [keySetName\(\)](#) and similar functions. If you need a copy of the name, consider using [keyGetName\(\)](#).

The name will be without owner, see [keyGetFullName\(\)](#) if you need the name with its owner.

[keyName\(\)](#) returns "" when there is no `keyName`. The reason is

```
key=keyNew(0);
keySetName(key, "");
keyName(key); // you would expect "" here
keyDel(key);
```

Note

Note that the `Key` structure keeps its own size field that is calculated by library internal calls, so to avoid inconsistencies, you must never use the pointer returned by [keyName\(\)](#) method to set a new value. Use [keySetName\(\)](#) instead.

Parameters

<code>key</code>	the key object to work with
------------------	-----------------------------

Returns

a pointer to the keyname which must not be changed.
 "" when there is no (a empty) keyname
 0 on NULL pointer

See Also

[keyGetNameSize\(\)](#) for the string length
[keyGetFullName\(\)](#), [keyGetFullNameSize\(\)](#) to [get](#) the full name
[keyGetName\(\)](#) as alternative to [get](#) a [copy](#)
[keyOwner\(\)](#) to [get](#) a pointer to owner

Note

unlike in the C version, it is safe to change the returned string.

8.2.3.20 `size_t kdb::Key::getNameSize () const [inline]`

Bytes needed to store the key name without owner.

For an empty key name you need one byte to store the ending NULL. For that reason 1 is returned.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

number of bytes needed, including ending NULL, to store key name without owner
 1 if there is no key Name
 -1 on NULL pointer

See Also

[keyGetName\(\)](#), [keyGetFullNameSize\(\)](#)

8.2.3.21 `size_t kdb::Key::getReferenceCounter () const` [inline]

Return how many references the key has.

The references will be incremented on successful calls to [ksAppendKey\(\)](#) or [ksAppend\(\)](#).

Note

[keyDup\(\)](#) will reset the references for dapped key.

For your own applications you can use [keyIncRef\(\)](#) and [keyDecRef\(\)](#) for reference counting. Keys with zero references will be deleted when using [keyDel\(\)](#).

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

the number of references
 -1 on null pointer

See Also

[keyIncRef\(\)](#) and [keyDecRef\(\)](#)

8.2.3.22 `std::string kdb::Key::getString () const` [inline]

Returns

the string directly from the key.

It should be the same as [get\(\)](#).

Returns

empty string on null pointers

Exceptions

<i>KeyException</i>	on null key or not a valid size
<i>KeyTypeMismatch</i>	if key holds binary data and not a string

Note

unlike in the C version, it is safe to change the returned string.

See Also

[isString\(\)](#), [getBinary\(\)](#)

8.2.3.23 `size_t kdb::Key::getStringSize () const` `[inline]`

Returns the number of bytes needed to store the key value, including the NULL terminator.

It returns the correct size, independent of the [Key](#) Type. If it is a binary there might be '\0' values in it.

For an empty string you need one byte to store the ending NULL. For that reason 1 is returned. This is not true for binary data, so there might be returned 0 too.

A binary key has no '\0' termination. String types have it, so to there length will be added 1 to have enough space to store it.

This method can be used with `malloc()` before [keyGetString\(\)](#) or [keyGetBinary\(\)](#) is called.

```
char *buffer;
buffer = malloc (keyGetValueSize (key));
// use this buffer to store the value (binary or string)
// pass keyGetValueSize (key) for maxSize
```

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

the number of bytes needed to store the key value
 1 when there is no data and type is not binary
 0 when there is no data and type is binary
 -1 on null pointer

See Also

[keyGetString\(\)](#), [keyGetBinary\(\)](#), [keyValue\(\)](#)

8.2.3.24 `const void * kdb::Key::getValue () const` `[inline]`

Return a pointer to the real internal `key` value.

This is a much more efficient version of [keyGetString\(\)](#) [keyGetBinary\(\)](#), and you should use it if you are responsible enough to not mess up things. You are not allowed to modify anything in the returned string. If you need a copy of the Value, consider to use [keyGetString\(\)](#) or [keyGetBinary\(\)](#) instead.

Returns

the value of the key

See Also

[getBinary\(\)](#)

8.2.3.25 `bool kdb::Key::isBelow (const Key & k) const` `[inline]`

Check if the key check is below the key key or not.

Example:

```
key user/sw/app
check user/sw/app/key
```

returns true because check is below key

Example:

```
key user/sw/app
check user/sw/app/folder/key
```

returns also true because check is indirect below key

Parameters

<i>key</i>	the key object to work with
<i>check</i>	the key to find the relative position of

Returns

1 if check is below key
0 if it is not below or if it is the same key

See Also

[keySetName\(\)](#), [keyGetName\(\)](#), [keyIsDirectBelow\(\)](#)

8.2.3.26 `bool kdb::Key::isBelowOrSame (const Key & k) const` `[inline]`**8.2.3.27** `bool kdb::Key::isBinary () const` `[inline]`

Check if a key is binary type.

The function checks if the key is a binary. Opposed to string values binary values can have '\0' inside the value and may not be terminated by a null character. Their disadvantage is that you need to pass their size.

Make sure to use this function and don't test the binary type another way to ensure compatibility and to write less error prone programs.

Returns

1 if it is binary
0 if it is not
-1 on NULL pointer

See Also

[keyGetBinary\(\)](#), [keySetBinary\(\)](#)

Parameters

<i>key</i>	the key to check
------------	------------------

8.2.3.28 `bool kdb::Key::isDirectBelow (const Key & k) const` `[inline]`

Check if the key check is direct below the key key or not.

Example:

```
key user/sw/app
check user/sw/app/key
```

returns true because check is below key

Example:

```
key user/sw/app
check user/sw/app/folder/key
```

does not return true, because there is only a indirect relation

Parameters

<i>key</i>	the key object to work with
<i>check</i>	the key to find the relative position of

Returns

1 if check is below key
0 if it is not below or if it is the same key
-1 on null pointer

See Also

[keyIsBelow\(\)](#), [keySetName\(\)](#), [keyGetName\(\)](#)

8.2.3.29 `bool kdb::Key::isInactive () const` `[inline]`

Check whether a key is inactive or not.

In elektra terminology any key is inactive if the it's basename starts with '.'. Inactive keys must not have any meaning to applications, they are reserved for users and administrators.

To remove a whole hierarchy in elektra, don't forget to pass option_t::KDB_O_INACTIVE to [kdbGet\(\)](#) to receive the inactive keys in order to remove them.

Otherwise you should not fetch these keys.

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

1 if the key is inactive, 0 otherwise
-1 on NULL pointer or when key has no name

8.2.3.30 `bool kdb::Key::isString () const` `[inline]`

Check if a key is string type.

String values are null terminated and are not allowed to have any '\0' characters inside the string.

Make sure to use this function and don't test the string type another way to ensure compatibility and to write less error prone programs.

Returns

1 if it is string
0 if it is not
-1 on NULL pointer

See Also

[keyGetString\(\)](#), [keySetString\(\)](#)

Parameters

<i>key</i>	the key to check
------------	------------------

8.2.3.31 `bool kdb::Key::isSystem () const [inline]`

Name starts with "system".

Return values

<i>true</i>	if it is a system key
<i>false</i>	otherwise

8.2.3.32 `bool kdb::Key::isUser () const [inline]`

Name starts with "user".

Return values

<i>true</i>	if it is a user key
<i>false</i>	otherwise

8.2.3.33 `bool kdb::Key::isValid () const [inline]`**Returns**

if the key is valid

An invalid key has no name. The name of valid keys either start with user or system.

Return values

<i>true</i>	if the key has a valid name
<i>false</i>	if the key has an invalid name

See Also

[getName\(\)](#), [isUser\(\)](#), [isSystem\(\)](#)

8.2.3.34 `const Key kdb::Key::nextMeta () [inline]`

Iterate to the next meta information.

Keys have an internal cursor that can be reset with [keyRewindMeta\(\)](#). Every time [keyNextMeta\(\)](#) is called the cursor is incremented and the new current Name of Meta Information is returned.

You'll get a NULL pointer if the meta information after the end of the [Key](#) was reached. On subsequent calls of [keyNextMeta\(\)](#) it will still return the NULL pointer.

The `key` internal cursor will be changed, so it is not const.

Note

That the resulting key is guaranteed to have a value, because meta information has no binary or null pointer semantics.

You must not delete or change the returned key, use [keySetMeta\(\)](#) if you want to delete or change it.

Parameters

<code>key</code>	the key object to work with
------------------	-----------------------------

Returns

a key representing meta information

0 when the end is reached

0 on NULL pointer

See Also

[ksNext\(\)](#) for pedant in iterator interface of [KeySet](#)

See Also

[rewindMeta\(\)](#), [currentMeta\(\)](#)

8.2.3.35 kdb::Key::operator bool () const [inline]

This is for loops and lookups only.

For loops it checks if there are still more keys. For lookups it checks if a key could be found.

Warning

you should not construct or use null keys

Returns

false on null keys

true otherwise

8.2.3.36 bool kdb::Key::operator!= (const Key & k) const [inline]

Compare the name of two keys.

Returns

a number less than, equal to or greater than zero if k1 is found, respectively, to be less than, to match, or be greater than k2.

The comparison is based on a strcmp of the keynames, and iff they match a strcmp of the owner will be used to distinguish. If even this matches the keys are found to be exactly the same and 0 is returned. These two keys can't be used in the same [KeySet](#).

[keyCmp\(\)](#) defines the sorting order for a [KeySet](#).

The following 3 points are the rules for null values. They only take account when none of the preceding rules matched.

- A null pointer will be found to be smaller than every other key. If both are null pointers, 0 is returned.
- A null name will be found to be smaller than every other name. If both are null names, 0 is returned.
- No owner will be found to be smaller then every other owner. If both don't have a owner, 0 is returned.

Note

the owner will only be used if the names are equal.

Often is enough to know if the other key is less then or greater then the other one. But Sometimes you need more precise information, see [keyRel\(\)](#).

Given any Keys k1 and k2 constructed with [keyNew\(\)](#), following equation hold true:

```
// keyCmp(0,0) == 0
// keyCmp(k1,0) == 1
// keyCmp(0,k2) == -1
```

You can write similar equation for the other rules.

Here are some more examples with equation:

```
Key *k1 = keyNew("user/a", KEY_END);
Key *k2 = keyNew("user/b", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0

Key *k1 = keyNew("user/a", KEY_OWNER, "markus", KEY_END);
Key *k2 = keyNew("user/a", KEY_OWNER, "max", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0
```

Warning

Do not try to strcmp the [keyName\(\)](#) yourself because the used strcmp implementation is allowed to differ from simple ascii comparison.

Parameters

<i>k1</i>	the first key object to compare with
<i>k2</i>	the second key object to compare with

See Also

[ksAppendKey\(\)](#), [ksAppend\(\)](#) will compare keys when appending
[ksLookup\(\)](#) will compare keys during searching

Return values

<i>true</i>	!= 0
-------------	------

8.2.3.37 `ckdb::Key * kdb::Key::operator*() const` `[inline]`

Passes out the raw key pointer.

This pointer can be used to directly change the underlying key object.

Note

that the ownership remains in the object

8.2.3.38 `void kdb::Key::operator++(int) const` `[inline]`

Increment the viability of a key object.

This function is intended for applications using their own reference counter for key objects. With it you can increment the reference and thus avoid destruction of the object in a subsequent [keyDel\(\)](#).

```
Key *k;
keyInc (k);
function_that_keyDec(k);
// work with k
keyDel (k); // now really free it
```

The reference counter can't be incremented once it reached SSIZE_MAX. In that situation nothing will happen and SSIZE_MAX will be returned.

Note

[keyDup\(\)](#) will reset the references for dupped key.

Returns

the value of the new reference counter
-1 on null pointer
SSIZE_MAX when maximum exceeded

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

See Also

[keyGetRef\(\)](#), [keyDecRef\(\)](#), [keyDel\(\)](#)

8.2.3.39 `void kdb::Key::operator++() const` `[inline]`

Increment the viability of a key object.

This function is intended for applications using their own reference counter for key objects. With it you can increment the reference and thus avoid destruction of the object in a subsequent [keyDel\(\)](#).

```
Key *k;
keyInc (k);
function_that_keyDec(k);
// work with k
keyDel (k); // now really free it
```

The reference counter can't be incremented once it reached SSIZE_MAX. In that situation nothing will happen and SSIZE_MAX will be returned.

Note

[keyDup\(\)](#) will reset the references for dupped key.

Returns

the value of the new reference counter
 -1 on null pointer
 SSIZE_MAX when maximum exceeded

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

See Also

[keyGetRef\(\)](#), [keyDecRef\(\)](#), [keyDel\(\)](#)

8.2.3.40 **Key & kdb::Key::operator+= (const std::string & *newAddBaseName*)** `[inline]`

Add a new basename.

See Also

[keyAddBaseName\(\)](#)

8.2.3.41 **Key & kdb::Key::operator+= (const char * *newAddBaseName*)** `[inline]`

(const std::string &)

Add a new basename.

See Also

[keyAddBaseName\(\)](#)

(const std::string &)

8.2.3.42 **void kdb::Key::operator-- (int) const** `[inline]`

Decrement the viability of a key object.

The references will be decremented for [ksPop\(\)](#) or successful calls of [ksLookup\(\)](#) with the option KDB_O_POP. It will also be decremented with an following [keyDel\(\)](#) in the case that an old key is replaced with another key with the same name.

The reference counter can't be decremented once it reached 0. In that situation nothing will happen and 0 will be returned.

Note

[keyDup\(\)](#) will reset the references for dupped key.

Returns

the value of the new reference counter
 -1 on null pointer
 0 when the key is ready to be freed

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

See Also

[keyGetRef\(\)](#), [keyDel\(\)](#), [keyIncRef\(\)](#)

8.2.3.43 void kdb::Key::operator-- () const [inline]

Decrement the viability of a key object.

The references will be decremented for [ksPop\(\)](#) or successful calls of [ksLookup\(\)](#) with the option KDB_O_POP. It will also be decremented with an following [keyDel\(\)](#) in the case that an old key is replaced with another key with the same name.

The reference counter can't be decremented once it reached 0. In that situation nothing will happen and 0 will be returned.

Note

[keyDup\(\)](#) will reset the references for dupped key.

Returns

the value of the new reference counter
-1 on null pointer
0 when the key is ready to be freed

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

See Also

[keyGetRef\(\)](#), [keyDel\(\)](#), [keyIncRef\(\)](#)

8.2.3.44 Key & kdb::Key::operator=(const std::string & newSetBaseName) [inline]

Set a new basename.

See Also

[keySetBaseName\(\)](#)

8.2.3.45 Key & kdb::Key::operator=(const char * newSetBaseName) [inline]

(const std::string &)

Set a new basename.

See Also

[keySetBaseName\(\)](#)

(const std::string &)

8.2.3.46 `bool kdb::Key::operator< (const Key & other) const [inline]`

Compare the name of two keys.

Returns

a number less than, equal to or greater than zero if k1 is found, respectively, to be less than, to match, or be greater than k2.

The comparison is based on a strcmp of the keynames, and iff they match a strcmp of the owner will be used to distinguish. If even this matches the keys are found to be exactly the same and 0 is returned. These two keys can't be used in the same [KeySet](#).

[keyCmp\(\)](#) defines the sorting order for a [KeySet](#).

The following 3 points are the rules for null values. They only take account when none of the preceding rules matched.

- A null pointer will be found to be smaller than every other key. If both are null pointers, 0 is returned.
- A null name will be found to be smaller than every other name. If both are null names, 0 is returned.
- No owner will be found to be smaller than every other owner. If both don't have a owner, 0 is returned.

Note

the owner will only be used if the names are equal.

Often is enough to know if the other key is less then or greater then the other one. But Sometimes you need more precise information, see [keyRel\(\)](#).

Given any Keys k1 and k2 constructed with [keyNew\(\)](#), following equation hold true:

```
// keyCmp(0,0) == 0
// keyCmp(k1,0) == 1
// keyCmp(0,k2) == -1
```

You can write similar equation for the other rules.

Here are some more examples with equation:

```
Key *k1 = keyNew("user/a", KEY_END);
Key *k2 = keyNew("user/b", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0

Key *k1 = keyNew("user/a", KEY_OWNER, "markus", KEY_END);
Key *k2 = keyNew("user/a", KEY_OWNER, "max", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0
```

Warning

Do not try to strcmp the [keyName\(\)](#) yourself because the used strcmp implementation is allowed to differ from simple ascii comparison.

Parameters

<i>k1</i>	the first key object to compare with
<i>k2</i>	the second key object to compare with

See Also

[ksAppendKey\(\)](#), [ksAppend\(\)](#) will compare keys when appending
[ksLookup\(\)](#) will compare keys during searching

Return values

<i>true</i>	< 0
-------------	-----

8.2.3.47 `bool kdb::Key::operator<= (const Key & other) const` `[inline]`

Compare the name of two keys.

Returns

a number less than, equal to or greater than zero if k1 is found, respectively, to be less than, to match, or be greater than k2.

The comparison is based on a strcmp of the keynames, and iff they match a strcmp of the owner will be used to distinguish. If even this matches the keys are found to be exactly the same and 0 is returned. These two keys can't be used in the same [KeySet](#).

[keyCmp\(\)](#) defines the sorting order for a [KeySet](#).

The following 3 points are the rules for null values. They only take account when none of the preceding rules matched.

- A null pointer will be found to be smaller than every other key. If both are null pointers, 0 is returned.
- A null name will be found to be smaller than every other name. If both are null names, 0 is returned.
- No owner will be found to be smaller then every other owner. If both don't have a owner, 0 is returned.

Note

the owner will only be used if the names are equal.

Often is enough to know if the other key is less then or greater then the other one. But Sometimes you need more precise information, see [keyRel\(\)](#).

Given any Keys k1 and k2 constructed with [keyNew\(\)](#), following equation hold true:

```
// keyCmp(0,0) == 0
// keyCmp(k1,0) == 1
// keyCmp(0,k2) == -1
```

You can write similar equation for the other rules.

Here are some more examples with equation:

```
Key *k1 = keyNew("user/a", KEY_END);
Key *k2 = keyNew("user/b", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0

Key *k1 = keyNew("user/a", KEY_OWNER, "markus", KEY_END);
Key *k2 = keyNew("user/a", KEY_OWNER, "max", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0
```

Warning

Do not try to strcmp the [keyName\(\)](#) yourself because the used strcmp implementation is allowed to differ from simple ascii comparison.

Parameters

<i>k1</i>	the first key object to compare with
<i>k2</i>	the second key object to compare with

See Also

[ksAppendKey\(\)](#), [ksAppend\(\)](#) will compare keys when appending
[ksLookup\(\)](#) will compare keys during searching

Return values

<i>true</i>	≤ 0
-------------	----------

8.2.3.48 Key & kdb::Key::operator= (ckdb::Key * *k*) [inline]

Assign a C key.

Will call del() on the old key.

8.2.3.49 Key & kdb::Key::operator= (const Key & *k*) [inline]

Assign a key.

Will call del() on the old key.

8.2.3.50 Key & kdb::Key::operator= (const std::string & *newName*) [inline]

Assign the name of a key.

See Also

[keySetName](#)

8.2.3.51 Key & kdb::Key::operator= (const char * *newName*) [inline]

(const std::string &newName)

Assign a C key.

Will call del() on the old key. (const std::string &newName)

8.2.3.52 bool kdb::Key::operator== (const Key & *k*) const [inline]

Compare the name of two keys.

Returns

a number less than, equal to or greater than zero if k1 is found, respectively, to be less than, to match, or be greater than k2.

The comparison is based on a strcmp of the keynames, and iff they match a strcmp of the owner will be used to distinguish. If even this matches the keys are found to be exactly the same and 0 is returned. These two keys can't be used in the same [KeySet](#).

[keyCmp\(\)](#) defines the sorting order for a [KeySet](#).

The following 3 points are the rules for null values. They only take account when none of the preceding rules matched.

- A null pointer will be found to be smaller than every other key. If both are null pointers, 0 is returned.
- A null name will be found to be smaller than every other name. If both are null names, 0 is returned.
- No owner will be found to be smaller than every other owner. If both don't have a owner, 0 is returned.

Note

the owner will only be used if the names are equal.

Often is enough to know if the other key is less then or greater then the other one. But Sometimes you need more precise information, see [keyRel\(\)](#).

Given any Keys k1 and k2 constructed with [keyNew\(\)](#), following equation hold true:

```
// keyCmp(0,0) == 0
// keyCmp(k1,0) == 1
// keyCmp(0,k2) == -1
```

You can write similar equation for the other rules.

Here are some more examples with equation:

```
Key *k1 = keyNew("user/a", KEY_END);
Key *k2 = keyNew("user/b", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0

Key *k1 = keyNew("user/a", KEY_OWNER, "markus", KEY_END);
Key *k2 = keyNew("user/a", KEY_OWNER, "max", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0
```

Warning

Do not try to strcmp the [keyName\(\)](#) yourself because the used strcmp implementation is allowed to differ from simple ascii comparison.

Parameters

<i>k1</i>	the first key object to compare with
<i>k2</i>	the second key object to compare with

See Also

[ksAppendKey\(\)](#), [ksAppend\(\)](#) will compare keys when appending
[ksLookup\(\)](#) will compare keys during searching

Return values

<i>true</i>	<code>== 0</code>
-------------	-------------------

8.2.3.53 bool kdb::Key::operator> (const Key & other) const [inline]

Compare the name of two keys.

Returns

a number less than, equal to or greater than zero if k1 is found, respectively, to be less than, to match, or be greater than k2.

The comparison is based on a strcmp of the keynames, and iff they match a strcmp of the owner will be used to distinguish. If even this matches the keys are found to be exactly the same and 0 is returned. These two keys can't be used in the same [KeySet](#).

[keyCmp\(\)](#) defines the sorting order for a [KeySet](#).

The following 3 points are the rules for null values. They only take account when none of the preceding rules matched.

- A null pointer will be found to be smaller than every other key. If both are null pointers, 0 is returned.
- A null name will be found to be smaller than every other name. If both are null names, 0 is returned.
- No owner will be found to be smaller then every other owner. If both don't have a owner, 0 is returned.

Note

the owner will only be used if the names are equal.

Often is enough to know if the other key is less then or greater then the other one. But Sometimes you need more precise information, see [keyRel\(\)](#).

Given any Keys k1 and k2 constructed with [keyNew\(\)](#), following equation hold true:

```
// keyCmp(0,0) == 0
// keyCmp(k1,0) == 1
// keyCmp(0,k2) == -1
```

You can write similar equation for the other rules.

Here are some more examples with equation:

```
Key *k1 = keyNew("user/a", KEY_END);
Key *k2 = keyNew("user/b", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0

Key *k1 = keyNew("user/a", KEY_OWNER, "markus", KEY_END);
Key *k2 = keyNew("user/a", KEY_OWNER, "max", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0
```


Warning

Do not try to strcmp the [keyName\(\)](#) yourself because the used strcmp implementation is allowed to differ from simple ascii comparison.

Parameters

<i>k1</i>	the first key object to compare with
<i>k2</i>	the second key object to compare with

See Also

[ksAppendKey\(\)](#), [ksAppend\(\)](#) will compare keys when appending
[ksLookup\(\)](#) will compare keys during searching

Return values

<i>true</i>	> 0
-------------	-----

8.2.3.54 `bool kdb::Key::operator>= (const Key & other) const [inline]`

Compare the name of two keys.

Returns

a number less than, equal to or greater than zero if k1 is found, respectively, to be less than, to match, or be greater than k2.

The comparison is based on a strcmp of the keynames, and iff they match a strcmp of the owner will be used to distinguish. If even this matches the keys are found to be exactly the same and 0 is returned. These two keys can't be used in the same [KeySet](#).

[keyCmp\(\)](#) defines the sorting order for a [KeySet](#).

The following 3 points are the rules for null values. They only take account when none of the preceding rules matched.

- A null pointer will be found to be smaller than every other key. If both are null pointers, 0 is returned.
- A null name will be found to be smaller than every other name. If both are null names, 0 is returned.
- No owner will be found to be smaller than every other owner. If both don't have a owner, 0 is returned.

Note

the owner will only be used if the names are equal.

Often is enough to know if the other key is less then or greater then the other one. But Sometimes you need more precise information, see [keyRel\(\)](#).

Given any Keys k1 and k2 constructed with [keyNew\(\)](#), following equation hold true:

```
// keyCmp(0,0) == 0
// keyCmp(k1,0) == 1
// keyCmp(0,k2) == -1
```

You can write similar equation for the other rules.

Here are some more examples with equation:

```

Key *k1 = keyNew("user/a", KEY_END);
Key *k2 = keyNew("user/b", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0

Key *k1 = keyNew("user/a", KEY_OWNER, "markus", KEY_END
);
Key *k2 = keyNew("user/a", KEY_OWNER, "max", KEY_END);

// keyCmp(k1,k2) < 0
// keyCmp(k2,k1) > 0

```

Warning

Do not try to strcmp the [keyName\(\)](#) yourself because the used strcmp implementation is allowed to differ from simple ascii comparison.

Parameters

<i>k1</i>	the first key object to compare with
<i>k2</i>	the second key object to compare with

See Also

[ksAppendKey\(\)](#), [ksAppend\(\)](#) will compare keys when appending
[ksLookup\(\)](#) will compare keys during searching

Return values

<i>true</i>	≥ 0
-------------	----------

8.2.3.55 `ckdb::Key * kdb::Key::release () [inline]`

Passes out the raw key pointer.

Note

that the ownership is moved outside.

The key will stay empty.

8.2.3.56 `void kdb::Key::rewindMeta () const [inline]`

Rewind the internal iterator to first meta data.

Use it to set the cursor to the beginning of the [Key](#) Meta Infos. [keyCurrentMeta\(\)](#) will then always return NULL afterwards. So you want to [keyNextMeta\(\)](#) first.

```

Key *key;
const Key *meta;

keyRewindMeta (key);
while ((meta = keyNextMeta (key)) != 0)
{
    printf ("name: %s, value: %s", keyName(meta), (const char*)
        keyValue(meta));
}

```

Parameters

<i>key</i>	the key object to work with
------------	-----------------------------

Returns

- 0 on success
- 0 if there is no meta information for that key ([keyNextMeta\(\)](#) will always return 0 in that case)
- 1 on NULL pointer

See Also

[keyNextMeta\(\)](#), [keyCurrentMeta\(\)](#)
[ksRewind\(\)](#) for pedant in iterator interface of [KeySet](#)

See Also

[nextMeta\(\)](#), [currentMeta\(\)](#)

8.2.3.57 `template<class T> void kdb::Key::set (T x) [inline]`

Set a key value.

Set the value for *key* as *newStringValue*.

The function will allocate and save a private copy of *newStringValue*, so the parameter can be freed after the call.

String values will be saved in backend storage, when *kdbSetKey()* will be called, in UTF-8 universal encoding, regardless of the program's current encoding, when *iconv* plugin is present.

Note

The type will be set to `KEY_TYPE_STRING`. When the type of the key is already a string type it won't be changed.

Parameters

<i>key</i>	the key to set the string value
<i>newStringValue</i>	NULL-terminated text string to be set as <i>key</i> 's value

Returns

the number of bytes actually saved in private struct including final NULL

Return values

<i>1</i>	if <i>newStringValue</i> is a NULL pointer, this will make the string empty (string only containing null termination)
<i>-1</i>	if <i>key</i> is a NULL pointer

See Also

[keyGetString\(\)](#), [keyValue\(\)](#), [keyString\(\)](#)

This method tries to deserialize the string to the given type.

8.2.3.58 `void kdb::Key::setBaseName (const std::string & baseName) [inline]`

Sets a base name for a key.

Sets *baseName* as the new basename for *key*.

All text after the last ' / ' in the *key* keyname is erased and *baseName* is appended.

So lets suppose *key* has name "system/dir1/dir2/key1". If *baseName* is "key2", the resulting key name will be "system/dir1/dir2/key2". If *baseName* is empty or NULL, the resulting key name will be "system/dir1/dir2".

Warning

You should not change a keys name once it belongs to a keyset. See `ksSort()` for more information.

TODO: does not work with .. and .

Parameters

<i>key</i>	the key object to work with
<i>baseName</i>	the string used to overwrite the basename of the key

Returns

the size in bytes of the new key name
-1 on NULL pointers

See Also

[keyAddBaseName\(\)](#)
[keySetName\(\)](#) to [set](#) a new name

Exceptions

<i>KeyInvalidName</i>	if the name is not valid
-----------------------	--------------------------

8.2.3.59 `size_t kdb::Key::setBinary (const void * newBinary, size_t dataSize) [inline]`

Set the value of a key as a binary.

A private copy of *newBinary* will allocated and saved inside *key*, so the parameter can be deallocated after the call.

Binary values might be encoded in another way then string values depending on the plugin. Typically character encodings should not take place on binary data. Consider using a string key instead.

When *newBinary* is a NULL pointer the binary will be freed and 0 will be returned.

Note

The meta data "binary" will be set to mark that the key is binary from now on. When the key is already binary the meta data won't be changed. This will only happen in the successful case, but not when -1 is returned.

Parameters

<i>key</i>	the object on which to set the value
<i>newBinary</i>	is a pointer to any binary data or NULL to free the previous set data
<i>dataSize</i>	number of bytes to copy from <i>newBinary</i>

Returns

the number of bytes actually copied to internal struct storage
 0 when the internal binary was freed and is now a null pointer
 -1 if key is a NULL pointer
 -1 when dataSize is 0 (but newBinary not NULL) or larger than SSIZE_MAX

See Also

[keyGetBinary\(\)](#)
[keyIsBinary\(\)](#) to check if the type is binary
[keyGetString\(\)](#) and [keySetString\(\)](#) as preferred alternative to binary

8.2.3.60 `template<class T > void kdb::Key::setMeta (const std::string & metaName, T x) [inline]`

Set a new Meta-Information.

Will set a new Meta-Information pair consisting of metaName and newMetaString.

Will add a new Pair for Meta-Information if metaName was not added up to now.

It will modify a existing Pair of Meta-Information if the the metaName was inserted already.

It will remove a meta information if newMetaString is 0.

Parameters

<i>key</i>	the key object to work with
<i>metaName</i>	the name of the meta information where you want to change the value
<i>newMetaString</i>	the new value for the meta information

Returns

-1 on error if key or metaName is 0, out of memory or names are not valid
 0 if the Meta-Information for metaName was removed
 size (>0) of newMetaString if Meta-Information was successfully added

See Also

[keyGetMeta\(\)](#)

See Also

[getMeta\(\)](#), [copyMeta\(\)](#), [copyAllMeta\(\)](#)

8.2.3.61 `void kdb::Key::setName (const std::string & newName) [inline]`

Set a new name to a key.

A valid name is of the forms:

- system/something
- user/something
- user:username/something

The last form has explicitly set the owner, to let the library know in which user folder to save the key. A owner is a user name. If it is not defined (the second form) current user is used.

You should always follow the guidelines for key tree structure creation.

A private copy of the key name will be stored, and the `newName` parameter can be freed after this call.

..., . and / will be handled correctly. A valid name will be build out of the (valid) name what you pass, e.g. `user///sw/./sw//././MyApp -> user/sw/MyApp`

On invalid names, NULL or "" the name will be "" afterwards.

Warning

You shall not change a key name once it belongs to a keyset.

Return values

size	in bytes of this new key name including ending NULL
0	if newName is an empty string or a NULL pointer (name will be empty afterwards)
-1	if newName is invalid (name will be empty afterwards)

Parameters

<i>key</i>	the key object to work with
<i>newName</i>	the new key name

See Also

[keyNew\(\)](#), [keySetOwner\(\)](#)
[keyGetName\(\)](#), [keyGetFullName\(\)](#), [keyName\(\)](#)
[keySetBaseName\(\)](#), [keyAddBaseName\(\)](#) to manipulate a name

Exceptions

<i>KeyInvalidName</i>	if the name is not valid
-----------------------	--------------------------

8.2.3.62 `void kdb::Key::setString (std::string newString) [inline]`

Set the value for `key` as `newStringValue`.

The function will allocate and save a private copy of `newStringValue`, so the parameter can be freed after the call.

String values will be saved in backend storage, when `kdbSetKey()` will be called, in UTF-8 universal encoding, regardless of the program's current encoding, when `iconv` plugin is present.

Note

The type will be set to `KEY_TYPE_STRING`. When the type of the key is already a string type it won't be changed.

Parameters

<i>key</i>	the key to set the string value
<i>newStringValue</i>	NULL-terminated text string to be set as <code>key</code> 's value

Returns

the number of bytes actually saved in private struct including final NULL

Return values

1	if newStringValue is a NULL pointer, this will make the string empty (string only containing null termination)
-1	if key is a NULL pointer

See Also

[keyGetString\(\)](#), [keyValue\(\)](#), [keyString\(\)](#)

The documentation for this class was generated from the following file:

- `src/bindings/cpp/include/key.hpp`

8.3 kdb::KeySet Class Reference

A keyset holds together a set of keys.

```
#include <keyset.hpp>
```

Public Member Functions

- [KeySet](#) ()
- [KeySet](#) (ckdb::KeySet *k)
- [KeySet](#) (const [KeySet](#) &other)
- [KeySet](#) (size_t alloc, va_list ap)
Create a new keyset.
- [KeySet](#) (size_t alloc,...)
Create a new keyset.
- [~KeySet](#) ()
Deconstruct a keyset.
- ckdb::KeySet * [release](#) ()
- ckdb::KeySet * [getKeySet](#) () const
Passes out the raw keyset pointer.
- void [setKeySet](#) (ckdb::KeySet *k)
Take ownership of passed keyset.
- [KeySet](#) & [operator=](#) ([KeySet](#) const &other)
- size_t [size](#) () const
The size of the keyset.
- ckdb::KeySet * [dup](#) () const
Duplicate a keyset.
- void [copy](#) (const [KeySet](#) &other)
Copy a keyset.
- void [clear](#) ()
Clear the keyset.
- ssize_t [append](#) (const [Key](#) &toAppend)
append a key

- `ssize_t append (const KeySet &toAppend)`
append a keyset
- `Key head () const`
- `Key tail () const`
- `void rewind () const`
- `Key next () const`
- `Key current () const`
- `void setCursor (cursor_t cursor) const`
- `cursor_t getCursor () const`
- `Key pop ()`
- `KeySet cut (Key k)`
- `Key lookup (const Key &k, const option_t options=KDB_O_NONE) const`
- `Key lookup (std::string const &name, const option_t options=KDB_O_NONE) const`
Lookup a key by name.

8.3.1 Detailed Description

A keyset holds together a set of keys.

Methods to manipulate KeySets. A KeySet is a sorted set of keys. So the correct name actually would be KeyMap.

With `ksNew()` you can create a new KeySet.

You can add keys with `ksAppendKey()` in the keyset. Using `ksAppend()` you can append a whole keyset.

Note

Because the key is not copied, also the pointer to the current metadata `keyNextMeta()` will be shared.

`ksGetSize()` tells you the current size of the keyset.

With `ksRewind()` and `ksNext()` you can navigate through the keyset. Don't expect any particular order, but it is assured that you will get every key of the set.

KeySets have an `internal cursor`. This is used for `ksLookup()` and `kdbSet()`.

KeySet has a fundamental meaning inside elektra. It makes it possible to get and store many keys at once inside the database. In addition to that the class can be used as high level datastructure in applications. With `ksLookup-ByName()` it is possible to fetch easily specific keys out of the list of keys.

You can easily create and iterate keys:

```
#include <kdb.h>

// create a new keyset with 3 keys
// with a hint that about 20 keys will be inside
KeySet *myConfig = ksNew(20,
    keyNew ("user/name1", 0),
    keyNew ("user/name2", 0),
    keyNew ("user/name3", 0),
```



```

        KS_END);
// append a key in the keyset
ksAppendKey(myConfig, keyNew("user/name4", 0));

Key *current;
ksRewind(myConfig);
while ((current=ksNext(myConfig))!=0)
{
    printf("Key name is %s.\n", keyName (current));
}
ksDel (myConfig); // delete keyset and all keys appended

```

Invariant

always holds an underlying elektra keyset.

Note

that the cursor is mutable, so it might be changed even in const functions as described.

8.3.2 Constructor & Destructor Documentation

8.3.2.1 kdb::KeySet::KeySet () [inline]

Creates a new empty keyset with no keys

Allocate, initialize and return a new [KeySet](#) object.

Objects created with [ksNew\(\)](#) must be destroyed with [ksDel\(\)](#).

You can use a various long list of parameters to preload the keyset with a list of keys. Either your first and only parameter is 0 or your last parameter must be [KEY_END](#).

So, terminate with [ksNew\(0\)](#) or [ksNew\(20, ..., KS_END\)](#)

For most uses

```

KeySet *keys = ksNew(0);
// work with it
ksDel (keys);

```

goes ok, the alloc size will be 16, defined in [kdbprivate.h](#). The alloc size will be doubled whenever size reaches alloc size, so it also performs out large keysets.

But if you have any clue how large your keyset may be you should read the next statements.

If you want a keyset with length 15 (because you know of your application that you normally need about 12 up to 15 keys), use:

```

KeySet * keys = ksNew (15,
    keyNew ("user/sw/app/fixedConfiguration/key01", KEY_SWITCH_VALUE,
"value01", 0),
    keyNew ("user/sw/app/fixedConfiguration/key02", KEY_SWITCH_VALUE,
"value02", 0),
    keyNew ("user/sw/app/fixedConfiguration/key03", KEY_SWITCH_VALUE,
"value03", 0),
    // ...
    keyNew ("user/sw/app/fixedConfiguration/key15", KEY_SWITCH_VALUE,
"value15", 0),
    KS_END);
// work with it
ksDel (keys);

```

If you start having 3 keys, and your application needs approximately 200-500 keys, you can use:

```

KeySet * config = ksNew (500,
    keyNew ("user/sw/app/fixedConfiguration/key1", KEY_SWITCH_VALUE,
"value1", 0),
    keyNew ("user/sw/app/fixedConfiguration/key2", KEY_SWITCH_VALUE,
"value2", 0),

```

```

    keyNew ("user/sw/app/fixedConfiguration/key3", KEY_SWITCH_VALUE,
    "value3", 0),
    KS_END); // don't forget the KS_END at the end!
// work with it
ksDel (config);

```

Alloc size is 500, the size of the keyset will be 3 after ksNew. This means the keyset will reallocate when appending more than 497 keys.

The main benefit of taking a list of variant length parameters is to be able to have one C-Statement for any possible [KeySet](#).

Due to ABI compatibility, the [KeySet](#) structure is only declared in kdb.h, and not defined. So you can only declare pointers to KeySets in your program. See <http://tldp.org/HOWTO/Program-Library-HOWTO/shared-libraries.html#AEN135>

See Also

[ksDel\(\)](#) to free the [KeySet](#) afterwards
[ksDup\(\)](#) to duplicate an existing [KeySet](#)

Parameters

<i>alloc</i>	gives a hint for the size how many Keys may be stored initially
--------------	---

Returns

a ready to use [KeySet](#) object
0 on memory error

8.3.2.2 kdb::KeySet::KeySet (ckdb::KeySet * k) [inline]

Takes ownership of keyset!

Keyset will be destroyed at destructor you cant continue to use keyset afterwards!

Use [KeySet::release\(\)](#) to avoid destruction.

Parameters

<i>k</i>	the keyset to take the ownership from
----------	---------------------------------------

See Also

[release\(\)](#)
[setKeySet\(\)](#)

8.3.2.3 kdb::KeySet::KeySet (const KeySet & other) [inline]

Duplicate a keyset.

This keyset will be a duplicate of the other afterwards.

Note

that they still reference to the same Keys, so if you change key values also the keys in the original keyset will be changed.

See Also

[dup](#)**8.3.2.4** `kdb::KeySet::KeySet (size_t alloc, va_list ap) [inline], [explicit]`

Create a new keyset.

Parameters

<i>alloc</i>	minimum number of keys to allocate
<i>ap</i>	variable arguments list

8.3.2.5 `kdb::KeySet::KeySet (size_t alloc, ...) [inline], [explicit]`

Create a new keyset.

Parameters

<i>alloc</i>	minimum number of keys to allocate
<i>...</i>	variable argument list

8.3.2.6 `kdb::KeySet::~~KeySet () [inline]`

Deconstruct a keyset.

A destructor for [KeySet](#) objects.

Cleans all internal dynamic attributes, decrement all reference pointers to all keys and then [keyDel\(\)](#) all contained Keys, and free(s) the release the [KeySet](#) object memory (that was previously allocated by [ksNew\(\)](#)).

Parameters

<i>ks</i>	the keyset object to work with
-----------	--------------------------------

Returns

0 when the keyset was freed
-1 on null pointer

See Also

[ksNew\(\)](#)**8.3.3** Member Function Documentation**8.3.3.1** `ssize_t kdb::KeySet::append (const Key & toAppend) [inline]`

append a key

Parameters

<i>toAppend</i>	key to append
-----------------	---------------

Returns

number of keys in the keyset

Appends a [Key](#) to the end of `ks`.

A pointer to the key will be stored, and not a private copy. So a future [ksDel\(\)](#) on `ks` may [keyDel\(\)](#) the `toAppend` object, see [keyGetRef\(\)](#).

The reference counter of the key will be incremented, and thus `toAppend` is not const.

Note

Because the key is not copied, also the pointer to the current metadata [keyNextMeta\(\)](#) will be shared.

If the keyname already existed, it will be replaced with the new key.

The [KeySet](#) internal cursor will be set to the new key.

Returns

the size of the [KeySet](#) after insertion
 -1 on NULL pointers
 -1 if insertion failed, the key will be deleted then.

Parameters

<i>ks</i>	KeySet that will receive the key
<i>toAppend</i>	Key that will be appended to <code>ks</code>

See Also

[ksAppend\(\)](#), [keyNew\(\)](#), [ksDel\(\)](#)
[keyIncRef\(\)](#)

8.3.3.2 `ssize_t kdb::KeySet::append (const KeySet & toAppend) [inline]`

append a keyset

Parameters

<i>toAppend</i>	keyset to append
-----------------	------------------

Returns

number of keys in the keyset

Append all `toAppend` contained keys to the end of the `ks`.

`toAppend` [KeySet](#) will be left unchanged.

If a key is both in `toAppend` and `ks`, the [Key](#) in `ks` will be overridden.

Note

Because the key is not copied, also the pointer to the current metadata [keyNextMeta\(\)](#) will be shared.

Postcondition

Sorted [KeySet](#) `ks` with all keys it had before and additionally the keys from `toAppend`

Returns

the size of the [KeySet](#) after transfer
 -1 on NULL pointers

Parameters

<i>ks</i>	the KeySet that will receive the keys
<i>toAppend</i>	the KeySet that provides the keys that will be transfered

See Also

[ksAppendKey\(\)](#)

8.3.3.3 void kdb::KeySet::clear () [inline]

Clear the keyset.

Keyset will have no keys afterwards.

8.3.3.4 void kdb::KeySet::copy (const KeySet & other) [inline]

Copy a keyset.

Parameters

<i>other</i>	other keyset to copy
--------------	----------------------

Copy a keyset.

Most often you may want a duplicate of a keyset, see [ksDup\(\)](#) or append keys, see [ksAppend\(\)](#). But in some situations you need to copy a keyset to a existing keyset, for that this function exists.

You can also use it to clear a keyset when you pass a NULL pointer as *source*.

Note that all keys in *dest* will be deleted. Afterwards the content of the source will be added to the destination and the [ksCurrent\(\)](#) is set properly in *dest*.

A flat copy is made, so the keys will not be duplicated, but there reference counter is updated, so both keysets need to be [ksDel\(\)](#).

Note

Because the key is not copied, also the pointer to the current metadata [keyNextMeta\(\)](#) will be shared.

```
int f (KeySet *ks)
{
    KeySet *c = ksNew (20, ..., KS_END);
    // c receives keys
    ksCopy (ks, c); // pass the keyset to the caller

    ksDel (c);
    // caller needs to ksDel (ks)
}
```

Parameters

<i>source</i>	has to be an initialized source KeySet or NULL
<i>dest</i>	has to be an initialized KeySet where to write the keys

Returns

1 on success
 0 if dest was cleared successfully (source is NULL)
 -1 on NULL pointer

See Also

[ksNew\(\)](#), [ksDel\(\)](#), [ksDup\(\)](#)
[keyCopy\(\)](#) for copying keys

8.3.3.5 Key kdb::KeySet::current () const [inline]

Return the current [Key](#).

The pointer is NULL if you reached the end or after [ksRewind\(\)](#).

Note

You must not delete the key or change the key, use [ksPop\(\)](#) if you want to delete it.

Parameters

ks	the keyset object to work with
--------------------	--------------------------------

Returns

pointer to the [Key](#) pointed by [ks](#)'s cursor
 0 on NULL pointer

See Also

[ksNext\(\)](#), [ksRewind\(\)](#)

8.3.3.6 KeySet kdb::KeySet::cut (Key k) [inline]

Cuts out a keyset at the cutpoint.

Searches for the cutpoint inside the [KeySet](#) [ks](#). If found it cuts out everything which is below (see [keyIsBelow\(\)](#)) this key. If not found an empty keyset is returned.

The cursor will stay at the same key as it was before. If the cursor was inside the region of cutted (moved) keys, the cursor will be set to the key before the cutpoint.

Returns

a new allocated [KeySet](#) which needs to be deleted with [ksDel\(\)](#). The keyset consists of all keys (of the original keyset [ks](#)) below the cutpoint. If the key cutpoint exists, it will also be appended.

Return values

0	on null pointers, no key name or allocation problems
---	--

Parameters

<i>ks</i>	the keyset to cut. It will be modified by removing all keys below the cutpoint. The cutpoint itself will also be removed.
<i>cutpoint</i>	the point where to cut out the keyset

8.3.3.7 `ckdb::KeySet * kdb::KeySet::dup () const` `[inline]`

Duplicate a keyset.

Returns

a copy of the keys

Return a duplicate of a keyset.

Objects created with `ksDup()` must be destroyed with `ksDel()`.

Memory will be allocated as needed for dynamic properties, so you need to `ksDel()` the returned pointer.

A flat copy is made, so the keys will not be duplicated, but there reference counter is updated, so both keysets need `ksDel()`.

Parameters

<i>source</i>	has to be an initialised source KeySet
---------------	--

Returns

a flat copy of source on success

0 on NULL pointer

See Also

[ksNew\(\)](#), [ksDel\(\)](#)

[keyDup\(\)](#) for [Key](#) duplication

8.3.3.8 `cursor.t kdb::KeySet::getCursor () const` `[inline]`

Get the [KeySet](#) internal cursor.

Use it to get the cursor of the actual position.

Warning

Cursors are getting invalid when the key was `ksPop()`ed or `ksLookup()` with `KDB_O_POP` was used.

8.3.3.9 `ckdb::KeySet * kdb::KeySet::getKeySet () const` `[inline]`

Passes out the raw keyset pointer.

Returns

pointer to internal ckdb [KeySet](#)

See Also

[release\(\)](#)

[setKeySet\(\)](#)

8.3.3.10 Key kdb::KeySet::head () const [inline]

Returns

alphabetical first key

Return the first key in the [KeySet](#).

The KeySets cursor will not be effected.

If [ksCurrent\(\)==ksHead\(\)](#) you know you are on the first key.

Parameters

<i>ks</i>	the keyset object to work with
-----------	--------------------------------

Returns

the first [Key](#) of a keyset

0 on NULL pointer or empty keyset

See Also

[ksTail\(\)](#) for the last [Key](#)

[ksRewind\(\)](#), [ksCurrent\(\)](#) and [ksNext\(\)](#) for iterating over the [KeySet](#)

8.3.3.11 Key kdb::KeySet::lookup (const Key & key, const option_t options = KDB_O_NONE) const [inline]

Look for a [Key](#) contained in *ks* that matches the name of the *key*.

Note

That the internal key cursor will point to the found key

8.3.3.12 Key kdb::KeySet::lookup (std::string const & name, const option_t options = KDB_O_NONE) const [inline]

Lookup a key by name.

Parameters

<i>name</i>	the name to look for
<i>options</i>	some options to pass

Returns

the found key

See Also

[lookup](#) (const [Key](#) &[Key](#), const [option_t](#) options)

Note

That the internal key cursor will point to the found key

8.3.3.13 Key kdb::KeySet::next () const [inline]

Returns the next [Key](#) in a [KeySet](#).

KeySets have an internal cursor that can be reset with [ksRewind\(\)](#). Every time [ksNext\(\)](#) is called the cursor is incremented and the new current [Key](#) is returned.

You'll get a NULL pointer if the key after the end of the [KeySet](#) was reached. On subsequent calls of [ksNext\(\)](#) it will still return the NULL pointer.

The `ks` internal cursor will be changed, so it is not const.

Note

You must not delete or change the key, use [ksPop\(\)](#) if you want to delete it.

Parameters

<code>ks</code>	the keyset object to work with
-----------------	--------------------------------

Returns

the new current [Key](#)
 0 when the end is reached
 0 on NULL pointer

See Also

[ksRewind\(\)](#), [ksCurrent\(\)](#)

8.3.3.14 KeySet & kdb::KeySet::operator= (KeySet const & other) [inline]

Duplicate a keyset.

This keyset will be a duplicate of the other afterwards.

Note

that they still reference to the same Keys, so if you change key values also the keys in the original keyset will be changed.

8.3.3.15 Key kdb::KeySet::pop () [inline]

Remove and return the last key of `ks`.

The reference counter will be decremented by one.

The KeySets cursor will not be effected if it did not point to the popped key.

Note

You need to [keyDel\(\)](#) the key afterwards, if you don't append it to another keyset. It has the same semantics like a key allocated with [keyNew\(\)](#) or [keyDup\(\)](#).

```
ks1=ksNew(0);
ks2=ksNew(0);

k1=keyNew("user/name", KEY_END); // ref counter 0
ksAppendKey(ks1, k1); // ref counter 1
ksAppendKey(ks2, k1); // ref counter 2
```

```

k1=ksPop (ks1); // ref counter 1
k1=ksPop (ks2); // ref counter 0, like after keyNew()

ksAppendKey(ks1, k1); // ref counter 1

ksDel (ks1); // key is deleted too
ksDel (ks2);
*
```

Returns

the last key of `ks`
 NULL if `ks` is empty or on NULL pointer

Parameters

<code>ks</code>	KeySet to work with
-----------------	-------------------------------------

See Also

[ksAppendKey\(\)](#), [ksAppend\(\)](#)
[commandList\(\)](#) for an example

8.3.3.16 `ckdb::KeySet * kdb::KeySet::release () [inline]`

If you dont want destruction of keyset at the end you can release the pointer.

8.3.3.17 `void kdb::KeySet::rewind () const [inline]`

Rewinds the [KeySet](#) internal cursor.

Use it to set the cursor to the beginning of the [KeySet](#). [ksCurrent\(\)](#) will then always return NULL afterwards. So you want to [ksNext\(\)](#) first.

```

ksRewind (ks);
while ((key = ksNext (ks)) !=0) {}
```

Parameters

<code>ks</code>	the keyset object to work with
-----------------	--------------------------------

Returns

0 on success
 -1 on NULL pointer

See Also

[ksNext\(\)](#), [ksCurrent\(\)](#)

8.3.3.18 `void kdb::KeySet::setCursor (cursor_t cursor) const [inline]`

Set the [KeySet](#) internal cursor.

Use it to set the cursor to a stored position. [ksCurrent\(\)](#) will then be the position which you got with.

Warning

Cursors may get invalid when the key was [ksPop\(\)](#)ed or [ksLookup\(\)](#) was used together with KDB_O_POP.

```
cursor_t cursor;
...
// key now in any position here
cursor = ksGetCursor (ks);
while ((key = keyNextMeta (ks)) != 0) {}
ksSetCursor (ks, cursor); // reset state
ksCurrent(ks); // in same position as before
```

An invalid cursor will set the keyset to its beginning like [ksRewind\(\)](#). When you set an invalid cursor [ksCurrent\(\)](#) is 0 and [ksNext\(\)](#) == [ksHead\(\)](#).

Parameters

<i>cursor</i>	the cursor to use
<i>ks</i>	the keyset object to work with

Returns

0 when the keyset is [ksRewind\(\)](#)ed
 1 otherwise
 -1 on NULL pointer

See Also

[ksNext\(\)](#), [ksGetCursor\(\)](#)

8.3.3.19 void kdb::KeySet::setKeySet (ckdb::KeySet * k) [inline]

Take ownership of passed keyset.

Parameters

<i>k</i>	the keyset to take ownership from
----------	-----------------------------------

See Also

[release\(\)](#)
[getKeySet\(\)](#)

8.3.3.20 size_t kdb::KeySet::size () const [inline]

The size of the keyset.

Returns

the number of keys in the keyset

8.3.3.21 Key kdb::KeySet::tail () const [inline]

Returns

alphabetical last key

Return the last key in the [KeySet](#).

The KeySets cursor will not be effected.

If [ksCurrent\(\)==ksTail\(\)](#) you know you are on the last key. [ksNext\(\)](#) will return a NULL pointer afterwards.

Parameters

ks	the keyset object to work with
--------------------	--------------------------------

Returns

the last [Key](#) of a keyset

0 on NULL pointer or empty keyset

See Also

[ksHead\(\)](#) for the first [Key](#)

[ksRewind\(\)](#), [ksCurrent\(\)](#) and [ksNext\(\)](#) for iterating over the [KeySet](#)

The documentation for this class was generated from the following file:

- `src/bindings/cpp/include/keyset.hpp`